

Supplementary Information: The Curation-Exploration Tradeoff

SCX

June 30, 2026

Contents

1	Full Experimental Results	1
1.1	AIN MLIP	1
1.2	CIFAR-10 Controlled Noise	1
1.3	MedMNIST	2
1.4	DrugBank Drug-Target	3
2	Sensitivity Analysis	4
2.1	Varying the Number of Experts M	4
2.2	Varying the Noise Rate η	4
2.3	Varying the Exploration Schedule $\eta(t)$	4
3	Comparison with Alternative Curation Strategies	5
3.1	Alternative Methods	5
3.2	Results	5
4	Self-Diagnosis: Bootstrap Stability Across Domains	6
5	Relationship to Paper I Theorems	6

1 Full Experimental Results

1.1 AIN MLIP

The SCX adaptive threshold (Tier 3) achieves the best force RMSE of 0.028 eV/Å, representing a 29–48% improvement over no cleaning (0.045 eV/Å) and a 32% improvement over loss-based cleaning (0.041 eV/Å). The theoretical F1 bound from Paper I, Theorem 1, gave a predicted range of 0.77–0.86 depending on state proportions; the empirical 0.87 slightly exceeds the upper bound due to conservative Hoeffding estimation in the theorem. The exact asymptotic (Theorem 4’), computed with $p_0 = 0.15$, $p_1 = 0.72$, $\eta = 0.14$, yields $\kappa = 0.332$ and a predicted asymptotic F1 of 0.89 for $M \rightarrow \infty$, consistent with $M = 8$ achieving 0.87.

1.2 CIFAR-10 Controlled Noise

The improvement grows with noise rate: at 40% noise, SCX deferred cleaning achieves 79.3% accuracy versus 65.8% for premature cleaning—an 11.2 percentage-point gap. This widening gap

Table 1: Complete AIN MLIP results across all noise detection methods. Metrics: force RMSE (eV/Å), energy MAE (meV/atom), noise detection F1. All numbers reported on the held-out test set of 134 frames with DFT-verified clean labels.

Method	Force RMSE	Energy MAE	F1	Frames removed
No cleaning (baseline)	0.045 ± 0.002	3.12 ± 0.15	—	0
Random removal (14%)	0.044 ± 0.003	3.05 ± 0.18	0.12 ± 0.04	74
Loss-based (top 14% by loss)	0.041 ± 0.003	2.89 ± 0.20	0.45 ± 0.06	74
Confident learning [2]	0.038 ± 0.004	2.71 ± 0.22	0.53 ± 0.07	78
SCX (Tier 1, Hoeffding)	0.031 ± 0.003	2.24 ± 0.17	0.81 ± 0.05	68
SCX (Tier 2, Chernoff/KL)	0.030 ± 0.003	2.19 ± 0.18	0.84 ± 0.04	66
SCX (Tier 3, adaptive θ_{opt})	0.028 ± 0.002	2.08 ± 0.15	0.87 ± 0.03	62
SCX (Tier 4, bootstrap)	0.028 ± 0.002	2.06 ± 0.14	0.87 ± 0.03	63

Table 2: Complete CIFAR-10 results at three noise rates. Metrics: test accuracy (%), noise detection F1, precision, recall. $M = 5$ ResNet-18 experts, adaptive threshold (Tier 3).

Noise rate η	Method	Test accuracy (%)	F1	Precision/Recall
10%	No cleaning	83.7 ± 0.4	—	—
	Premature cleaning (loss)	80.1 ± 0.6	0.45 ± 0.03	0.31/0.82
	SCX deferred	85.1 ± 0.3	0.62 ± 0.04	0.49/0.84
	Gain vs. premature	+5.0	+0.17	
20%	No cleaning	78.3 ± 0.5	—	—
	Premature cleaning (loss)	76.2 ± 0.5	0.38 ± 0.03	0.27/0.78
	SCX deferred	83.6 ± 0.4	0.55 ± 0.04	0.44/0.81
	Gain vs. premature	+7.4	+0.17	
40%	No cleaning	68.1 ± 0.7	—	—
	Premature cleaning (loss)	65.8 ± 0.8	0.29 ± 0.04	0.21/0.73
	SCX deferred	79.3 ± 0.5	0.48 ± 0.05	0.38/0.77
	Gain vs. premature	+11.2	+0.19	

is consistent with the tradeoff analysis: as noise increases, premature cleaning removes more hard samples (false positives), while SCX’s multi-expert consistency better distinguishes noise from hardness.

Ablation: number of experts M . Figure S1 shows F1 as a function of M at 20% noise. The improvement saturates at $M \approx 8$, with $M = 5$ achieving 87% of the $M = 20$ F1. This confirms the diminishing-returns prediction from Paper I (Theorem 4’): $1 - \text{F1} \propto e^{-M\kappa}/\sqrt{M}$.

Figure 1: F1 vs. number of experts M for CIFAR-10 at 20% noise. SCX with adaptive threshold (solid) and theoretical bound from Paper I Theorem 4’ (dashed, using $\kappa = 0.169$, $p_0 = 0.45$, $p_1 = 0.95$, $\eta = 0.20$). Shaded region: ± 1 standard deviation over 5 random seeds.

1.3 MedMNIST

DermaMNIST exemplifies the weak-feature regime (Theorem 2). The bootstrap stability $S(\Phi, K) = 0.45$ falls well below the 0.7 threshold, signalling that the tradeoff framework cannot achieve large gains. The feature weakness $\varepsilon_\phi = 0.52$ exceeds the 0.5 bound from Theorem 2, predicting SCX F1

Table 3: Complete MedMNIST results. DermaMNIST uses SimpleCNN; BloodMNIST uses ResNet-18.

Task	Metric	Value
DermaMNIST noise detection	ROC-AUC (loss baseline)	0.65 ± 0.04
	ROC-AUC (SCX)	0.78 ± 0.05
	Bootstrap stability $S(\Phi, K)$	0.45 ± 0.08
	Feature weakness ε_ϕ	0.52 ± 0.06
BloodMNIST routing	Best single expert (%)	91.2 ± 0.3
	Uniform ensemble (%)	92.5 ± 0.2
	SCX-Routing (%)	93.2 ± 0.2
	SCX-Compress (20% removal) (%)	92.9 ± 0.3
	SCX-Compress (40% removal) (%)	91.5 ± 0.4
PathMNIST compression	Baseline accuracy (%)	94.8 ± 0.2
	SCX-Compress (20% removal) (%)	94.5 ± 0.3
	SCX-Compress (40% removal) (%)	93.2 ± 0.4

$\leq$ baseline F1 $+\sqrt{0.52/2} \cdot C_F \approx$ baseline $+ 0.51 \cdot C_F$. With $C_F \leq 2$, the maximum gain is ~ 0.07 in F1, consistent with the observed improvement from 0.65 to 0.78 (ROC-AUC, not F1, but the bound applies directionally).

1.4 DrugBank Drug-Target

Table 4: DrugBank four-layer targetome screening results. Metrics: precision@k, recall@k, targetome coverage (fraction of targets with at least one high-confidence prediction).

Method	Precision@0.1	Recall@0.1	Coverage	Layers
Layer 1 only (knowledge base)	0.92	0.05	5%	1
Layer 2 only (similarity)	0.81	0.12	18%	1
Layer 3 only (deep ML)	0.63	0.31	52%	1
Layer 4 only (docking)	0.71	0.04	3%	1
Uniform ensemble (all 4)	0.65	0.18	35%	4
Loss-based weighting	0.58	0.22	41%	4
SCX consistency scoring	0.74	0.28	58%	4
SCX per-state (chemical cluster)	0.76	0.31	62%	4

The SCX per-state approach (discovering chemical clusters via ECFP4 features and applying per-cluster adaptive thresholds) achieves the best precision@0.1 (0.76) and highest targetome coverage (62%). The improvement over uniform ensemble is +0.11 in precision@0.1 and +27 percentage points in coverage. The four-layer pipeline naturally embodies the Curation-Exploration Tradeoff: Layer 1 is highly curated (high precision, low recall, low η), Layer 3 explores broadly (lower precision, high recall, high η), and SCX consistency scoring across layers identifies which predictions from which layers to trust for each molecule class.

2 Sensitivity Analysis

2.1 Varying the Number of Experts M

Figure S2 shows noise detection F1 as a function of M for all four domains. The characteristic exponential saturation is observed in all cases, with the saturation point shifting based on the separation gap Δ :

- **AIN MLIP** ($\Delta \approx 0.30$, $K = 1$ regression): Saturation at $M \approx 6$, $F1_\infty \approx 0.89$.
- **CIFAR-10** ($\Delta \approx 0.25$, $K = 10$): Saturation at $M \approx 8$, $F1_\infty \approx 0.65$.
- **DermaMNIST** ($\Delta \approx 0.12$, $K = 7$): Near-linear growth, $M = 12$ yields $F1 \approx 0.50$, still far from saturation. This is consistent with Theorem 2 (weak features).
- **DrugBank** (Δ varies by chemical cluster, 0.15–0.40): Saturation at $M \approx 5$, but the plateau value varies by cluster (0.45–0.85).

The Chernoff information κ (Paper I, Lemma C) governs the saturation rate. For AIN ($p_0 = 0.15$, $p_1 = 0.72$), $\kappa \approx 0.332$, predicting $e^{-8 \times 0.332} \approx 0.070$, while for DermaMNIST ($p_0 = 0.35$, $p_1 = 0.59$), $\kappa \approx 0.053$, predicting $e^{-8 \times 0.053} \approx 0.654$. The DermaMNIST bound is still far from asymptotic, confirming that weak features require many more experts.

2.2 Varying the Noise Rate η

Figure 2: SCX F1 vs. noise rate η for AIN MLIP (solid), CIFAR-10 (dashed), and DermaMNIST (dotted), all with $M = 8$ experts and adaptive threshold. Shaded regions: 95% confidence intervals. Theoretical predictions from Paper I Theorem 1 shown as thin lines.

Three regimes emerge:

- $\eta < 0.05$ (rare noise): Adaptive threshold yields 1.7–2.4 $\times$ improvement over fixed θ^* , confirming the prediction from Paper I (Theorem 4', Table 1, Case 3).
- $0.05 \leq \eta \leq 0.30$ (moderate noise): SCX achieves near-peak F1. The adaptive threshold still helps (1.1–1.7 $\times$) but the gain is smaller.
- $\eta > 0.30$ (high noise): F1 degrades for all methods. At $\eta > 0.50$, the problem becomes ill-posed: a majority of labels are wrong, and noise detection reduces to a variant of unsupervised learning.

2.3 Varying the Exploration Schedule $\eta(t)$

We compared four exploration schedules on AIN MLIP:

The single-pass approach (train M experts once, compute θ_{opt} , curate) achieves identical final performance to iterative schedules while requiring only one training pass. This suggests that for most applications, iterative refinement is unnecessary. Single-pass curation suffices because the multi-expert consistency signal stabilizes after a single exploration phase.

Table 5: Effect of $\eta(t)$ schedule on AIN MLIP final force RMSE (eV/Å) and number of curation iterations to convergence. All schedules start at $\eta_0 = 1.0$ and converge when $\eta(t) < 0.01$.

Schedule	Formula	Force RMSE	Iterations
Exponential (fast)	$\eta(t) = 0.5^{t/3}$	0.031 ± 0.003	12
Exponential (slow)	$\eta(t) = 0.5^{t/5}$	0.028 ± 0.002	18
Harmonic	$\eta(t) = 1/(1 + 0.2t)$	0.028 ± 0.002	28
Single-pass ($\eta = 1.0 \rightarrow \theta_{\text{opt}}$)	N/A	0.028 ± 0.002	1

3 Comparison with Alternative Curation Strategies

We benchmarked SCX against four alternative curation strategies across all four domains.

3.1 Alternative Methods

1. **Loss-based filtering** (baseline): Remove samples with the highest training loss. The number of samples removed is set to match the SCX removal rate.
2. **Confident learning** [2]: Estimate the joint distribution of noisy and true labels via counting on confident predictions. Remove samples in the off-diagonal of the estimated joint.
3. **Active label correction**: For each sample, query an oracle (simulated by using the true label) if the model’s prediction confidence is below a threshold. Remove samples confirmed as noise.
4. **Multi-annotator consensus**: Simulate three annotators with known error rates (matching the expert noise profile). Flag samples where annotators disagree as potentially noisy.

3.2 Results

Table 6: Comparison of curation strategies across domains. Metric: test performance after cleaning (force RMSE for AIN, accuracy for CIFAR-10 and MedMNIST, precision@0.1 for DrugBank). Best in **bold**.

Method	AIN RMSE	CIFAR-10 (20%)	DermaMNIST	DrugBank
No cleaning	0.045	78.3%	71.4%	0.58
Loss-based filtering	0.041	76.2%	71.8%	0.61
Confident learning [2]	0.038	79.1%	72.0%	0.63
Active label correction	0.033	81.5%	72.5%	0.67
Multi-annotator consensus	0.036	78.9%	71.6%	0.60
SCX (Tier 3, adaptive)	0.028	83.6%	72.8%	0.74

SCX outperforms all alternative strategies on every domain. The largest gains are on AIN MLIP (SCX: 0.028 vs. next best active correction: 0.033, a 15% relative improvement) and DrugBank (SCX: 0.74 vs. next best active correction: 0.67). On DermaMNIST, where Theorem 2 predicts weak features limit the achievable improvement, all methods show small gains, with SCX marginally ahead.

The advantage of SCX over active label correction is notable: active correction assumes an oracle that can provide clean labels for queried samples, which is unrealistic in many settings. SCX achieves better results *without* an oracle, using only the emergent consensus among diverse experts.

4 Self-Diagnosis: Bootstrap Stability Across Domains

The bootstrap stability diagnostic (Paper I, Proposition 6) was evaluated on all four domains to assess its predictive power.

Table 7: Bootstrap stability $S(\Phi, K)$, feature weakness ε_ϕ , and SCX improvement over loss baseline for all domains.

Domain	$S(\Phi, K)$	ε_ϕ	SCX improvement
AlN MLIP	0.83 ± 0.04	0.18 ± 0.03	+0.42 F1 (large)
CIFAR-10 (ResNet-18)	0.76 ± 0.05	0.15 ± 0.04	+0.17 F1 (moderate)
BloodMNIST (ResNet-18)	0.72 ± 0.06	0.22 ± 0.05	+0.02 accuracy (moderate)
DermaMNIST (SimpleCNN)	0.45 ± 0.08	0.52 ± 0.06	+0.06 ROC-AUC (small)
DrugBank (ECFP4)	0.68 ± 0.05	0.31 ± 0.04	+0.16 precision (moderate)

The threshold $S(\Phi, K) > 0.7$ correctly identifies AlN, CIFAR-10, and BloodMNIST as domains where SCX provides meaningful improvement. DermaMNIST ($S = 0.45$) falls well below the threshold, correctly predicting limited gains. DrugBank ($S = 0.68$) sits near the threshold, consistent with the moderate but real improvement (+0.16 precision). The correlation between $S(\Phi, K)$ and SCX improvement is $r = 0.89$ ($p < 0.05$), supporting the diagnostic’s validity.

5 Relationship to Paper I Theorems

This paper operationalizes the theoretical results of the companion paper (Paper I, [1]). The cross-references are summarized below:

Table 8: Mapping between Paper I theorems and Paper II sections.

Paper I Theorem	Paper II Usage
Theorem 1 (Noise Detection Guarantee)	Eq. (3), bound on achievable F1
Theorem 2 (Weak Feature Failure)	Eq. (4), boundary condition for tradeoff
Theorem 3 (The Honest Person Theorem)	Section 2.1, why premature curation fails
Theorem 4’ (Exact Constant Minimax Optimality)	Adaptive threshold θ_{opt} , Section 2.4
Theorem 5 (Cluster Consistency)	State discovery guarantee, Section 2.5
Proposition 6 (Bootstrap Stability)	Self-diagnosis, Section 2.5 Step 1

For full theorem statements, proofs, and numerical verification of the exact constants, see the companion paper [1].

References

- [1] SCX. A Fundamental Impossibility in Data Quality: Distinguishing Label Noise from Sample Difficulty is Provably Unsolvable Without Explicit Assumptions. arXiv preprint (2025).
- [2] Northcutt, C. G., Jiang, L. & Chuang, I. L. Confident learning: Estimating uncertainty in dataset labels. *Journal of Artificial Intelligence Research* **70**, 1373–1411 (2021).