

SCX OOD Detection Lower Bound

SCX

2026

Abstract

Proves a finite-expert lower bound for out-of-distribution detection via Spring gatekeeper scores. Minimum detectable OOD shift: $\delta_{\text{min}} = \sigma / \sqrt{M}$. False alarm and detection power both controlled.

0.1 OOD Detection Lower Bound

The CUSUM statistic above monitors temporal drift across a rollout. Spring also induces a one-sample OOD test: a sample is suspicious when its gatekeeper score is atypical relative to scores stored in the in-distribution memory. To avoid collision with the CUSUM statistic S_t , write $G_t(x)$ for the per-sample Spring gatekeeper score; this is the score denoted $S_t(x)$ in the Spring gating notes.

Let $\mathfrak{M}_t$ be the Spring memory bank of accepted in-distribution samples. For M expert modules, let $Z_{m,t}(x) \in [0, 1]$ be expert m 's normalized validity score for sample x at time t , and define the aggregate gatekeeper score

$$G_t(x) = \frac{1}{M} \sum_{m=1}^M Z_{m,t}(x). \quad (1)$$

The memory bank estimates the in-distribution score center $\mu_{\text{in},t} = \mathbb{E}[G_t(X) \mid X \sim P_{\text{in},t}]$ by $\hat{\mu}_{\text{in},t}$.

Definition 0.1 (Spring OOD Hypothesis Test). OOD detection is the binary hypothesis test

$$H_0 : X \sim P_{\text{in},t}, \quad H_1 : X \sim P_{\text{out},t}, \quad (2)$$

with OOD shift

$$\delta_{\text{OOD}} = |\mathbb{E}[G_t(X) \mid H_1] - \mu_{\text{in},t}|. \quad (3)$$

For a tolerance $\epsilon > 0$, the ideal calibrated Spring detector is

$$\varphi_{\epsilon,t}(x) = |G_t(x) - \mu_{\text{in},t}| > \epsilon, \quad (4)$$

where $\varphi_{\epsilon,t}(x) = 1$ means “declare OOD.” In deployment, $\mu_{\text{in},t}$ is replaced by $\hat{\mu}_{\text{in},t}$ from $\mathfrak{M}_t$; the calibration error is handled in Corollary ??.

Theorem 0.2 (SCX OOD Detection Lower Bound). *Condition on the Spring memory state $\mathfrak{M}_t$. Assume that $Z_{1,t}(X), \dots, Z_{M,t}(X)$ are independent and take values in $[0, 1]$. Under H_0 , suppose $\mathbb{E}[G_t(X) \mid H_0] = \mu_{\text{in},t}$; under H_1 , suppose $|\mathbb{E}[G_t(X) \mid H_1] - \mu_{\text{in},t}| = \delta_{\text{OOD}}$. Then, writing $p_i(\epsilon) = \mathbb{P}(\varphi_{\epsilon,t}(X) = 1 \mid H_i)$, the test in (??) satisfies*

$$p_0(\epsilon) \leq 2 \exp(-2M\epsilon^2), \quad (5)$$

$$p_1(\epsilon) \geq 1 - \exp(-2M(\delta_{\text{OOD}} - \epsilon)^2) \quad (6)$$

for every $0 < \epsilon < \delta_{\text{OOD}}$. In particular, if the OOD shift has margin $\delta_{\text{OOD}} \geq 2\epsilon$, then

$$\mathbb{P}(\text{detect OOD} \mid \delta_{\text{OOD}} \geq 2\epsilon) \geq 1 - \exp(-2M\epsilon^2). \quad (7)$$

Proof. The bounds follow directly from Hoeffding's inequality applied to the M independent gatekeeper scores. $\square$