

State-Conditioned Expertise: Data Quality Across Domains

From Uncertainty Principle to Global Infrastructure

SCX

July 2, 2026

Abstract

Data quality is widely recognized as the dominant factor in machine learning performance, yet the field has lacked a rigorous framework for distinguishing label errors from intrinsic sample difficulty. The SCX (State-Conditioned eXpertise) framework fills this gap through a layered architecture whose centerpiece—Theorem 3, the SCX Uncertainty Principle—proves that noise and difficulty are observationally indistinguishable without structural assumptions. This hard boundary cascades into a series of consequences that span from mathematical foundations to geopolitical infrastructure: (i) a noise detection guarantee with exponential convergence and exact constant minimax optimality (Theorems 1 and 4’); (ii) an optimal active learning sampling theorem with closed-form rate ρ^* derived from the Cercis Score (Theorem 5); (iii) an audit protocol adoption game exhibiting a non-proliferation Nash equilibrium with late-mover multi-armed bandit dynamics (Theorem 6); (iv) a cross-domain partition preservation bound via the Situs conditional mutual information (Theorem 7); (v) a two-layer storage paradox with a phase transition at critical audit maturity α^* ; (vi) the Audit Sword principle establishing universal verifiability of all SCX quality claims; (vii) an efficiency cascade yielding $8.4\times$ inference throughput and 30–40% training step reduction on clean data; (viii) an economic geography of data quality revealing Matthew-effect concentration of audit infrastructure and manufacturing revival through the non-remoteability of physical audit; and (ix) a vision of the human future in which SCX automates the clerical dimension of data work, restoring the scientist to science. We close with five open problems that define the next generation of SCX theory. The SCX framework is made available for academic research; the mathematical theorems are in the public domain.

1 Introduction

Data quality has emerged as the single most impactful factor in machine learning performance, consistently dominating architectural improvements across multiple benchmarks [?, ?]. In materials science, cleaning 14% of training frames improves force prediction error by 29–48%, a gain that no architecture change can replicate [?]. In computer vision, data cleaning yields 12–19 times the performance improvement of architecture modification [?]. These findings challenge the field’s traditional emphasis on model-centric innovation and demand a principled framework for understanding *when* and *how* data errors can be detected.

The fundamental challenge is that label noise and sample difficulty are mathematically indistinguishable from observational data alone. This is not a practical limitation; it is a theorem. In 1948, Claude Shannon proved that a noisy channel has a capacity—a hard limit on how much information it can carry. The SCX framework enters this lineage through a single result: **Theorem 3**—the SCX Uncertainty Principle—proves that noise and difficulty are observationally identical without structural assumptions. From this hard boundary, the rest of the framework unfolds as a cascade of conse-

quences: if you cannot perfectly separate noise from difficulty, every quality claim must be independently auditable (the Audit Sword); if auditability is structurally guaranteed, data quality infrastructure becomes a shared public good (the Telegraph Public Good); if depth was always a compensation mechanism for noise, clean data enables shallow architectures with massive efficiency gains (the Efficiency Cascade); and if audit accuracy is a function of accumulated history, the economic geography of data quality concentrates in manufacturing-dense economies (the Matthew Effect). Each step is a logical consequence of the previous one. Together they form a complete arc from a mathematical impossibility to a vision of global infrastructure.

Scope and narrative structure. This review synthesizes the SCX framework across six established domains and eight new theoretical and conceptual contributions that extend its reach beyond data cleaning into active learning, game theory, cross-domain transfer, storage economics, system efficiency, economic geography, and the human future of data work. The paper is organized as follows. Section 2 presents the mathematical foundations, centered on Theorem 3 as the flagship uncertainty principle and extending to Theorems 5–7.

Section 3 surveys the six-domain instantiation. Section 4 presents the Storage Paradox. Section 5 presents the Audit Sword and Telegraph Public Good. Section 6 presents the Efficiency Cascade. Section 7 presents the Economic Geography of SCX. Section 8 reflects on the Human Future. Section 9 outlines five open problems for next-generation SCX theory. Section 10 concludes.

Relationship to companion papers. This review synthesizes results from the SCX paper series: Paper I establishes the mathematical foundations (Theorems 1–4') [?]; Paper II develops the curation-exploration trade-off [?]; Paper III provides cross-domain experimental validation [?]; Paper IV demonstrates MLIP expert merging [?]; Theorems 5–7 appear as standalone contributions in the present framework; the Storage Paradox, Efficiency Cascade, Economic Geography, and Human Future analyses are drawn from the Yajie Protocol architecture document and companion essays.

A vision for global infrastructure. Data quality is not a competitive advantage—it is a shared scientific challenge. The SCX theory reveals universal mathematical laws governing the detectability of data errors, laws that operate identically whether the data comes from a laboratory in Shanghai, a hospital in Nairobi, or a factory in Stuttgart. These laws cannot be owned, embargoed, or weaponized; they can only be discovered, verified, and applied. We make the complete framework available for academic research and evaluation, and call upon the international scientific community to build upon this foundation together.

2 Mathematical Foundations

The SCX framework rests on a growing body of theorems that together form a complete theory of data quality assessment. We present the core results with an emphasis on the narrative arc: the Uncertainty Principle (Theorem 3) is the hard boundary from which everything else follows.

2.1 Theorem 3: The SCX Uncertainty Principle

The starting point—and the flagship result—of SCX theory is a negative theorem that establishes fundamental limits on what any data-cleaning algorithm can achieve. It occupies the same structural position in data quality that Shannon’s channel capacity theorem occupies in communication, that Gödel’s incompleteness theorems occupy in formal proof, and that Heisenberg’s uncertainty principle occupies in quantum measurement: it

identifies a hard boundary, and in doing so, prevents wasted effort on its wrong side.

Theorem 1 (SCX Uncertainty Principle; Theorem 3 of [?]). For any $K \geq 2$ classification problem, any $M \geq 1$ experts, and any finite state space $\mathcal{S}$, there exist two data-generating processes $\mathcal{P}_{\text{noise}}$ and $\mathcal{P}_{\text{hard}}$ such that:

- Under $\mathcal{P}_{\text{noise}}$, label errors are caused by label noise at rate η_{err} ;
- Under $\mathcal{P}_{\text{hard}}$, all observed labels equal the true labels, but some samples are intrinsically difficult (ambiguity rate η_{amb});
- Under the symmetry condition $\eta_{\text{err}} = \eta_{\text{amb}} = \eta$, the two processes produce *identical* observable joint distributions over $(x, y, \{f_m(x)\})$;
- Any algorithm has expected error $\geq \eta\rho/2$ on the ambiguous subset, where ρ is its proportion.

This theorem says something precise and uncomfortable: data cleaning without stated assumptions is scientifically meaningless. Every noise-detection algorithm is making claims it cannot justify unless it declares its structural assumptions—and the theorem tells you exactly what those assumptions must be. The six SCX assumptions (A1–A6) are not arbitrary technical conditions; they are the minimal sufficient set for identifiability. Remove any one, and the ambiguity returns for some data distribution.

Geometric interpretation. Theorem 3 admits a clean fiber-bundle structure. Let $\mathcal{P}$ be the space of data-generating processes and $\mathcal{Q}$ the space of observable joint distributions. The map $\pi : \mathcal{P} \rightarrow \mathcal{Q}$ that forgets the internal noise/difficulty structure is surjective but not injective: the fiber $\pi^{-1}(Q)$ over any observable distribution contains at least two distinct processes—the noise-world and the hardness-world. The fiber has the structure of a convex set parametrized by the noise rate, with extreme points corresponding to pure-noise and pure-hardness interpretations. This geometric picture makes precise what “observationally indistinguishable” means.

The cascade from Theorem 3. The Uncertainty Principle generates a chain of consequences that structures the entire SCX framework:

1. If noise and difficulty are indistinguishable without assumptions, then every data quality claim must declare its assumptions (A1–A6).
2. If assumptions are declared, then every quality claim must be **independently auditable**—otherwise the claim is unfalsifiable and Theorem 3’s boundary is merely shifted, not resolved. This is the **Audit Sword** (Section ??).

3. If auditability is structurally guaranteed, data quality infrastructure functions as a shared channel whose fragmentation harms everyone—the **Telegraph Public Good** (Section ??).
4. If data can be audited, depth is revealed as a compensation mechanism for noise, not an architectural necessity. Clean data enables shallow architectures with massive efficiency gains—the **Efficiency Cascade** (Section ??).
5. If audit accuracy accumulates with history, the economic geography of data quality concentrates in regions that combine manufacturing density with audit infrastructure—the **Matthew Effect** (Section ??).

2.2 The Detection Guarantee and Minimax Optimality

Under assumptions A1–A6, the SCX framework provides a noise detector with provable guarantees. Train M experts on disjoint data subsets, compute consensus scores $C(x) = \frac{1}{M} \sum_{m=1}^M \mathbf{1}\{\ell(f_m(x), y) > \tau\}$, and flag samples when $C(x)$ exceeds a state-dependent threshold.

Theorem 2 (Noise Detection Guarantee; Theorem 1 of [?]). Under A1–A6, the SCX noise detector achieves:

$$F1 \geq 1 - \frac{1}{\eta} \sum_{s \in S} \rho_s \cdot \exp(-2M\Delta_s^2), \quad (1)$$

where ρ_s is the state proportion and Δ_s is the separation gap.

The F1 score converges to 1 exponentially in the number of experts M . This exponential rate is not arbitrary. Using Chernoff-Stein and Bahadur-Rao large-deviation asymptotics, SCX achieves exact constant minimax optimality (Theorem 4’): $\lim_{M \rightarrow \infty} e^{M\kappa} \sqrt{2\pi M} \cdot (1 - F1_{\text{SCX}}) = C_{\min}/\eta$, and no algorithm can achieve a smaller constant. The $O(1/M)$ adaptive threshold shift $\theta^\dagger$ is essential—naïve threshold selection can be 1.7–2.4× suboptimal.

Theorem 2 quantifies the dependence on feature quality: when the feature mapping ϕ satisfies $I(\phi(X); S) \leq \delta$, the F1 improvement is bounded by $C_F \sqrt{\delta/2}$. When features are strong ($\delta/\log K < 0.2$), SCX substantially improves over baselines; when weak ($\delta/\log K > 0.5$), it degrades gracefully to baseline—and signals this degradation through bootstrap stability diagnostics (Proposition 6).

2.3 Theorem 5: Optimal Active Learning via the Cercis Score

The SCX framework extends beyond noise detection into optimal data acquisition. The Cercis Score $S(x) =$

$Q(x) + \eta N(x)$ decomposes sample value into a quality component $Q(x)$ (model uncertainty) and a novelty component $N(x)$ (information diversity). Under the assumption that expected F_1 gain is affine in the Cercis Score, the optimal sampling distribution takes Gibbs form $p^*(x) \propto \pi(x) \exp(\alpha S(x)/\lambda)$.

The key result is a **closed-form optimal sampling rate**:

Theorem 3 (Optimal Sampling Rate; Theorem 5). The sampling rate that maximizes risk-adjusted expected F_1 gain in active learning is

$$\rho^* = \frac{\eta}{1 + \eta} \frac{\sum_{k=1}^K w_k \mathbb{V}[\Gamma_k]}{\sum_{k=1}^K w_k \bar{\Gamma}_k}, \quad (2)$$

where $\bar{\Gamma}_k$ and $\mathbb{V}[\Gamma_k]$ are the per-class mean and variance of expected F_1 gain, and w_k are class weights.

The formula provides a principled, hyperparameter-free answer to the long-standing question of how many samples to label. The factor $\eta/(1 + \eta)$ acts as a saturation weight: $\eta \rightarrow 0$ (pure quality) yields $\rho^* \rightarrow 0$; $\eta \rightarrow \infty$ (pure novelty) yields $\rho^* \rightarrow \sum w_k \mathbb{V}[\Gamma_k] / \sum w_k \bar{\Gamma}_k$. The plug-in estimator is $\sqrt{n}$ -consistent, and experiments on CIFAR-10, AG News, and Cora confirm 15–40% annotation cost reduction relative to fixed-budget heuristics.

2.4 Theorem 6: Audit Protocol Adoption Game

When two mutually incompatible audit protocols compete for adoption, each adopter faces a strategic decision shaped by intrinsic protocol quality, network effects, and the information revealed by earlier adopters. We model this as a sequential-move game where early movers choose under prior uncertainty and late movers employ a UCB-based Multi-Armed Bandit (MAB).

Theorem 4 (Non-Proliferation Equilibrium; Theorem 6). The audit protocol adoption game admits a mixed-strategy Nash equilibrium. When the network-effect coefficient satisfies $\gamma < \gamma_{\text{crit}}$, where

$$\gamma_{\text{crit}} = \frac{1}{N_0 - 1} \left(\Delta\mu^0 + \sqrt{\frac{2 \log N_1}{N_0 + 1}} + \frac{\sigma_\xi}{\sqrt{N_0}} \right)^{-1}, \quad (3)$$

the equilibrium is unique. The expected adoption-share distortion relative to full information is bounded by $O(\sqrt{\log N_1/N_0})$.

The analysis reveals an **information-externality channel**: early adopters internalize that their choice generates public signals consumed by the late-mover MAB, distorting equilibrium adoption away from the full-information benchmark. The Cercis quality–novelty decomposition maps naturally onto the MAB’s

exploitation-exploration trade-off: Q corresponds to intrinsic protocol quality μ_P , and N corresponds to the UCB exploration bonus.

When $\gamma < \gamma_{\text{crit}}$, the informational motive dominates and a unique mixed equilibrium emerges. When $\gamma \geq \gamma_{\text{crit}}$, multiple equilibria appear—typically two pure-strategy cascades (all-adopt- A , all-adopt- B) corresponding to information cascades where the MAB never explores the alternative. The critical threshold thus defines the boundary between diversified and monopolized audit infrastructure.

2.5 Theorem 7: Cross-Domain Partition Preservation

A state partition optimized on a source domain will degrade when deployed on a target domain with shifted distribution. Theorem 7 provides a tight bound on this degradation.

Theorem 5 (Cross-Domain Partition Preservation; Theorem 7). Let π be a state partition with K cells. The cross-domain information loss satisfies

$$\Delta I(\pi) \leq \sum_{k=1}^K \left[L_k \cdot \text{diam}(\Omega_k) \cdot \text{Wass}(P_X^{(s)}, P_X^{(t)}) + \delta_k + \varepsilon_k \right] \quad (4)$$

where L_k is the Lipschitz constant of the cell-conditional density, $\text{diam}(\Omega_k)$ is cell diameter, Wass is the 1-Wasserstein distance between source and target marginals, δ_k is discretization error, and ε_k is irreducible concept-shift error.

The bound cleanly separates three sources of degradation: geometric (smoothness $\times$ granularity $\times$ domain gap), statistical (discretization loss from finite partition resolution), and irreducible (changes in $P(Y|X)$ across domains). It is asymptotically tight—a sequence of domain shifts can be constructed for which the ratio approaches 1. The optimal partition granularity K^* under domain shift satisfies $K^* \asymp (\text{Wass} \cdot \bar{L}/L_H)^{-d/(d+1)}$, formalizing the intuitive principle that larger domain gaps call for coarser partitions.

3 Domain Applications

The SCX framework has been instantiated across six domains. We summarize each domain’s problem formulation, SCX adaptation, validated results, and honest limitations. Full details appear in Papers I–IV [?, ?, ?, ?] and the engineering guide [?].

3.1 Life Sciences

Drug discovery. The SCX-Drug module implements a four-layer DTI prediction pipeline combining molecu-

lar featurization, knowledge base lookup (DrugBank), drug-likeness assessment, and deep learning prediction. Applied to 15K molecules against 5K human targets (75M drug-target pairs), it identified TOP-10 validated targets including KRAS G12D (score 97/100) and WRN helicase (96/100). The four-layer architecture embodies the curation-exploration tradeoff: Layer 1 (knowledge base, $\sim 5\%$ coverage) is highly curated; Layer 3 (deep ML) explores all remaining pairs; Layer 4 (docking) curates only top- N candidates.

Medical imaging. On DermaMNIST, SCX-Noise achieves ROC-AUC 0.72–0.82 for detecting 10–40% injected label noise vs. 0.61–0.65 for baselines. On BloodMNIST, SCX-Routing achieves 93.2% accuracy (+2.0% over best single expert). On PathMNIST (100K images), SCX-Compress achieves $< 2\%$ accuracy loss at 20–40% compression. The DermaMNIST result represents the weak-feature regime ($\varepsilon_\phi \approx 0.52$), where Theorem 2 correctly predicts degraded but non-zero improvement.

3.2 Advanced Manufacturing

Machine-learned interatomic potentials. The AlN ACE dataset is the flagship SCX validation. With $M = 12$ experts, theoretical $F1 \geq 0.87$ matches empirical $F1$ of 0.87. After cleaning (removing 14% of frames as noisy), force RMSE improves 29–48%. Random removal yields 2–5% improvement—the gap measures the curation-exploration tradeoff. Cross-validation on Si, Cu, and MgO confirms generalization ($F1$ 0.82–0.91). The separation gap Δ_s varies primarily with the distinctness of the noise-generating physical process.

Quality control and process optimization. SCX for manufacturing quality control and process optimization has been proposed with theoretical justification; systematic experimental validation is pending.

3.3 Embodied Intelligence

Robot learning. SCX applies to demonstration data where genuine suboptimal trajectories (valuable for learning from failure) must be distinguished from noise (sensor error, teleoperation glitches). The key insight from Theorem 3: premature filtering of “bad” trajectories is provably harmful. SCX requires exploration—training on all demonstrations—before curation. Systematic validation on Meta-World and RoboNet benchmarks is forthcoming.

Sim-to-real transfer. Multi-policy experts trained on different domain randomization distributions identify whether failure is a simulation artifact (experts disagree) or a genuine dynamic challenge (experts agree).

3.4 Artificial Intelligence

Computer vision. On CIFAR-10 with 10% injected noise, SCX achieves F1=0.62 vs. 0.45 for loss-based baseline. The theoretical Hoeffding bound ($F1 \geq 0.18$) is satisfied but conservative; the Chernoff bound ($\exp(-20\kappa) \approx 5 \times 10^{-5}$) is tight for $M \geq 20$. Premature cleaning removes 15% clean + 5% noisy; deferred cleaning removes 18% noisy + 2% clean, achieving 4–7% higher test accuracy.

LLM data curation. The SCX-LLM analysis reveals six structural insights: (i) Theorem 3 provides a rigorous proof of LLM hallucination inevitability through multi-seed independent decoding, with $\eta\rho/2$ as a provable lower bound that no model scaling can erase; (ii) BPE tokenization is revealed as frequency-driven SCX state crystallization in a degenerate regime; (iii) Sitrus vs. RoPE maps position encoding onto the SCX state concept; (iv) depth is a product of the noisy era—SCX explains when it can be replaced rather than why it works; (v) the LLM is architecturally a Yajie system missing its audit layer—softmax outputs are unaudited self-assessments, while Yajie provides state-conditioned reliability scores from independent multi-expert consistency; (vi) the Audit Sword guarantees that every SCX curation decision is traceable to an explicit state, judge set, and reliability estimate.

3.5 Autonomous Driving

SCX for autonomous driving perception leverages the natural multi-sensor setup: LiDAR, camera, and radar serve as complementary experts whose consensus identifies sensor-specific noise artifacts vs. genuine perception challenges. The provable guarantees have direct safety relevance: the F1 lower bound provides a mathematical floor on data quality that is certifiable, enabling verifiable bounds on undetected noise samples and regulatory documentation with explicit assumptions.

3.6 Natural Sciences

SCX treats DFT with different functionals (PBE, SCAN, HSE06) or different codes (VASP, Quantum ESPRESSO) as an expert ensemble, identifying when the DFT “oracle” is trustworthy vs. unreliable. In chemistry, SCX assesses reaction prediction data where successful reactions are overrepresented and failed reactions rarely published. In cryo-EM, different particle-picking algorithms form the expert ensemble. In climate science, CMIP6 multi-model ensembles provide natural experts for spatial and temporal reliability assessment.

3.7 Cross-Domain Patterns

When SCX works: (1) Strong domain features ($\varepsilon_\phi < 0.5$). (2) Moderate noise rates (5–30%). (3) Available oracle or consensus labels for calibration.

When SCX struggles: (1) Weak features cause graceful degradation to baseline, with self-diagnosis via bootstrap ARI. (2) Very sparse noise ($\eta < 1\%$) requires $M \sim 170$ for non-vacuous bounds. (3) Unknown K_S without domain expertise.

The Curation-Exploration Tradeoff. Across all six domains, premature data cleaning destroys the signal needed to distinguish noise from hardness. The tradeoff has three regimes governed by feature strength ε_ϕ : under-exploration (curation is random), optimal exploration (maximal improvement), and over-exploration (noise becomes imprinted). Strong features ($\varepsilon_\phi < 0.3$) admit exponential decay of the exploration rate; intermediate features require harmonic decay; weak features ($\varepsilon_\phi > 0.7$) indicate that features, not cleaning, need improvement.

4 The Storage Paradox

A persistent assumption in the data economy is that more data yields more value. SCX challenges this assumption through a two-layer paradox.

4.1 First Layer: Clean Data Beats More Data

When data quality is unverifiable, the rational strategy is to store more—more training data dilutes individual errors, more backups protect against corruption, more redundant sensors provide cross-validation. SCX inverts this logic.

Theorem 6 (Storage Paradox). There exists a signal-to-noise threshold $\rho^* \in (0, 1)$ such that for any dataset with $\rho(\mathcal{D}) < \rho^*$: $V(\mathcal{D}_{\text{clean}}) > V(\mathcal{D})$. The clean subset is more valuable than the full noisy dataset. As SCX audit precision improves (Theorem 1), noise identification converges to completeness, and the value gap $V(\mathcal{D}_{\text{clean}}) - V(\mathcal{D})$ increases.

Noise contributes two penalties: an irreducible Bayes error term proportional to $(1 - \rho)$, and capacity consumption that displaces signal representation. SCX multi-expert consensus identifies noise with exponentially decaying false-positive rate, enabling safe deletion of flagged records while retaining the provably clean core.

4.2 Second Layer: The Jevons Counter-Effect

This first layer—audit $\rightarrow$ denoising $\rightarrow$ storage decline—assumes exogenous data generation. But audit maturity also improves downstream AI models, which lowers the cost of data generation, which induces *more* data production. This is the data version of the Jevons paradox: efficiency improvements stimulate demand.

Define audit maturity $\alpha(t) = \theta(t)/\theta_{\max}$. The total storage evolves as:

$$\frac{dS}{dt} = \underbrace{\kappa'(\alpha)\dot{\alpha}(t) \int_0^t \mathcal{G}(\tau)d\tau}_{\text{historical compression } (<0)} + \underbrace{\kappa(\alpha(t))\mathcal{G}(t)}_{\text{new net inflow } (>0)} \quad (5)$$

where $\kappa(\alpha) = \rho + (1 - \rho)(1 - \alpha)^\eta$ is the post-audit retention ratio and $\mathcal{G}(t) = \mathcal{G}_0(1 + \beta\alpha(t)^\gamma)$ is the audit-stimulated generation rate.

Definition 1 (Phase Transition α^*). The critical audit maturity α^* satisfies $\Phi_{\text{quality}}(\alpha^*) = \Phi_{\text{quantity}}(\alpha^*)$, where $\Phi_{\text{quality}} = \partial(\text{data quality})/\partial\alpha$ is the storage compression force and $\Phi_{\text{quantity}} = \partial(\text{data quantity})/\partial\alpha$ is the storage expansion force.

Three dynamic stages:

1. **Pre-audit** ($\alpha \approx 0$): Storage grows linearly with generation. No compression is possible.
2. **Preemptive inflation** ($0 < \alpha < \alpha^*$): Audit stimulates data generation faster than it enables compression. A game-theoretic “front-running” equilibrium emerges: each data holder accelerates generation to preempt competitors, even though the data currently contains noise. Total storage *accelerates* upward.
3. **Denoising-dominant** ($\alpha > \alpha^*$): Compression force exceeds generation stimulus. Historical data is systematically compressed to golden standards. Generation returns to fundamentals. Storage converges to $\sum_i |\mathcal{D}_{\text{clean},i}^*|$.

After Spring convergence, the Golden Standard retention principle applies: when multi-expert consensus produces no new anomalies across consecutive audit cycles ($\Delta\mathcal{E}_t \approx \emptyset$), the audit has exhausted the dataset’s quality information. Only the golden standard—the verified clean subset with cryptographic hash and audit metadata—need be retained. It is mathematically equivalent to the original dataset for all downstream tasks while stripped of all identified noise.

The paradox’s policy implication is non-obvious: a mid-term storage surge is a *sign of audit success*, not failure—it reflects actors accumulating data reserves for the coming high-precision audit era. Regulators should distinguish phases: penalize storage in Phase 2 at the risk of suppressing audit adoption; incentivize golden-standard retention in Phase 3 as the metric of audit maturity.

5 The Audit Sword and the Telegraph Public Good

5.1 The Audit Sword

Theorem 3 establishes a hard boundary. The Audit Sword is the institutional consequence.

Principle (Audit Sword). Any organization deploying the SCX framework for data quality assessment renders its assessment results *unconditionally verifiable* by any independent third party. The verification procedure requires no cooperation, permission, or access privileges from the audited organization. The false-positive rate of verification is bounded by the same concentration inequality as Theorem 1, decaying exponentially in the number of audit modules. Military use is not exempt.

This principle is not a policy choice—it is a mathematical consequence of the SCX architecture. Because the protocol’s specification is published and its reference implementation is open-source, any third party with access to the dataset and audit modules can independently re-execute the consensus computation and compare outputs. The verifier need not trust the original auditor; she need only trust the mathematics.

The Audit Sword generates a **mutual deterrence structure**: every deployer is simultaneously a potential auditor and a potential auditee. The deterrence operates without monitoring, enforcement, or adjudication. The mathematics itself is the verifier, the inspector, and the adjudicator.

A commercial corollary—the “Bible-and-Pope” model—follows naturally. The Spring self-evolution engine is released as open-source software (the “Bible”): any organization can deploy it, audit its own data, and accumulate a private evidentiary corpus. The Yajie API provides cross-client calibration (the “Pope”): it maps local quality scores onto a globally comparable reference scale. Organizations that decline the API are not technically incapacitated, but become epistemically isolated—unable to compare their data quality to industry benchmarks. The open-source engine manufactures the demand for calibration; the calibration service resolves it.

5.2 The Telegraph Public Good

Shannon’s original context was the telegraph: a physical channel carrying electrical pulses, subject to thermal noise, connecting distant points. The telegraph was infrastructure—a public utility, maintained for universal access. Data quality assessment occupies the same structural position today.

Training data is the channel through which empirical knowledge flows into models. If that channel is corrupted—if its error rate is unknown, if its provenance is opaque, if its quality claims are unverifiable—every downstream model inherits the corruption. Data infrastructure, like telegraph infrastructure, is a public good: it benefits all users, its value increases with universal adoption, and its fragmentation (competing proprietary audit standards, incompatible quality metrics) imposes coordination costs on everyone.

The Audit Sword makes this infrastructure possible by making quality claims independently verifiable. Without the Audit Sword, data quality would remain an experience good—asserted by the producer, costly for the consumer to verify. With the Audit Sword, data quality becomes a search good—revealed through adoption decisions and independently verifiable audit outputs. The transformation is structural: the market for data quality shifts from reputation-based to proof-based.

The Non-Proliferation Equilibrium. A world with N competing audit protocols, each partially trusted and partially contested, is a world where audit results are themselves subject to meta-audit disputes. The development of a competing protocol by one jurisdiction would trigger reciprocal development by rivals, fragmenting the global audit infrastructure along geopolitical lines. In this environment, the rational choice for most actors is to accept the dominance of a single protocol—provided it is governed under adequate neutrality safeguards—rather than pursue audit autarky. The Non-Proliferation Equilibrium formalizes this logic: when the incumbent’s time-accumulated accuracy advantage exceeds the fragmentation cost, universal adoption is the unique Nash equilibrium.

6 The SCX Efficiency Cascade

The depth-as-noisy-era-product insight has direct architectural consequences. When Yajie audit reduces training data noise to $\eta < 1\%$, Theorem 1 guarantees that M shallow explicit experts voting in parallel can match the detection F1 of a deep monolithic model. This transforms the inference compute graph from a serial chain of L_{deep} layers to M parallel chains of L_{shallow} layers ($L_{\text{shallow}} \ll L_{\text{deep}}$), with a lightweight voting step costing $O(\log M)$.

Three mechanisms compose the **SCX Efficiency Cascade**:

1. **Parallel shallow inference.** For $M = 5$ experts at 18 layers replacing one 152-layer model, the theoretical throughput improvement approaches $152/18 \approx 8.4\times$. This expert parallelism is *embarassingly parallel*: each expert processes the identical input independently, requiring zero inter-expert com-

munication until final voting. The architecture is naturally suited to distributed multi-GPU deployment without high-bandwidth interconnects.

2. **Accelerated training convergence.** Clean data ($\eta < 1\%$) yields lower-variance gradient estimates, enabling larger stable learning rates. Models trained on Yajie-audited data require 30–40% fewer training steps to reach target validation performance. For LLM-scale training consuming millions of GPU-hours, a one-third reduction represents absolute energy savings comparable to the annual electricity consumption of a mid-sized datacenter.
3. **Cost-aware judge orchestration.** Tiered scheduling routes most samples through cheap proxy judges, reserving expensive full-model judges for boundary cases where proxy consensus is ambiguous.

A structural corollary: SCX decouples model quality from extreme single-accelerator capability. Eight 18-layer experts on eight mid-range GPUs can match the effective inference capacity of a single 152-layer model on a flagship accelerator, lowering the minimum viable compute threshold and broadening the set of environments in which high-quality AI can be trained and deployed.

The efficiency cascade represents a rare alignment between performance optimization and energy minimization in a field oriented toward ever-larger models and ever-more-specialized hardware. It makes a falsifiable prediction: on Yajie-audited clean data, the F1 gap between ResNet-18 and ResNet-152 should be ≤ 0.03 ; on original noisy data, the gap should be ≥ 0.08 .

7 SCX Economic Geography

SCX audit infrastructure does not distribute uniformly. Its spatial concentration is governed by three interacting mechanisms.

7.1 The Matthew Effect of Audit Accumulation

The SCX audit quality $\theta(t)$ is a monotonically increasing function of cumulative audit volume $|\mathcal{E}_t|$. Each audit cycle enriches the shared evidentiary corpus, making the next audit more accurate, which attracts more audits. The feedback loop is self-reinforcing: $\theta \uparrow \rightarrow Q \uparrow \rightarrow M \uparrow \rightarrow \text{users} \uparrow \rightarrow |\mathcal{E}| \uparrow \rightarrow \theta \uparrow$. This is cumulative causation in the sense of Myrdal and Arthur—a process in which initial advantages compound rather than equilibrate.

Theorem 3’s unidentifiability amplifies this dynamic. A late entrant cannot distinguish whether its inferior model performance is due to worse data, worse algorithms, or both—because the audit infrastructure

needed to diagnose data quality is precisely what the late entrant lacks. The incumbent accumulated its CEC in an environment without competitors; the entrant must compete while simultaneously building the tool needed to understand its competitive position.

In the SCX era, two economies possess all structural prerequisites for becoming audit infrastructure poles: (1) large, diverse data pools; (2) independent compute infrastructure; (3) sufficiently large internal markets for initial CEC accumulation; (4) legal frameworks for translating audit outputs into regulatory compliance. China and the United States satisfy all four conditions; other economies face structural gaps in at least two.

7.2 Manufacturing Revival Through Non-Remoteable Audit

A critical asymmetry distinguishes physically embedded data from purely digital data. The audit quality of data from physical production—sensor measurements, quality control readings, laboratory instrument outputs—depends on physical sensor configuration, on-site domain expert judgment, and hardware calibration that cannot be replicated remotely: $\partial\theta_{\text{audit}}/\partial(\text{distance}) < 0$.

This non-remoteability has a direct economic consequence: a manufacturing facility with deployed SCX audit infrastructure (sensors + measurement stations + domain experts + Spring instance) produces not only physical goods but an inseparable byproduct—mathematically verifiable quality certification. Relocating production to a lower-cost region without equivalent audit infrastructure means losing this certification and the quality premium it commands in global markets.

The traditional manufacturing competitiveness narrative locates advantage in labor cost differentials. SCX cuts this transmission mechanism. When quality is independently verifiable, the competitive parameter shifts from Δw (labor cost gap) to $\Delta\theta_{\text{audit}}$ (audit quality gap). A high-wage facility with complete SCX audit may out-compete a low-wage facility without audit, because the market pays more for verifiable quality than it saves in labor costs.

Manufacturing-Audit Complex (MAC). Physical production lines, audit infrastructure, and energy-compute supply chains form a symbiotic triad with complementarity lock-in: removing any component degrades audit quality discontinuously. The resulting Nash equilibrium is asymmetric: manufacturing-dense economies deploy audit; service-dense economies, whose data lacks independent physical ground truth, face a structural disadvantage in auditability itself. Physical data has an SI-traceable reference (sensor calibration chains); service data has socially constructed ground truth. The auditability gradient is ontological, not algorithmic.

7.3 Geopolitical Implications

SCX introduces three emerging dimensions of national capability whose accumulation is structurally slower than traditional dimensions. Military projection, reserve currency status, and ideological influence can be accelerated by political will and fiscal mobilization. CEC size, audit accuracy, and operational standard-setting authority cannot—they require physical sensor deployment and incompressible time accumulation.

The energy-compute-audit triangle sharpens this asymmetry. Audit infrastructure requires stable, high-density power for GPU clusters. Economies that already possess industrial-grade power infrastructure for manufacturing can supply audit compute at lower marginal cost, reinforcing the MAC concentration. Conversely, audit capability is exposed to energy supply disruption as a new channel of geopolitical vulnerability.

A critical caveat: all economic geography conclusions depend on the non-remoteability hypothesis. If quantum sensing or high-fidelity digital twin technology enables remote physical measurement at precision matching on-site sensors, the MAC logic weakens and service-dense economies gain an alternative pathway. The hypothesis is empirical, not axiomatic.

8 Human Future

Every technology that automates a cognitive function reorganizes the human activities that generated it. Writing externalized memory; arithmetic externalized calculation; SCX externalizes data quality assessment.

Before SCX, a researcher who wanted to know whether a dataset was reliable had to develop domain-specific heuristics, manually inspect samples, cross-reference against trusted sources, and ultimately rely on professional judgment accumulated through years of experience. SCX makes data quality assessment a measurement rather than an art. The researcher is freed to do what SCX cannot: decide what question to ask, interpret what the cleaned data implies, and exercise the judgment that determines whether a finding is important or merely significant.

The pattern is consistent across domains. The data scientist is freed from cleaning pipelines to focus on model architecture. The clinician is freed from questioning trial data reliability to focus on treatment appropriateness. The policy analyst is freed from debating statistical quality to focus on policy design. The journalist is freed from verifying whether a leaked dataset has been manipulated to focus on what the verified data reveals.

What remains for humans, in each case, is what humans have always done that machines cannot: creativity (seeing a problem no one has formulated), judgment

(weighing incommensurable values), and taste (recognizing quality where quality cannot be reduced to a metric). None of these can be specified as optimization problems. SCX does not encroach on them; it clears the ground beneath them.

The telegraph analogy. The telegraph did not eliminate the human element from communication—it eliminated the human element from the part of communication that did not require humanity. The messenger who spent days riding between cities was freed to become a telegraph operator, a wire-service journalist, or something else no one had yet imagined. SCX does for data work what the telegraph did for communication: it automates the clerical part. The data scientist who previously spent eighty percent of their time on data cleaning becomes a scientist again.

Trust without relationships. SCX changes what trust is about. When audit results are independently verifiable, trust is no longer a prerequisite for collaboration. Two organizations that have never worked together, that operate in different jurisdictions, with no overlapping professional networks, can nevertheless verify each other’s data quality claims through shared audit infrastructure whose conclusions are mathematically bounded. Trust that previously required years of relationship can be established through a single audit cycle. This is a structural change in the economics of collaboration: when trust is cheap, collaboration expands to include anyone who can submit their data to independent audit. The barrier to entry in quality-sensitive markets shifts from the cost of building a reputation to the cost of maintaining data quality sufficient to survive scrutiny. Diligence scales better than reputation.

Longevity and the happy problem. The application of SCX auditing to medicine carries a demographic consequence. Theorem 1’s exponential noise detection reduces false positives in drug screening, which reduces failed clinical trials, which lowers R&D costs per approved drug, which accelerates the pharmaceutical pipeline—producing, over decades, a measurable extension of the human lifespan. The same mechanism operates in clinical diagnosis: multi-expert consensus reduces misdiagnosis rates through the same mathematics that reduces label errors in training data.

The objection that an aging population is a crisis confuses two distinct phenomena. An aging population produced by people living longer in sickness is a burden. An aging population produced by people living longer in health—by the extension not merely of lifespan but of healthspan—is a gift. It is, in a phrase demography has never had occasion to use before, a happy problem. Spring’s golden-standard retention formalizes this: by retaining only data that survives the most stringent consensus threshold, Spring stores fewer patient records but each stored record carries a confidence bound that

makes it actionable. The clinical consequence is treatment decisions based on evidence whose reliability can be stated in quantitative terms. This flips the medical AI narrative from “more data” to “better data”—and better data saves lives.

9 Next-Generation Theorems: Five Open Problems

The SCX framework, as presented in this review, resolves the foundational question of detectability (Theorem 3), provides optimal detection (Theorems 1, 4’), quantifies degradation (Theorem 2), extends to active acquisition (Theorem 5), strategic adoption (Theorem 6), and cross-domain transfer (Theorem 7). Ten structural open problems—five foundational and five emerging—define the next generation of SCX theory.

Open Problem 1 (Temporal Drift — Sharp Critical Rate). When data distributions shift over time, the Spring self-evolution loop must track a moving target. The Lyapunov analysis establishes that geometric ergodicity holds when the distribution drift rate v satisfies $v < (1 - \rho)/\Delta t$, and divergence occurs above this threshold. The *sharp* critical drift rate—the exact boundary between convergence and divergence rather than a conservative sufficient condition—remains open. Resolving it would provide a design constraint for deployed SCX systems: the maximum sustainable rate of distribution change before audit feedback destabilizes.

Open Problem 2 (Adversarial Phase Transition — The S_{crit} Boundary). Theorem 3 establishes that noise and difficulty are indistinguishable. Adversarial perturbations occupy an ambiguous position: they are neither classical noise (they are structured, not random) nor genuine difficulty (they exploit model-specific vulnerabilities). The Cercis framework suggests that adversarial samples have a detectable geometric signal in ∇S space when model smoothness holds, and become a genuine third class when it fails. A critical Cercis threshold S_{crit} is conjectured to separate the reducible and irreducible adversarial regimes. Its analytic form—as a function of model architecture, Lipschitz constant, and data geometry—is the central open problem connecting SCX to adversarial robustness.

Open Problem 3 (Causal Confounding — Unconditional Weak Separability). Theorem 3’s unidentifiability extends to causal discovery: label noise and unobserved confounding produce observationally equivalent Cercis deviation patterns. Under the Situs causal-skeleton preservation hypothesis, a test statistic with non-trivial power to separate noise from confounding exists (conditionally rigorous). Whether any such statistic exists *without* this hypothesis—unconditional weak

separability—is open. Resolution would determine whether SCX-based causal discovery requires the Situs skeleton hypothesis as a necessary condition, or merely as a sufficient one.

Open Problem 4 (Federated Zero-Knowledge — Situs Homomorphic Composition). Federated auditing requires computing multi-expert consensus across jurisdictions without exposing raw data. The Situs homomorphic composition operator would, if it exists, guarantee ε -bounded information loss while preserving audit completeness. The operator’s existence is an open problem with a clear theoretical path. Its resolution would enable privacy-preserving cross-jurisdictional audit without raw data access—a capability essential for the Telegraph Public Good under data sovereignty constraints.

Open Problem 5 (Non-Uniform Noise — Beyond A4). All existing SCX guarantees assume uniform label noise (Assumption A4). In many real-world settings—genomic sequencing, financial transactions, sensor networks—noise is systematic, correlated, or adversarial. Extending Theorems 1–4’ to non-uniform noise classes, and characterizing the degradation in the detection exponent as a function of noise structure, would expand the framework’s applicability to domains where A4 is the binding constraint.

Open Problem 6 (Quantum Audit — Entanglement-Enhanced Detection). Whether quantum entanglement among expert measurements can surpass the classical Chernoff detection exponent of Theorem 1, enabling a quadratic speedup $F1 \geq 1 - Ce^{-M^2\Delta^2}$ in the convergence rate under favorable entanglement structure.

Open Problem 7 (Audit Entropy Second Law — Monotonicity of $\mathcal{H}(\mathcal{E}_t)$). Whether a suitably defined audit entropy $\mathcal{H}(\mathcal{E}_t)$ over the cumulative evidentiary corpus is monotonically non-decreasing under Spring self-evolution, establishing a thermodynamic arrow that bounds the reversibility of audit quality degradation.

Open Problem 8 (Audit Complexity Lower Bound — Minimal M for Target Error). Whether a fundamental lower bound $M_{\min}(\varepsilon, \eta, K)$ governs the minimum number of experts required to achieve target detection error ε , as a function of noise rate η and state-space cardinality K , paralleling Shannon’s channel capacity bound in communication.

Open Problem 9 (Recursive Audit Fixed Point — The Self-Audit Operator). Whether the self-audit operator $\mathcal{A} : \mathcal{D} \rightarrow \mathcal{D}$ mapping a dataset to its audited subset possesses a non-trivial fixed point $\mathcal{A}(\mathcal{D}^*) = \mathcal{D}^*$, characterizing the limit where iterative quality assessment yields no further information gain.

Open Problem 10 (Alignment Audit — Extending Theorem 3 to Normative Error). Whether multi-expert consensus with value-diverse judge sets extends Theorem 3’s exponential detection guarantee from factual labeling error to normative misalignment in LLM outputs, establishing a rigorous bridge between data quality and AI alignment.

These ten problems span two generations of SCX theory. The temporal drift threshold (Problem 1) governs the deployment horizon; the adversarial boundary (Problem 2) defines the robustness frontier; the causal separability question (Problem 3) determines whether SCX can contribute to scientific discovery beyond data cleaning; the federated composition operator (Problem 4) enables the institutional architecture of the Telegraph Public Good; and non-uniform noise (Problem 5) expands the domain coverage. Five emerging directions extend the framework into new terrain: quantum audit (Problem 6) asks whether entanglement beats the classical detection exponent; the audit entropy second law (Problem 7) seeks a thermodynamic arrow for audit quality; the audit complexity lower bound (Problem 8) parallels Shannon’s capacity limit for expert ensembles; the recursive audit fixed point (Problem 9) characterizes the limit of iterative self-assessment; and alignment audit (Problem 10) extends Theorem 3 from factual to normative error. Together they constitute a complete research program for the next generation of SCX theory.

10 Conclusion

This review has traced the SCX framework from its mathematical foundations—centered on the Uncertainty Principle (Theorem 3) as the hard boundary distinguishing noise from difficulty—through its domain applications across six fields, and outward to its broader consequences: the Storage Paradox revealing that audit maturity can reduce storage demand, the Audit Sword establishing universal verifiability, the Efficiency Cascade demonstrating 8.4× throughput and 30–40% training reduction on clean data, the Economic Geography showing Matthew-effect concentration and manufacturing revival through non-remoteable audit, and the Human Future in which SCX restores the scientist to science.

The arc is coherent. Shannon taught us that noise has structure. Theorem 3 taught us that the structure has a hard boundary. The Audit Sword taught us that the boundary requires verifiability. The Efficiency Cascade taught us that verifiability reveals depth as compensation, not architecture. The Telegraph Public Good teaches us that, once data quality is auditable, maintaining it as shared infrastructure is not an ethical preference—it is the equilibrium toward which rational

actors converge.

The SCX framework is released as an open academic contribution: all theorems are fully proven, all assumptions are explicitly stated, and all protocols are documented for independent reproduction. The mathematical results—being theorems rather than products—cannot be restricted by licensing or export control. They are available to any researcher, institution, or practitioner who wishes to apply them. The ten open problems of Section ?? define the next milestones; we hope the broader scientific community will join in their resolution.

References

- [1] Northcutt, C. G., Jiang, L. & Chuang, I. L. Confident learning: Estimating uncertainty in dataset labels. *Journal of Artificial Intelligence Research* **70**, 1373–1411 (2021).
- [2] Lipinski, C. A. *et al.* Experimental and computational approaches to estimate solubility and permeability. *Advanced Drug Delivery Reviews* **23**, 3–25 (1997).
- [3] Bickerton, G. R. *et al.* Quantifying the chemical beauty of drugs. *Nature Chemistry* **4**, 90–98 (2012).
- [4] Mendez, D. *et al.* ChEMBL: towards direct deposition of bioassay data. *Nucleic Acids Research* **47**, D930–D940 (2019).
- [5] Liu, T. *et al.* BindingDB: a web-accessible database of experimentally determined protein–ligand binding affinities. *Nucleic Acids Research* **35**, D198–D201 (2007).
- [6] Bahadur, R. R. & Rao, R. R. On deviations of the sample mean. *Annals of Mathematical Statistics* **31**, 1015–1027 (1960).
- [7] Chernoff, H. A measure of asymptotic efficiency for tests of a hypothesis based on the sum of observations. *Annals of Mathematical Statistics* **23**, 493–507 (1952).
- [8] Hoeffding, W. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association* **58**, 13–30 (1963).
- [9] Pinsker, M. S. *Information and Information Stability of Random Variables and Processes*. Holden-Day (1964).
- [10] Vershynin, R. *High-Dimensional Probability*. Cambridge University Press (2018).
- [11] Shamir, O. & Tishby, N. On the reliability of clustering stability in the large sample regime. *NIPS 2008* (2009).
- [12] Landis, J. R. & Koch, G. G. The measurement of observer agreement for categorical data. *Biometrics* **33**, 159–174 (1977).
- [13] Dawid, A. P. & Skene, A. M. Maximum likelihood estimation of observer error-rates using the EM algorithm. *Journal of the Royal Statistical Society Series C* **28**, 20–28 (1979).
- [14] Pollard, D. Strong consistency of k-means clustering. *Annals of Statistics* **9**, 135–140 (1981).
- [15] Le Cam, L. *Asymptotic Methods in Statistical Decision Theory*. Springer (1986).
- [16] Teicher, H. Identifiability of finite mixtures. *Annals of Mathematical Statistics* **34**, 1265–1269 (1963).
- [17] Hoeffding, W. A class of statistics with asymptotically normal distribution. *Annals of Mathematical Statistics* **19**, 293–325 (1948).
- [18] Myrdal, G. *Economic Theory and Under-Developed Regions*. Duckworth (1957).
- [19] Arthur, W. B. Positive feedbacks in the economy. *Scientific American* **262**(2), 92–99 (1990).
- [20] McFadden, D. Conditional logit analysis of qualitative choice behavior. In *Frontiers in Econometrics*, Academic Press (1974).
- [21] Auer, P., Cesa-Bianchi, N. & Fischer, P. Finite-time analysis of the multiarmed bandit problem. *Machine Learning* **47**(2), 235–256 (2002).
- [22] Bikhchandani, S., Hirshleifer, D. & Welch, I. A theory of fads, fashion, custom, and cultural change as informational cascades. *Journal of Political Economy* **100**(5), 992–1026 (1992).
- [23] Ben-David, S., Blitzer, J., Crammer, K. & Pereira, F. Analysis of representations for domain adaptation. *NeurIPS* (2007).
- [24] Shen, J., Qu, Y., Zhang, W. & Yu, Y. Wasserstein distance guided representation learning for domain adaptation. *AAAI* (2018).
- [25] Ganin, Y. & Lempitsky, V. Domain-adversarial training of neural networks. *Journal of Machine Learning Research* **17**(1), 2096–2030 (2016).
- [26] Li, L., Walsh, T. J. & Littman, M. L. Towards a unified theory of state abstraction for MDPs. *ISAIM* (2006).

- [27] Cover, T. M. & Thomas, J. A. *Elements of Information Theory*, 2nd ed. Wiley (2006).
- [28] Villani, C. *Optimal Transport: Old and New*. Springer (2008).
- [29] Farrell, H. & Newman, A. L. Weaponized interdependence. *International Security* **44**(1), 42–79 (2019).
- [30] Drezner, D. W. Targeted sanctions in a world of global finance. *International Interactions* **41**(4), 755–764 (2015).
- [31] Benveniste, A., Métivier, M. & Priouret, P. *Adaptive Algorithms and Stochastic Approximations*. Springer (1990).
- [32] Kairouz, P., McMahan, H. B. *et al.* Advances and open problems in federated learning. *Foundations and Trends in Machine Learning* **14**(1–2), 1–210 (2021).
- [12] SCX. SCX-Noise: Multi-expert label noise detection for medical imaging. *Technical Report* (2026).
- [13] SCX. SCX-Routing: State-conditioned expert routing. *Technical Report* (2026).
- [14] SCX. SCX-Compress: State-conditioned data compression. *Technical Report* (2026).
- [15] Yang, J., Shi, R. & Ni, B. MedMNIST: Classification and segmentation benchmark for medical images. *arXiv preprint arXiv:2110.14795* (2021).
- [16] SCX. The Yajie Protocol: Technology lock-in, audit sovereignty, and the non-proliferation logic of data quality assessment. *Technical Report* (2026).

References

- [1] SCX. State-conditioned expertise: Mathematical foundations for data quality assessment. *arXiv preprint* (2026).
 - [2] SCX. The curation-exploration tradeoff: A theoretical and empirical analysis. *arXiv preprint* (2026).
 - [3] SCX. Cross-domain experimental validation of the SCX framework. *arXiv preprint* (2026).
 - [4] SCX. SCX for machine-learned interatomic potentials: Expert merging and noise detection. *arXiv preprint* (2026).
 - [5] SCX. SCX for LLM data curation: State-conditioned multi-judge calibration. *In preparation* (2026).
 - [6] SCX. Unified theorem document: Theorems 1–7 of the SCX framework. *Technical Report* (2026).
 - [7] SCX. SCX engineering guide: Practical deployment across domains. *Technical Report* (2026).
 - [8] SCX. Deep mathematical structure of the SCX framework. *Companion analysis* (2026).
 - [9] SCX. SCX-Life: Domain-specific instantiation for life sciences. *Technical Report* (2026).
 - [10] SCX. TOP-10 drug target identification via the SCX-Drug pipeline. *Technical Report* (2026).
 - [11] SCX. TargetProfileState: State-conditioned drug-target modeling. *Technical Report* (2026).
- Intellectual Property Statement.** The SCX software framework, including all algorithms, implementations, and experiment scripts described in this paper, was independently developed by the author and is released for academic research and evaluation purposes. Commercial use, proprietary deployment, and integration into production systems are subject to separate licensing terms available from the author. The mathematical theorems and proofs presented herein are in the public domain and cannot be the subject of patent claims. The AIN interatomic potential function and associated DFT reference data used in the materials science case study were provided by the SCX and are used with permission; these data remain the intellectual property of the SCX.
- Data Availability.** The AIN v3 dataset was provided by the SCX and is available upon reasonable request and with permission of the data owner. CIFAR-10, DermaMNIST, and MedMNIST are publicly available benchmark datasets. DrugBank is available from go.drugbank.com.
- Code Availability.** The SCX core framework is available at [repository]. Experiment reproduction scripts are included in the repository.
- Competing Interests.** The author declares the following competing interests: the SCX software framework may be commercialized through a future entity. The author has no competing interests with respect to the AIN data, which belong to the SCX.