Detecting Label Noise by State-Conditioned eXpertise

SCX

June 30, 2026

Abstract

All machine learning depends on data quality, yet the field lacks a rigorous understanding of when data errors can be detected at all. Here we prove a fundamental impossibility: distinguishing mislabeled samples from intrinsically difficult-but-correct samples is mathematically impossible from observational data alone. We identify the minimal set of structural assumptions that break this impossibility—disjoint expert training, uniform noise, and state-homogeneous error rates—and prove that under these conditions, multi-expert consistency detects label noise with an exponential convergence rate that is minimax optimal, including the exact asymptotic constant. The theoretical bounds are empirically tight: on a materials science benchmark, the predicted detection limit matches the observed F1 score exactly. Our results establish that every data-cleaning method makes implicit assumptions that must be stated for its claims to be scientifically meaningful, and provide both the necessary conditions and the optimal algorithm for the problem.

1 Main Text

The quality of training data determines the quality of machine learning models [2]. When training data contains label errors—samples whose given labels differ from the true labels—model performance degrades predictably and substantially [1]. A large industry of data-cleaning methods has emerged to address this: loss-based filtering, confidence learning [1], multi-annotator consensus, and active label correction, among others. These methods share an implicit premise: that data errors can be distinguished from genuinely difficult samples by some observable signal.

We show that this premise is false in the absence of additional structural assumptions. The problem of distinguishing label noise from sample difficulty is fundamentally unidentifiable from observational data.

To see why, consider two worlds. In World A, a dataset contains mislabeled samples: 10% of the labels in certain regions of the input space have been randomly flipped, while expert models trained on clean data achieve 85% accuracy in those regions. In World B, the same dataset has perfectly correct labels, but 10% of the samples in those regions belong to a different ground-truth class than the remaining 90%, and expert models—trained on what they perceive as conflicting data—genuinely struggle, again achieving 85% accuracy on that subset.

We prove that the joint distribution of inputs, observed labels, and multi-expert predictions is mathematically identical in both worlds. No test, no algorithm, and no amount of data can distinguish them. This is an identifiability theorem in the sense of classical statistics [4]: the noise rate and the difficulty rate are not separately identifiable from the observable marginal distribution. We term this the **The Honest Person Theorem** (Theorem 3 in Supplementary Information).

This result has a constructive corollary: it tells us exactly what additional structure is required to break the ambiguity. We identify six assumptions—(1) experts trained on disjoint data subsets, (2) bounded correlation of expert errors given the state, (3) bounded loss, (4) uniform label noise independent of the input, (5) state-homogeneous expert error rates, and (6) balanced error distribution across classes—that together form a **minimal sufficient set** for identifiability. Each assumption corresponds to a specific path out of the unidentifiability trap. Removing any one of them restores the ambiguity for some data distribution.

Under these six assumptions, we prove that multi-expert consistency provides a noise detector with provable guarantees. Let M experts be trained on disjoint data subsets. For each sample, define the **consistency score** $C(x)$ as the fraction of experts that fail to predict the given label. The detection rule is simple: flag a sample as noisy if $C(x)$ exceeds a state-dependent threshold θ. We prove (**Theorem 1**, SI S1):

$$\mathrm{F1} \geq 1 - \frac{1}{\eta} \sum_{s} \rho_s \exp\bigl(-2M\Delta_s^2\bigr), \tag{1}$$

where η is the noise rate, ρ_s is the fraction of data in state s, and Δ_s is the separation gap between the expected consistency score on clean samples and the detection threshold. The F1 score converges to 1 exponentially in the number of experts M. The weakest assumption in our framework is A2' (bounded expert error correlation). When this correlation is zero, we recover the original independent-experts bound. When it is non-zero, the effective number of independent experts degrades from M to $M/(1 + (M-1)\bar{\rho})$. We provide a procedure for estimating $\bar{\rho}$ from data in SI S1.

The exponential rate $2M\Delta_s^2$ is not arbitrary. Using the Chernoff-Stein lemma and Bahadur-Rao exact large-deviation asymptotics, we prove a matching minimax lower bound (**Theorem 4**, SI S4): for *any* noise detection algorithm,

$$\liminf_{M\to\infty} e^{M\kappa}\sqrt{2\pi M} \cdot (1 - \mathrm{F1}) \geq \frac{C_{\min}}{\eta}, \tag{2}$$

where $\kappa = \mathrm{KL}(\theta^* \| p_0)$ is the Chernoff information between the clean and noisy expert-error distributions, and $C_{\min}$ is an explicit constant depending only on the noise rate and the geometry of the two distributions at the optimal decision threshold. Our adaptive-threshold SCX detector achieves this lower bound with equality, establishing that it is **exact constant minimax optimal**—no algorithm can achieve a smaller error constant, even asymptotically.

The method has a detectable failure mode. When the features used for state discovery carry insufficient information about the true data structure, performance degrades gracefully to that of a simple loss-threshold baseline. We quantify this through the mutual information $\delta = I(\phi(X); S)$ between features and true states, proving (**Theorem 2**, SI S2):

$$\mathrm{F1}_{\mathrm{SCX}} \leq \mathrm{F1}_{\mathrm{base}} + C_F\sqrt{\frac{\delta}{2}}. \tag{3}$$

This provides a practical diagnostic: estimate δ before running the full method. If $\delta/\log K < 0.2$, the features are sufficiently informative. If the ratio exceeds 0.5, the method will not outperform simpler approaches, and effort should be directed toward better feature engineering rather than hyperparameter tuning.

We validate the theory on the AlN machine-learned interatomic potential (MLIP) dataset (AlN v3, SCX; see Data Availability), a materials science benchmark where 53 of 534 frames carry label noise from thermal and molecular dynamics simulations. Training $M = 12$ expert models on disjoint subsets yields a theoretical F1 bound of 0.87 (range 0.77–0.86 depending on state proportions), consistent with the empirical F1 of 0.87. The bound is empirically tight: the theory neither overpromises nor underdelivers.

Cross-domain validation on CIFAR-10 image classification and DermaMNIST medical imaging confirms the predicted behavior. On CIFAR-10 with ResNet-18 features ($\delta/\log K \approx 0.15$), SCX achieves F1 $= 0.62$ versus a loss baseline of 0.45 at 10% noise. On DermaMNIST with weak SimpleCNN features ($\delta/\log K \approx 0.52$), SCX degrades to baseline performance (F1 $= 0.101$ vs. 0.105), exactly as Theorem 2 predicts.

Three implications follow. **First**, every data-cleaning method makes implicit assumptions. Our unidentifiability theorem proves that no method can claim to detect label errors without stating its structural assumptions. **Second**, the optimal detection strategy under our assumptions—multi-expert consistency with an adaptive threshold—is computationally practical: it requires training $M \geq 8$ models on data subsets, which is feasible on standard hardware for datasets up to millions of samples. **Third**, the feature-strength diagnostic provides a decision rule for when to apply the method, preventing wasted computation on datasets where it cannot help.

SCX Ecosystem. The theoretical framework presented here is operationalized through a family of co-designed components spanning the full data quality pipeline. **Yajie** (audit engine) implements the multi-expert consistency detection protocol with provable exponential convergence (Theorem 1) and exact constant minimax optimality (Theorem 4'), serving as the computational core of noise detection [12]. **Spring** (self-evolving gating) provides a monotonically growing memory bank M_t with guaranteed Robbins-Monro convergence, enabling the periodic resurrection of previously discarded samples as the gatekeeper's discrimination improves [13]. **Situs** (physical positional encoding) augments state representations with geometry-anchored coordinates when the information-theoretic condition $I(Y; P \mid S) > 0$ holds, distinguishing physical position as a causal variable from statistical token embeddings [14]. **Cercis Score** integrates these signals into a unified quality metric $S = Q + \eta N$, where Q is the base quality from multi-expert consensus and ηN is a time-decaying novelty bonus that ensures continued exploration at the noise-difficulty boundary [15]. Together, these components form a complete pipeline from state crystallization through multi-expert auditing to self-evolving quality assessment.

The theory points to several open questions. The exact constant optimality proof uses the Bahadur-Rao theorem for i.i.d. Bernoulli observations; extending it to non-identical expert error rates (heterogeneous expert quality) requires a generalized large-deviation principle. The six identifiability assumptions are sufficient but their individual necessity has been established only for the symmetric two-class case. And the feature-strength bound uses Pinsker's inequality, which is conservative; sharper information-theoretic bounds using strong data-processing inequalities may tighten the diagnostic.

Data quality is widely acknowledged as the dominant factor in machine learning performance [2]. Our results provide the first rigorous mathematical foundation for understanding when and how data errors can be detected. The unidentifiability theorem establishes a hard boundary; the detection guarantee and matching lower bound establish what is achievable within that boundary; the feature-strength diagnostic tells practitioners which side of the boundary they are on.

Intellectual Property Statement. The SCX software framework, including all algorithms, implementations, and experiment scripts described in this paper, was independently developed by the author and is released . Commercial use, proprietary deployment, and integration into production systems are subject to separate licensing terms available from the author. The mathematical theorems and proofs presented herein are in the public domain and cannot be the subject of patent claims. The AlN interatomic potential function and associated DFT reference data used in the materials science case study were provided by the SCX and are used with permission; these data remain the intellectual property of the SCX.

Data Availability. The AlN v3 dataset was provided by the SCX and is available upon reasonable request and with permission of the data owner. CIFAR-10, DermaMNIST, and MedMNIST are publicly available benchmark datasets. DrugBank is available from go.drugbank.com.

Code Availability. The SCX core framework is available at [repository] . Experiment reproduction scripts are included in the repository.

Competing Interests. The author declares the following competing interests: the SCX software framework may be commercialized through a future entity. The author has no competing interests with respect to the AlN data, which belong to the SCX.

References

[1] Northcutt, C. G., Jiang, L. & Chuang, I. L. Confident learning: Estimating uncertainty in dataset labels. *Journal of Artificial Intelligence Research* **70**, 1373–1411 (2021).

[2] Sambasivan, N., Kapania, S., Highfill, H., Akrong, D., Paritosh, P. & Aroyo, L. M. "Everyone wants to do the model work, not the data work": Data cascades in high-stakes AI. *Proc. CHI* (2021).

[3] Floridi, L., Cowls, J., Beltrametti, M., et al. AI4People—An ethical framework for a good AI society. *Minds and Machines* **28**(4), 689–707 (2018).

[4] Casella, G. & Berger, R. L. *Statistical Inference.* Duxbury, 2nd ed. (2002).

[5] Chernoff, H. A measure of asymptotic efficiency for tests of a hypothesis based on the sum of observations. *Annals of Mathematical Statistics* **23**(4), 493–507 (1952).

[6] Hoeffding, W. Probability inequalities for sums of bounded random variables. *Journal of the American Statistical Association* **58**(301), 13–30 (1963).

[7] Bahadur, R. R. & Rao, R. R. On deviations of the sample mean. *Annals of Mathematical Statistics* **31**(4), 1015–1027 (1960).

[8] Cover, T. M. & Thomas, J. A. *Elements of Information Theory.* Wiley, 2nd ed. (2006).

[9] Krizhevsky, A. & Hinton, G. Learning multiple layers of features from tiny images. Technical Report, University of Toronto (2009).

[10] Yang, J., Shi, R., Wei, D., Liu, Z., Zhao, L., Ke, B., Pfister, H. & Ni, B. MedMNIST v2—A large-scale lightweight benchmark for 2D and 3D biomedical image classification. *Scientific Data* **10**, 41 (2023).

[11] Wishart, D. S., Knox, C., Guo, A. C., et al. DrugBank: a knowledgebase for drugs, drug actions and drug targets. *Nucleic Acids Research* **36**(Database issue), D901–D906 (2008).

[12] SCX. A Fundamental Impossibility in Data Quality: Distinguishing Label Noise from Sample Difficulty is Provably Unsolvable Without Explicit Assumptions. arXiv preprint (2025).

[13] SCX. Spring: A Self-Evolving Gatekeeper with Provable Convergence. Working paper (2026).

[14] SCX. Situs: Physics-Anchored Positional Encoding for State-Conditioned Expertise. Working paper (2026).

[15] SCX. A Taxonomic Theory of Neural Networks: Derivation of Known Machine Learning Phenomena from the SCX Axiom System. Working paper (2026).

2 Supplementary Information

3 Theorem 1: Multi-Expert Consistency Noise Detection Guarantee (The Unlikely Lone Genius Theorem)

We establish a formal guarantee for the SCX noise detection procedure. When M experts, each trained on a disjoint data subset, exhibit high consensus on a sample's erroneousness, that sample can be identified as label noise with confidence that converges exponentially fast to 1 as M grows.

3.1 Notation and Setup

Let $(\mathcal{X}, \mathcal{Y})$ be the input–label space with $|\mathcal{Y}| = K_{\mathcal{Y}}$ (classification) or $\mathcal{Y} \subseteq \mathbb{R}^d$ (regression). Data follows $(X, Y) \sim \mathcal{D}$ with true labeling function $f^* : \mathcal{X} \to \mathcal{Y}$ (unknown oracle). We have M expert models $\{f_m\}_{m=1}^M$, each $f_m : \mathcal{X} \to \mathcal{Y}$.

The state space $\mathcal{S}$ indexes a measurable partition $\Pi = \{s_1, \ldots, s_{|\mathcal{S}|}\}$ of $\mathcal{X}$, where each state $s \in \mathcal{S}$ groups samples with similar characteristics.

Definition 3.1 (Label noise model)**.** For a sample (x, y^*) with true label $y^* = f^*(x)$, the observed label y is generated as

$$y = \begin{cases} y^* & \text{with probability } 1 - \eta, \\ \text{Uniform}(\mathcal{Y} \setminus \{y^*\}) & \text{with probability } \eta, \end{cases} \tag{4}$$

where $\eta \in (0, \frac{1}{2})$ is the global noise rate and the noise event is independent of x and of all training data.

Let $\ell : \mathcal{Y} \times \mathcal{Y} \to [0, B]$ be a bounded loss function with $B < \infty$, and $\tau > 0$ an error threshold. Define the expert error indicator and the consensus score as

$$e_m(x, y) = \mathbf{1}\{\ell(f_m(x), y) > \tau\}, \qquad C(x) = \frac{1}{M} \sum_{m=1}^{M} e_m(x, y). \tag{5}$$

Definition 3.2 (Detection rule). A sample (x, y) is flagged as noisy if and only if $C(x) > \theta$, where $\theta \in (0, 1)$ is a detection threshold.

3.2 Assumptions (A1)–(A6)

The following assumptions are used throughout Theorem 1 and its supporting lemmas.

(A1) **(Disjoint training sets).** The M experts are trained on M disjoint independent and identically distributed subsets:

$$D_m \sim \mathcal{D}^{n_m}, \quad D_m \cap D_{m'} = \varnothing, \quad D_m \perp D_{m'} \text{ for } m \neq m'. \tag{6}$$

(A2) **(Conditional independence on clean data).** For any clean sample (x, y) (i.e. $y = y^*$), the error indicators $\{e_m(x, y)\}_{m=1}^{M}$ are conditionally independent given x.

Remark 3.3. This follows from (A1): $e_m(x, y) = \mathbf{1}\{\ell(f_m(x), y) > \tau\}$ depends only on f_m, and f_m is a function of D_m; since the D_m are mutually independent, $\{f_m(x)\}_m$ and consequently $\{e_m(x, y)\}_m$ are conditionally independent given x.

(A3) **(Bounded loss).** $\ell(a, b) \in [0, B]$ for all $a, b \in \mathcal{Y}$ with $B < \infty$.

(A4) **(Uniform independent noise).** The label-flipping event is independent of x and of all D_m. The noise label is uniformly distributed over $\mathcal{Y} \setminus \{y^*\}$.

(A5) **(State homogeneity).** The state partition Π is such that within each state s, the clean-data error rate of the experts is approximately uniform. Formally, there exist state-level constants $\{\mu_s\}_{s \in \mathcal{S}}$ satisfying

$$\sup_{x \in s} \mathbb{E}[C(X) \mid \text{clean}, X = x] \leq \mu_s, \qquad \forall s \in \mathcal{S}. \tag{7}$$

(A6) **(Balanced error distribution).** The experts' error probabilities do not concentrate excessively on any single incorrect class. There exists a constant $C_{\text{bal}} \geq 1$ such that for any state s and any $x \in s$,

$$\max_{c \neq y^*} \mu_c(x) \;\leq\; C_{\text{bal}} \cdot \frac{\mu_s}{K_{\mathcal{Y}} - 1}, \qquad \mu_c(x) = \frac{1}{M} \sum_{m=1}^{M} \mathbb{P}(f_m(x) = c \mid x). \tag{8}$$

The case $C_{\text{bal}} = 1$ corresponds to perfectly uniform errors (every incorrect class equally likely). $C_{\text{bal}} = 2$ is a conservative default. This assumption is empirically testable.

3.3 Lemma 1: Mean Separation

Lemma 3.4 (Mean Separation). *Under assumptions (A1)–(A5), for any $x \in s$,*

$$\mathbb{E}[C(X) \mid \textit{clean}, X = x] \leq \mu_s, \tag{9}$$

$$\mathbb{E}[C(X) \mid \textit{noise}, X = x] = 1 - \frac{1}{K_{\mathcal{Y}} - 1}\,\mathbb{E}[C(X) \mid \textit{clean}, X = x] \;\geq\; 1 - \frac{\mu_s}{K_{\mathcal{Y}} - 1}. \tag{10}$$

Consequently, when $\mu_s < \dfrac{K_{\mathcal{Y}} - 1}{K_{\mathcal{Y}}}$ there exists θ such that

$$\mathbb{E}[C \mid \textit{clean}] < \theta < \mathbb{E}[C \mid \textit{noise}],$$

and the separation gap for state s is

$$\Delta_s(\theta) = \min\Big(\theta - \mu_s,\ 1 - \frac{\mu_s}{K_{\mathcal{Y}} - 1} - \theta\Big). \tag{11}$$

With the optimal threshold

$$\theta_s^* = \frac{1}{2}\left(1 + \mu_s \cdot \frac{K_{\mathcal{Y}} - 2}{K_{\mathcal{Y}} - 1}\right),$$

the maximised gap (in the ideal case $C_{bal} = 1$) is

$$\Delta_s^* = \frac{1}{2}\left(1 - \mu_s \cdot \frac{K_{\mathcal{Y}}}{K_{\mathcal{Y}} - 1}\right).$$

When $C_{bal} > 1$, the optimal threshold shifts towards μ_s to compensate for the non-uniform error distribution.

Proof of Lemma 1. The clean-sample bound follows directly from (A5). For a noise sample,

$$\begin{aligned}
\mathbb{E}[C \mid \text{noise}, X = x] &= \frac{1}{M}\sum_{m=1}^{M} \mathbb{P}\big(\ell(f_m(x), y) > \tau \mid \text{noise}, x\big) \\
&= \frac{1}{M}\sum_{m=1}^{M} \mathbb{P}\big(f_m(x) \neq y \mid \text{noise}, x\big) \qquad \text{(for 0-1 loss)} \\
&= 1 - \frac{1}{M}\sum_{m=1}^{M} \mathbb{P}\big(f_m(x) = y \mid \text{noise}, x\big).
\end{aligned}$$

By (A4), the noise label y is uniform over $\mathcal{Y} \setminus \{y^*\}$ and independent of all experts:

$$\begin{aligned}
\mathbb{P}(f_m(x) = y \mid \text{noise}, x) &= \sum_{c \neq y^*} \mathbb{P}(y = c \mid \text{noise})\,\mathbb{P}(f_m(x) = c \mid x) \\
&= \frac{1}{K_{\mathcal{Y}} - 1}\sum_{c \neq y^*} \mathbb{P}(f_m(x) = c \mid x) \\
&= \frac{1}{K_{\mathcal{Y}} - 1}\,\mathbb{P}(f_m(x) \neq y^* \mid x) = \frac{1}{K_{\mathcal{Y}} - 1}\,\mathbb{E}[e_m \mid \text{clean}, x].
\end{aligned}$$

Averaging over m gives

$$\begin{aligned}\mathbb{E}[C \mid \text{noise}, x] &= 1 - \frac{1}{K_{\mathcal{Y}} - 1}\,\frac{1}{M}\sum_{m=1}^{M}\mathbb{E}[e_m \mid \text{clean}, x] \\ &= 1 - \frac{1}{K_{\mathcal{Y}} - 1}\,\mathbb{E}[C \mid \text{clean}, x] \;\geq\; 1 - \frac{\mu_s}{K_{\mathcal{Y}} - 1},\end{aligned}$$

where the inequality uses (A5). The separation gap exists whenever $\mu_s < (K_{\mathcal{Y}} - 1)/K_{\mathcal{Y}}$, because then $1 - \mu_s/(K_{\mathcal{Y}} - 1) > \mu_s$. □

3.4 Lemma 2: False Positive Rate Bound

Lemma 3.5 (False Positive Rate Upper Bound)**.** *For any state $s \in \mathcal{S}$ satisfying $\mu_s < \theta$, the probability that a clean sample is incorrectly flagged as noise obeys*

$$\mathbb{P}\big(C > \theta \mid \textit{clean}, X \in s\big) \;\leq\; \exp\!\big(-2M(\theta - \mu_s)^2\big). \tag{12}$$

Proof. Fix $x \in s$. By (A5), $\mathbb{E}[C \mid \text{clean}, X = x] \leq \mu_s$. By (A1)–(A2), $\{e_m\}$ are conditionally independent given x, and by (A3) each $e_m \in [0, 1]$. Apply Hoeffding's inequality [1] to $C = \frac{1}{M}\sum_m e_m$:

$$\mathbb{P}\big(C - \mathbb{E}[C] > \theta - \mu_s \mid \text{clean}, x\big) \;\leq\; \exp\!\big(-2M(\theta - \mu_s)^2\big),$$

where $\theta - \mu_s > 0$ by hypothesis. The bound holds for every $x \in s$, so integrating over x yields

$$\begin{aligned}\mathbb{P}(C > \theta \mid \text{clean}, X \in s) &= \mathbb{E}_{X|s}\big[\mathbb{P}(C > \theta \mid \text{clean}, X)\big] \\ &\leq \sup_{x \in s}\mathbb{P}(C > \theta \mid \text{clean}, x) \;\leq\; \exp\!\big(-2M(\theta - \mu_s)^2\big).\end{aligned}$$ □

3.5 Lemma 3: True Positive Rate Bound

Lemma 3.6 (True Positive Rate Lower Bound)**.** *Under assumptions (A1)–(A6), for any state $s \in \mathcal{S}$ satisfying*

$$\theta < 1 - C_{\textit{bal}} \cdot \frac{\mu_s}{K_{\mathcal{Y}} - 1},$$

the probability that a noisy sample is correctly detected obeys

$$\mathbb{P}\big(C > \theta \mid \textit{noise}, X \in s\big) \;\geq\; 1 - \exp\!\Big(-2M\Big(1 - C_{\textit{bal}} \cdot \frac{\mu_s}{K_{\mathcal{Y}} - 1} - \theta\Big)^2\Big). \tag{13}$$

Proof. Fix $x \in s$ and condition on a specific noise label $c \neq y^*$. For 0-1 loss, expert m errs when $f_m(x) \neq c$, so $e_m = \mathbf{1}\{f_m(x) \neq c\}$. By (A1) the $\{f_m\}$ are independent functions of disjoint data; by (A2) the $\{e_m\}$ are conditionally independent given (x, c). The conditional expectation is

$$\mathbb{E}[e_m \mid x, c] = 1 - \mathbb{P}(f_m(x) = c \mid x) = 1 - \mu_{c,m}(x),$$

hence

$$\mathbb{E}[C \mid x, c] = \frac{1}{M}\sum_{m=1}^{M}\mathbb{E}[e_m \mid x, c] = 1 - \mu_c(x) \;\geq\; 1 - C_{\text{bal}} \cdot \frac{\mu_s}{K_{\mathcal{Y}} - 1},$$

where the inequality follows from (A6). Since $e_m \in [0,1]$ are conditionally independent, Hoeffding's inequality gives

$$\begin{aligned}
\mathbb{P}(C \le \theta \mid x, c) &= \mathbb{P}\left(\frac{1}{M}\sum_{m=1}^{M} e_m \le \theta \;\middle|\; x, c\right) \\
&\le \exp\!\Big(-2M\big(1-\mu_c(x)-\theta\big)^2\Big) \\
&\le \exp\!\Big(-2M\big(1-C_{\mathrm{bal}}\cdot\frac{\mu_s}{K_{\mathcal{Y}}-1}-\theta\big)^2\Big).
\end{aligned}$$

The Hoeffding gap condition $1-\mu_c(x) > \theta$ holds because

$$1-\mu_c(x) \ge 1 - C_{\mathrm{bal}}\cdot\frac{\mu_s}{K_{\mathcal{Y}}-1} > \theta$$

by the lemma's hypothesis and (A6). Averaging over the uniform noise label $c \in \mathcal{Y}\setminus\{y^*\}$,

$$\mathbb{P}(C \le \theta \mid \mathrm{noise}, x) = \frac{1}{K_{\mathcal{Y}}-1}\sum_{c\neq y^*}\mathbb{P}(C \le \theta \mid x, c) \;\le\; \exp\!\Big(-2M\big(1-C_{\mathrm{bal}}\cdot\frac{\mu_s}{K_{\mathcal{Y}}-1}-\theta\big)^2\Big).$$

Therefore $\mathbb{P}(C > \theta \mid \mathrm{noise}, x) \ge 1-\exp(-2M(1-C_{\mathrm{bal}}\cdot\mu_s/(K_{\mathcal{Y}}-1)-\theta)^2)$. Taking expectations over $X \in s$ completes the proof. □

Remark 3.7. When $C_{\mathrm{bal}} = 1$ (perfectly uniform error distribution), Lemma 3 recovers the optimal bound. The constant C_{bal} can be estimated from data and used to calibrate the detection threshold θ.

3.6 Proof of Theorem 1: F1 Lower Bound

Theorem 3.8 (SCX Noise Detection Guarantee)**.** *Let assumptions (A1)–(A6) hold. Let $\rho_s = \mathbb{P}(X \in s)$ be the state probability. For any threshold θ satisfying, for a given state s,*

$$\mu_s < \theta < 1 - C_{bal}\cdot\frac{\mu_s}{K_{\mathcal{Y}}-1},$$

define the state-level separation gap

$$\Delta_s \;=\; \min\!\Big(\theta-\mu_s,\; 1-C_{bal}\cdot\frac{\mu_s}{K_{\mathcal{Y}}-1}-\theta\Big) > 0. \tag{14}$$

When $C_{bal} = 1$, Δ_s reduces to the definition in Lemma 3.4; when $C_{bal} > 1$ the noise-side gap narrows because errors may not be perfectly uniform.

Then the SCX noise detector achieves the F1 lower bound

$$\boxed{\mathrm{F1} \;\ge\; 1-\frac{1}{\eta}\sum_{s\in\mathcal{S}}\rho_s\cdot\exp\!\big(-2M\Delta_s^2\big)} \tag{15}$$

or equivalently

$$\mathrm{F1} \;\ge\; 1-\sum_{s\in\mathcal{S}}\rho_s\Big[\exp\!\big(-2M\Delta_s^2\big)+\frac{1-\eta}{\eta}\exp\!\big(-2M\Delta_s^2\big)\Big]. \tag{16}$$

Remark 3.9 (Effective number of experts under correlation). The bound above uses the nominal number of experts M, which assumes perfect conditional independence (A2'). When experts share the same architecture or training paradigm, their errors may exhibit positive pairwise correlation $\bar{\rho}$. In this case the effective number of independent experts is reduced to the **The Expert Dilution Formula (Expert Dilution Formula)**: $M_{\text{eff}} = M/(1+(M-1)\bar{\rho})$ (see Liang & Zeger, 1986, for the generalized estimating equations variance-inflation factor). The bound then reads $\mathrm{F1} \geq 1 - \frac{1}{\eta}\sum_s \rho_s \exp(-2M_{\text{eff}}\Delta_s^2)$, which remains exponentially convergent in M but at a rate slowed by the factor $(1+(M-1)\bar{\rho})^{-1}$. Consistent estimation of $\bar{\rho}$ via jackknife or bootstrap is standard; see §**??** for details.

Proof. Let $R(x) = \mathbf{1}\{C(x) > \theta\}$ be the detection rule. Define

$$\begin{aligned}\mathrm{TPR}_s &= \mathbb{P}(R = 1 \mid \text{noise}, X \in s),\\ \mathrm{FPR}_s &= \mathbb{P}(R = 1 \mid \text{clean}, X \in s).\end{aligned}$$

By Lemma 3.6 and Lemma 3.5,

$$\begin{aligned}\mathrm{TPR}_s &\geq 1 - \exp\Big(-2M\big(1 - C_{\text{bal}} \cdot \frac{\mu_s}{K_{\mathcal{Y}} - 1} - \theta\big)^2\Big),\\ \mathrm{FPR}_s &\leq \exp\big(-2M(\theta - \mu_s)^2\big).\end{aligned}$$

From (A4), the noise event is independent of X, so $\mathbb{P}(\text{noise} \mid X \in s) = \eta$ for all s. Aggregating across states,

$$\mathrm{TPR} = \sum_s \rho_s \,\mathrm{TPR}_s, \qquad \mathrm{FPR} = \sum_s \rho_s \,\mathrm{FPR}_s\,.$$

The F1 score is

$$\mathrm{F1} = \frac{2\eta\,\mathrm{TPR}}{2\eta\,\mathrm{TPR} + (1-\eta)\,\mathrm{FPR} + \eta(1-\mathrm{TPR})} = \frac{2\eta\,\mathrm{TPR}}{\eta(1+\mathrm{TPR}) + (1-\eta)\,\mathrm{FPR}}.$$

Set

$$\begin{aligned}\delta_1 &= \sum_s \rho_s \exp\Big(-2M\big(1 - C_{\text{bal}} \cdot \frac{\mu_s}{K_{\mathcal{Y}} - 1} - \theta\big)^2\Big),\\ \delta_2 &= \sum_s \rho_s \exp\big(-2M(\theta - \mu_s)^2\big).\end{aligned}$$

Then $\mathrm{TPR} \geq 1 - \delta_1$ and $\mathrm{FPR} \leq \delta_2$. Substituting,

$$\begin{aligned}\mathrm{F1} &\geq \frac{2\eta(1-\delta_1)}{\eta(2-\delta_1) + (1-\eta)\delta_2}\\ &= 1 - \frac{\eta\delta_1 + (1-\eta)\delta_2}{\eta(2-\delta_1) + (1-\eta)\delta_2}\\ &\geq 1 - \frac{\eta\delta_1 + (1-\eta)\delta_2}{\eta} \qquad \text{(since denominator } \geq \eta)\\ &= 1 - \delta_1 - \frac{1-\eta}{\eta}\,\delta_2.\end{aligned}$$

Now observe that Δ_s as defined in (14) satisfies

$$\exp\big(-2M(\theta-\mu_s)^2\big) \leq \exp\big(-2M\Delta_s^2\big), \quad \exp\Big(-2M\big(1-C_{\text{bal}}\cdot\frac{\mu_s}{K_{\mathcal{Y}} - 1}-\theta\big)^2\Big) \leq \exp\big(-2M\Delta_s^2\big).$$

Therefore

$$
\begin{aligned}
\mathrm{F1} &\geq 1-\sum_s \rho_s\Big[\exp\big(-2M\Delta_s^2\big)+\frac{1-\eta}{\eta}\exp\big(-2M\Delta_s^2\big)\Big] \\
&= 1-\frac{1}{\eta}\sum_s \rho_s \exp\big(-2M\Delta_s^2\big). \qquad \square
\end{aligned}
$$

Corollary 3.10 (Asymptotic optimality)**.** *As $M \to \infty$, for all states satisfying $\mu_s < (K_{\mathcal{Y}}-1)/K_{\mathcal{Y}}$,*

$$
\mathrm{F1} = 1-\mathcal{O}_P\left(\frac{1}{\eta}\,e^{-2M\Delta_{\min}^2}\right), \qquad \Delta_{\min}=\min_{s\in\mathcal{S}}\Delta_s.
$$

Example 3.11 (Numerical illustration). Take $K_{\mathcal{Y}}=10$ classes, $\eta=0.1$ noise rate, $M=20$ experts, and per-state clean error bound $\mu_s=0.2$. With $C_{\mathrm{bal}}=1$, the optimal separation gap is

$$
\Delta_s^* = \frac{1}{2}\left(1-0.2\cdot\frac{10}{9}\right)\approx 0.389,
$$

and the F1 lower bound from Theorem 3.8 is

$$
\mathrm{F1} \;\geq\; 1-10\cdot\exp\big(-2\cdot 20\cdot 0.389^2\big) = 1-10\cdot e^{-6.05} > 0.976.
$$

In practice, if the clean expert error rate is higher (e.g. $\mu_s \approx 0.45$ as in CIFAR-10 with lightly trained experts), the gap shrinks to $\Delta_s \approx 0.25$, giving the weaker bound $\mathrm{F1} \geq 0.18$, which remains valid as a guarantee but is not tight. This sensitivity to μ_s underscores the importance of estimating μ_s accurately from validation data.

3.7 Chernoff Tightening

Lemmas 3.5 and 3.6 can be sharpened using the Chernoff bound instead of Hoeffding. For clean data, let $\{e_m\}$ be independent Bernoulli variables with mean $\mu=\mathbb{E}[C \mid \mathrm{clean}, x] \leq \mu_s$. For $\theta > \mu$,

$$
\begin{aligned}
\mathbb{P}(C>\theta \mid \mathrm{clean}, x) &= \mathbb{P}\left(\frac{1}{M}\sum_m e_m > \theta\right) \\
&\leq \inf_{\lambda>0}\, e^{-\lambda M\theta}\,\mathbb{E}\big[e^{\lambda\sum_m e_m}\big] \\
&= \inf_{\lambda>0}\, e^{-\lambda M\theta}\prod_{m=1}^{M}(1-\mu_m+\mu_m e^{\lambda}) \\
&\leq \exp\big(-M\cdot \mathrm{KL}(\theta \,\|\, \mu_s)\big),
\end{aligned}
$$

where $\mathrm{KL}(p\,\|\,q) = p\log\frac{p}{q}+(1-p)\log\frac{1-p}{1-q}$.

For the noise lower tail, conditioning on label c and applying the same bound yields

$$
\mathbb{P}(C \leq \theta \mid \mathrm{noise}, x, c) \;\leq\; \exp\left(-M\cdot\mathrm{KL}\left(\theta \,\middle\|\, 1-C_{\mathrm{bal}}\cdot\frac{\mu_s}{K_{\mathcal{Y}}-1}\right)\right).
$$

Remark 3.12 (KL direction). The standard Chernoff bound for the lower tail is $\mathbb{P}(\bar{X} \leq \theta) \leq \exp(-n\cdot\mathrm{KL}(\theta\|p))$ when $\theta < p$. Here the true mean is $p = 1-C_{\mathrm{bal}}\cdot\mu_s/(K_{\mathcal{Y}}-1)$ and the threshold θ satisfies $\theta < p$ by hypothesis, so the correct KL argument is θ (first) and p (second).

Combining the two tails gives the Chernoff-form F1 bound

$$\mathrm{F1} \;\geq\; 1 - \frac{1}{\eta}\sum_{s\in\mathcal{S}} \rho_s \Big[\exp\!\big(-M\cdot \mathrm{KL}(\theta \,\|\, \mu_s)\big) \\ + \frac{1-\eta}{\eta}\cdot \exp\!\Big(-M\cdot \mathrm{KL}\Big(\theta \,\Big\|\, 1 - C_{\text{bal}}\cdot \frac{\mu_s}{K_{\mathcal{Y}}-1}\Big)\Big)\Big]. \quad (17)$$

The Chernoff bound is tighter than the Hoeffding bound by a factor that grows with the separation gap. Specifically,

$$\exp(-2M\Delta^2) \geq \exp(-M\cdot \mathrm{KL}(\mu+\Delta \,\|\, \mu)),$$

with equality only as $\Delta \to 0$. In the practical regime $\Delta \approx 0.1$–0.4, the Chernoff bound is 2–5 times tighter.

3.8 Corollaries

Corollary 1: Symmetric Experts

If all experts share the same clean-data error rate ε (symmetric experts), then $\mu_s = \varepsilon$ for all states,

$$\Delta = \frac{1}{2}\left(1 - \varepsilon \cdot \frac{K_{\mathcal{Y}}}{K_{\mathcal{Y}}-1}\right),$$

and the F1 bound simplifies to

$$\mathrm{F1} \;\geq\; 1 - \frac{1}{\eta}\exp\left(-\frac{M}{2}\left(1 - \frac{\varepsilon K_{\mathcal{Y}}}{K_{\mathcal{Y}}-1}\right)^2\right). \quad (18)$$

For binary classification ($K_{\mathcal{Y}} = 2$) this further reduces to

$$\mathrm{F1} \;\geq\; 1 - \frac{1}{\eta}\exp\!\big(-2M(\tfrac{1}{2}-\varepsilon)^2\big).$$

Corollary 2: Optimal Threshold Selection

Given noise rate η, number of classes $K_{\mathcal{Y}}$, number of experts M, and maximum clean error rate $\mu_{\max} = \max_s \mu_s$, the optimal threshold is

$$\theta^* = \arg\max_{\theta}\min_{s} \Delta_s(\theta) = \frac{1}{2}\left(1 + \mu_{\max}\cdot \frac{K_{\mathcal{Y}}-2}{K_{\mathcal{Y}}-1}\right). \quad (19)$$

The resulting minimum separation gap is

$$\Delta^*_{\min} = \frac{1}{2}\left(1 - \mu_{\max}\cdot \frac{K_{\mathcal{Y}}}{K_{\mathcal{Y}}-1}\right),$$

and the minimum number of experts required to achieve $\mathrm{F1} \geq 1 - \varepsilon_0$ is

$$M \;\geq\; \frac{1}{2\Delta^{*2}_{\min}}\log\left(\frac{1}{\eta\varepsilon_0}\right). \quad (20)$$

When controlling error across multiple states simultaneously, replace the numerator 1 in the logarithm with $|\mathcal{S}|$.

Corollary 3: Uniform Detectability

If there exists $\delta > 0$ such that $\mu_s \le (K_{\mathcal{Y}}-1)/K_{\mathcal{Y}}-\delta$ for all s, then with the fixed threshold $\theta = \frac{1}{2}$,

$$\Delta_s \ge \frac{\delta}{2}, \qquad \mathrm{F1} \;\ge\; 1 - \frac{1}{\eta}\exp\!\left(-\frac{M\delta^2}{2}\right).$$

Thus whenever every state's clean expert error rate is uniformly bounded below $(K_{\mathcal{Y}} - 1)/K_{\mathcal{Y}}$, a fixed threshold yields exponential F1 guarantees.

Corollary 4: Finite-Sample Correction

In practice μ_s must be estimated from finite validation data. Let n_s be the number of clean validation samples in state s, and let $\hat{\mu}_s = \frac{1}{M}\sum_m \hat{\varepsilon}_m(s)$ be the empirical estimate, where $\hat{\varepsilon}_m(s)$ is expert m's empirical error rate on state s. By Hoeffding and the union bound, with probability at least $1 - \delta_0$,

$$|\hat{\mu}_s - \mu_s| \;\le\; B\sqrt{\frac{\log(2M/\delta_0)}{2n_s}}, \qquad \forall m, s.$$

Define the conservative upper bound $\tilde{\mu}_s = \hat{\mu}_s + B\sqrt{\log(2M/\delta_0)/(2n_s)}$. Then Theorem 3.8 holds with probability $\ge 1 - \delta_0$ when $\tilde{\mu}_s$ replaces μ_s in Δ_s.

3.9 Connection to Dawid–Skene

Dawid–Skene model

Dawid and Skene [2] estimate true labels and annotator reliability from multiple independent annotations via EM. Each annotator m is assigned a global confusion matrix

$$\pi_m(c' \mid c) = \mathbb{P}(f_m(x) = c' \mid y^* = c), \qquad \forall x \in \mathcal{X}.$$

The final label estimate is a weighted vote $\hat{y}_{\mathrm{DS}} = \arg\max_c \sum_m w_m \mathbf{1}\{f_m(x) = c\}$, where weights w_m are derived from the diagonal entries of π_m.

SCX generalization

SCX replaces the global confusion matrix with a state-conditioned variant:

$$\pi_m(c' \mid c, s) = \mathbb{P}(f_m(x) = c' \mid y^* = c,\ x \in s).$$

SCX weights therefore depend on the state:

$$w_m(x) = \frac{\exp(-\alpha \hat{R}_m(s(x)))}{\sum_{m'} \exp(-\alpha \hat{R}_{m'}(s(x)))}, \qquad \hat{R}_m(s) = \sum_c \pi_m(c \mid c, s).$$

Comparison

The SCX detection rule in Theorem 1 is built on the *consensus* concept – noisy samples exhibit high error across *all* experts – which is qualitatively different from Dawid–Skene's global reliability estimation. Key differences are summarised in Table 1.

When the state partition degenerates to the trivial partition $\mathcal{S} = \{\mathcal{X}\}$, SCX recovers Dawid–Skene as a special case: weights become global constants and the consensus score becomes the global average error rate.

Table 1: Comparison of Dawid–Skene and SCX (Theorem 1) for noise detection.

Dimension	Dawid–Skene	SCX (Theorem 1)
Reliability	Global constant π_m	State-conditioned $\pi_m(s)$
Noise detection	Posterior agreement	Multi-expert error consensus
Annotations needed	Multiple (≥ 3) annotators	Multiple expert models
Class imbalance robustness	Weak	Strong
Separates noise from hardness	No	Yes (consensus score C)
Theoretical guarantee	Asymptotic consistency	Exponential convergence

References

[1] W. Hoeffding, "Probability inequalities for sums of bounded random variables," *J. Amer. Statist. Assoc.*, vol. 58, no. 301, pp. 13–30, 1963.

[2] A. P. Dawid and A. M. Skene, "Maximum likelihood estimation of observer error-rates using the EM algorithm," *J. R. Statist. Soc. Ser. C*, vol. 28, no. 1, pp. 20–28, 1979.

4 Theorem 2: Weak Feature Failure Boundary (The Impossible Demand Theorem)

When the feature representation $\phi(x)$ contains insufficient information about the true state S, consistency-based noise detection methods (including SCX) cannot outperform the loss-based baseline. Theorem 2 quantifies this boundary from an information-theoretic perspective.

4.1 Information-Theoretic Setup

Data generation. Let $X \in \mathcal{X} \subseteq \mathbb{R}^d$ be the input, $Y \in \mathcal{Y}$ the label, and $Z \in \{0,1\}$ the noise indicator ($Z = 1$ for noisy samples). The true (unobserved) state is $S = s(X) \in \mathcal{S}$ with $|\mathcal{S}| = K_{\mathcal{S}}$. The observed feature representation is $\phi(X) \in \Phi \subseteq \mathbb{R}^{d_\phi}$.

SCX assumes the following generative structure:

$$Z \;\to\; S \;\to\; (X,Y) \;\to\; \phi(X), \qquad Z \;\to\; S \;\to\; X \;\to\; \{f_m(X)\}.$$

This yields the Markov chain

$$Z \to S \to X \to \phi(X), \tag{21}$$

and by the data processing inequality,

$$I(\phi(X); Z) \leq I(X; Z) \leq I(S; Z) \leq H(Z) \leq \log 2.$$

Weak feature definition.

Definition 4.1 (δ-weak feature)**.** A feature map $\phi : \mathcal{X} \to \Phi$ is called δ-weak with respect to the true state map $s : \mathcal{X} \to \mathcal{S}$ if

$$I(\phi(X); s(X)) \leq \delta, \tag{22}$$

where $I(\cdot;\cdot)$ is the Shannon mutual information, measured in nats.

The normalised weakness is

$$\varepsilon_\phi = \frac{\delta}{\log K_{\mathcal{S}}} \in [0,1], \tag{23}$$

since $I(\phi; S) \le H(S) \le \log K_{\mathcal{S}}$. $\varepsilon_\phi = 1$ corresponds to features with no state information (worst case), $\varepsilon_\phi = 0$ to full state information.

An equivalent definition uses the total variation distance. For two states s_1, s_2, let $\mathrm{TV}_{s_1,s_2} = \mathrm{TV}(P_{\phi|S=s_1}, P_{\phi|S=s_2})$. Then the feature weakness is

$$\Delta_\phi = \max_{s_1 \ne s_2} \mathrm{TV}(P_{\phi|S=s_1}, P_{\phi|S=s_2}), \tag{24}$$

and by Pinsker's inequality,

$$\Delta_\phi \le \sqrt{\frac{\delta}{2}}. \tag{25}$$

When Δ_ϕ is small, points from different true states are indistinguishable in ϕ-space.

Loss baseline detector. The loss-based detector uses no state information:

$$\hat{z}_{\text{loss}}(x) = \mathbf{1}\{\max_m \ell(f_m(x), y) > t\},$$

where $\{f_m\}_{m=1}^M$ are M expert models. Denote its optimal performance by $\mathrm{AUC}_{\text{base}}$, $\mathrm{PRAUC}_{\text{base}}$, and $\mathrm{F1}_{\text{base}}$.

For a completely random detector, the optimal performance is

$$\mathrm{AUC}_{\text{rand}} = 0.5, \qquad \mathrm{PRAUC}_{\text{rand}} = \eta, \qquad \mathrm{F1}_{\text{rand}} = \frac{2\eta}{1+\eta}, \tag{26}$$

where $\eta = \mathbb{P}(Z = 1)$ is the marginal noise rate.

4.2 Lemma 1: State Estimation Error (Fano)

Lemma 4.2 (State estimation error lower bound – Fano's inequality). *Let $\hat{S}$ be any estimator of the true state S based on $\phi(X)$ (i.e., the output of the SCX state-discovery algorithm). If ϕ is δ-weak, then*

$$\mathbb{P}(\hat{S} \ne S) \;\ge\; \frac{H(S) - \delta - \log 2}{\log K_{\mathcal{S}}}. \tag{27}$$

Proof. Apply Fano's inequality [1, 2]:

$$\begin{aligned}\mathbb{P}(\hat{S} \ne S) &\ge \frac{H(S \mid \phi(X)) - \log 2}{\log K_{\mathcal{S}}} \\ &= \frac{H(S) - I(\phi(X); S) - \log 2}{\log K_{\mathcal{S}}} \\ &= \frac{H(S) - \delta - \log 2}{\log K_{\mathcal{S}}}.\end{aligned}$$

□

Corollary 4.3 (Uniform states). *If the true states are approximately uniform, $H(S) \approx \log K_{\mathcal{S}}$, then*

$$\mathbb{P}(\hat{S} \ne S) \;\ge\; 1 - \frac{\delta + \log 2}{\log K_{\mathcal{S}}}.$$

As $\delta \to 0$ with $K_{\mathcal{S}}$ fixed, $\mathbb{P}(\hat{S} \ne S) \to 1 - \frac{\log 2}{\log K_{\mathcal{S}}} > 0$. For large $K_{\mathcal{S}}$ this limit approaches 1.

Corollary 4.4 (Two-state case). *When $K_{\mathcal{S}} = 2$,*

$$\mathbb{P}(\hat{S} \neq S) \;\geq\; \frac{H(S) - \delta - \log 2}{\log 2}.$$

Remark 4.5 (Achievability). The converse of Fano's inequality guarantees the *existence* of an estimator $\hat{S}$ operating on ϕ such that

$$\mathbb{P}(\hat{S} \neq S) \;\leq\; \frac{\delta + \log 2}{\log K_{\mathcal{S}}}.$$

This existence bound does not guarantee that a particular algorithm (e.g. k-means) attains it. In practice, the SCX clustering error satisfies

$$\mathbb{P}(\hat{S} \neq S) \;\leq\; O\!\left(\frac{\delta}{\log K_{\mathcal{S}}}\right) + \varepsilon_{k\text{-means}},$$

where $\varepsilon_{k\text{-means}}$ captures the approximation error of k-means (initialisation, non-convexity, non-spherical clusters). When δ is small (weak features), *any* algorithm is fundamentally limited by the Fano lower bound.

4.3 Lemma 2: SCX Reliability Degradation

Lemma 4.6 (SCX reliability estimation degradation). *Let $C(s)$ be the consensus score for true state s, and let $C(\hat{s})$ be the consensus score for an estimated state $\hat{s}$ obtained by clustering ϕ. If ϕ is δ-weak, then*

$$\mathbb{E}\big[|C(\hat{S}) - C(S)|\big] \;\leq\; 2 \cdot \mathbb{P}(\hat{S} \neq S) + O\!\left(\frac{1}{\sqrt{n_{\min}}}\right), \tag{28}$$

where $n_{\min} = \min_{s \in \mathcal{S}} n_s$ is the minimum per-state sample size.

Proof. Decompose the expectation into correctly and incorrectly estimated states:

$$\begin{aligned}\mathbb{E}\big[|C(\hat{S}) - C(S)|\big] &= \mathbb{P}(\hat{S} = S)\,\mathbb{E}\big[|C(\hat{S}) - C(S)| \mid \hat{S} = S\big] \\ &\quad + \mathbb{P}(\hat{S} \neq S)\,\mathbb{E}\big[|C(\hat{S}) - C(S)| \mid \hat{S} \neq S\big].\end{aligned}$$

Case 1 ($\hat{S} = S$): When the state is correctly estimated, $C(\hat{S})$ is a consistent estimator of $C(S)$. By Chebyshev's inequality and $C(\cdot) \in [0, 1]$,

$$\mathbb{E}\big[|C(\hat{S}) - C(S)| \mid \hat{S} = S\big] \;\leq\; \frac{\sigma_C}{\sqrt{n_s}} \;\leq\; O\!\left(\frac{1}{\sqrt{n_{\min}}}\right),$$

where σ_C is the standard deviation of C.

Case 2 ($\hat{S} \neq S$): When the state estimate is incorrect, $C(\hat{S})$ is a mixture of the consensus scores of the true states:

$$C(\hat{S}) = \sum_{s \in \mathcal{S}} \mathbb{P}(S = s \mid \hat{S})\, C(s).$$

Since $C(\cdot) \in [0, 1]$, the pointwise difference satisfies $|C(\hat{S}) - C(S)| \leq 1$, and consequently $\mathbb{E}[|C(\hat{S}) - C(S)| \mid \hat{S} \neq S] \leq 2$.

Combining the two cases,

$$\begin{aligned}\mathbb{E}\big[|C(\hat{S}) - C(S)|\big] &\leq \mathbb{P}(\hat{S} = S) \cdot O(1/\sqrt{n_{\min}}) + \mathbb{P}(\hat{S} \neq S) \cdot 2 \\ &\leq 2 \cdot \mathbb{P}(\hat{S} \neq S) + O(1/\sqrt{n_{\min}}).\end{aligned}$$

□

Corollary 4.7 (Degradation to global consistency). *As $\delta \to 0$, Lemma 4.2 gives $\mathbb{P}(\hat{S} \neq S) \to 1 - \frac{\log 2}{\log K_{\mathcal{S}}}$. In this limit,*

$$C(\hat{S}) \xrightarrow{p} \sum_{s \in \mathcal{S}} \rho(s)\, C(s) = \bar{C},$$

where $\rho(s) = \mathbb{P}(S = s)$ and $\bar{C}$ is the global average consistency. When features carry no state information, all estimated states converge to the same consistency value $\bar{C}$.

Corollary 4.8 (Noise score degradation). *Under the same conditions and the additional assumption that the estimated states are approximately balanced ($\max_{\hat{s}} \rho(\hat{s}) / \min_{\hat{s}} \rho(\hat{s}) \leq R$ for some constant R), the SCX noise score degenerates to*

$$\mathrm{NS}(x) \propto r(x),$$

where $r(x)$ is the residual of sample x. The SCX detector therefore reduces to the loss-based baseline detector. If the estimated states are highly imbalanced, the degradation is approximate with an additional $O(1/\sqrt{n})$ bias from state proportion estimation error.

4.4 Proof of Theorem 2: AUC, PR-AUC, and F1 Bounds

Theorem 4.9 (Weak feature failure lower bound). *Let $\phi : \mathcal{X} \to \Phi$ be a δ-weak feature map with respect to the true state S, i.e. $I(\phi(X); S) \leq \delta$ (in nats). Let h_{SCX} be the SCX noise detector operating under the standard pipeline (clustering on ϕ, state-conditional reliability estimates, noise scoring). Assume the estimated states are approximately balanced ($\max_{\hat{s}} \rho(\hat{s}) / \min_{\hat{s}} \rho(\hat{s}) \leq R$) and the clustering algorithm does not introduce error beyond the Fano lower bound. Then:*

(a) ***AUC bound:***

$$\boxed{\mathrm{AUC}(h_{SCX}) \;\leq\; \mathrm{AUC}_{base} + \sqrt{\frac{\delta}{2}} \cdot \left(\frac{1}{\eta} + \frac{1}{1-\eta}\right)} \tag{29}$$

(b) ***PR-AUC bound:***

$$\boxed{\mathrm{PRAUC}(h_{SCX}) \;\leq\; \mathrm{PRAUC}_{base} + \sqrt{\frac{\delta}{2}} \cdot \left(\frac{1}{\eta} + \frac{1}{1-\eta}\right)} \tag{30}$$

(c) ***F1 bound:*** *Assume the detector operates in a regime where both precision and recall are bounded below by $\varepsilon > 0$ (typically $\varepsilon = 0.1$). Then there exists a constant $C_F = 2/\varepsilon^2$ such that*

$$\boxed{\mathrm{F1}(h_{SCX}) \;\leq\; \mathrm{F1}_{base} + C_F \cdot \sqrt{\frac{\delta}{2}}} \tag{31}$$

Here AUC_{base}, PRAUC_{base}, and $\mathrm{F1}_{base}$ are the performance metrics of the loss-based baseline detector that thresholds $\max_m \ell(f_m(x), y)$ without using ϕ or state information.

Proof. The proof proceeds in six steps.

Step 1: Construct an auxiliary distribution $\tilde{P}$. Define $\tilde{P}$ by forcing ϕ and S to be independent while preserving their marginal distributions:

$$\tilde{P}(\phi, S) = P(\phi)\, P(S). \tag{32}$$

All other conditional distributions (given S, the distribution of (X, Y), experts, etc.) are unchanged.

Step 2: Relate KL divergence and mutual information. The KL divergence between P and $\tilde{P}$ equals the mutual information between ϕ and S:

$$\mathrm{KL}(P \parallel \tilde{P}) = \iint P(\phi, S) \log \frac{P(\phi, S)}{P(\phi)P(S)}\, d\phi\, dS = I(\phi(X); S) = \delta.$$

By Pinsker's inequality [3],

$$\mathrm{TV}(P, \tilde{P}) \;\le\; \sqrt{\frac{\mathrm{KL}(P \parallel \tilde{P})}{2}} = \sqrt{\frac{\delta}{2}}. \tag{33}$$

Step 3: Transfer the TV bound to the prediction distribution. Let P_{pred} and $\tilde{P}_{\text{pred}}$ be the joint distributions of $(\hat{z}_{\text{SCX}}(X), Z)$ under P and $\tilde{P}$, respectively. Since the SCX noise score is a deterministic function of X (and hence of ϕ and S), the data processing inequality [2] yields

$$\mathrm{TV}(P_{\text{pred}}, \tilde{P}_{\text{pred}}) \;\le\; \mathrm{TV}(P, \tilde{P}) \;\le\; \sqrt{\frac{\delta}{2}}. \tag{34}$$

Step 4: Analyse SCX under $\tilde{P}$. Under $\tilde{P}$, $\phi \perp S$, so ϕ carries no information about S. By Corollary 4.8, the SCX noise score degenerates to the residual:

$$\mathrm{NS}_{\tilde{P}}(x) \propto \max_m \ell(f_m(x), y).$$

Hence the SCX detector under $\tilde{P}$ is equivalent to the loss-based baseline detector:

$$\mathrm{AUC}_{\tilde{P}}(h_{\text{SCX}}) = \mathrm{AUC}_{\text{base}}, \tag{35}$$

$$\mathrm{PRAUC}_{\tilde{P}}(h_{\text{SCX}}) = \mathrm{PRAUC}_{\text{base}}, \tag{36}$$

$$\mathrm{F1}_{\tilde{P}}(h_{\text{SCX}}) = \mathrm{F1}_{\text{base}}\,. \tag{37}$$

Step 5: Convert TV bounds to AUC/PR-AUC/F1 bounds. *AUC.* The AUC can be written as $\mathrm{AUC} = \mathbb{P}(\text{score}_{\text{noise}} > \text{score}_{\text{clean}})$, which involves two independent samples: one from $P(\text{score} \mid Z = 1)$ (noise) and one from $P(\text{score} \mid Z = 0)$ (clean). For product measures,

$$\begin{aligned}
&\mathrm{TV}\big(P(\cdot \mid Z=1) \times P(\cdot \mid Z=0),\ \tilde{P}(\cdot \mid Z=1) \times \tilde{P}(\cdot \mid Z=0)\big) \\
&\quad \le \mathrm{TV}(P(\cdot \mid Z=1), \tilde{P}(\cdot \mid Z=1)) + \mathrm{TV}(P(\cdot \mid Z=0), \tilde{P}(\cdot \mid Z=0)).
\end{aligned}$$

We transfer the marginal TV bound to the conditional distributions. For any event A,

$$\begin{aligned}
|P(A \mid Z=1) - \tilde{P}(A \mid Z=1)| &= \frac{|P(A \cap \{Z=1\}) - \tilde{P}(A \cap \{Z=1\})|}{\eta} \\
&\le \frac{\mathrm{TV}(P, \tilde{P})}{\eta} \;\le\; \frac{1}{\eta}\sqrt{\frac{\delta}{2}}.
\end{aligned}$$

Similarly, $\mathrm{TV}(P(\cdot \mid Z=0), \tilde{P}(\cdot \mid Z=0)) \leq \frac{1}{1-\eta}\sqrt{\frac{\delta}{2}}$. Combining,

$$
\begin{aligned}
|\,\mathrm{AUC}_P(h) - \mathrm{AUC}_{\tilde{P}}(h)| &\leq \mathrm{TV}(P(\cdot \mid Z=1), \tilde{P}(\cdot \mid Z=1)) + \mathrm{TV}(P(\cdot \mid Z=0), \tilde{P}(\cdot \mid Z=0)) \\
&\leq \sqrt{\frac{\delta}{2}} \cdot \left(\frac{1}{\eta} + \frac{1}{1-\eta}\right).
\end{aligned}
$$

PR-AUC. The PR-AUC involves an integral over thresholds of precision, $\mathbb{P}(Z=1 \mid \hat{z}=1)$. The same amplification argument applies at each threshold, yielding the identical bound (30). (A fully rigorous treatment requires additional steps to handle the integral over thresholds, but the bound at the level of the joint distribution is valid.)

F1. The F1 score is a function of the joint distribution $P(\hat{z}, Z)$ only:

$$
\mathrm{F1} = \frac{2 \cdot \mathrm{TP}}{2 \cdot \mathrm{TP} + \mathrm{FP} + \mathrm{FN}},
$$

where $\mathrm{TP} = \mathbb{P}(\hat{z}=1, Z=1)$, $\mathrm{FP} = \mathbb{P}(\hat{z}=1, Z=0)$, and $\mathrm{FN} = \mathbb{P}(\hat{z}=0, Z=1)$. Since F1 does not involve conditional sampling, the TV bound applies directly without amplification.

The F1 score is Lipschitz continuous in $(\mathrm{TP}, \mathrm{FP}, \mathrm{FN})$. Its partial derivatives satisfy

$$
\left|\frac{\partial\, \mathrm{F1}}{\partial\, \mathrm{TP}}\right| = \frac{2(\mathrm{FP} + \mathrm{FN})}{(2\,\mathrm{TP} + \mathrm{FP} + \mathrm{FN})^2} \leq \frac{2}{\min(\mathrm{TP}, 1)},
$$

and similarly for FP, FN. Let $p_{\min} = \min(2\,\mathrm{TP} + \mathrm{FP} + \mathrm{FN})$. The Lipschitz constant satisfies $C_F \leq 2/p_{\min}^2$.

In the typical operating range where precision and recall are both at least 0.1, direct calculation gives $C_F \leq 2$. More conservatively:

- Precision, Recall ≥ 0.1: $C_F \leq 3$,
- Precision, Recall ≥ 0.5: $C_F \leq 1$.

Hence

$$
|\,\mathrm{F1}_P(h) - \mathrm{F1}_{\tilde{P}}(h)| \;\leq\; C_F \cdot \mathrm{TV}(P_{\text{pred}}, \tilde{P}_{\text{pred}}) \;\leq\; C_F \cdot \sqrt{\frac{\delta}{2}}. \tag{38}
$$

Step 6: Combine the bounds. From (35) and (**??**),

$$
\begin{aligned}
\mathrm{AUC}_P(h_{\mathrm{SCX}}) &\leq \mathrm{AUC}_{\tilde{P}}(h_{\mathrm{SCX}}) + \sqrt{\frac{\delta}{2}} \cdot \left(\frac{1}{\eta} + \frac{1}{1-\eta}\right) \\
&= \mathrm{AUC}_{\text{base}} + \sqrt{\frac{\delta}{2}} \cdot \left(\frac{1}{\eta} + \frac{1}{1-\eta}\right).
\end{aligned}
$$

Similarly,

$$
\begin{aligned}
\mathrm{PRAUC}_P(h_{\mathrm{SCX}}) &\leq \mathrm{PRAUC}_{\text{base}} + \sqrt{\frac{\delta}{2}} \cdot \left(\frac{1}{\eta} + \frac{1}{1-\eta}\right), \\
\mathrm{F1}_P(h_{\mathrm{SCX}}) &\leq \mathrm{F1}_{\text{base}} + C_F \cdot \sqrt{\frac{\delta}{2}}.
\end{aligned}
$$

This completes the proof. □

Remark 4.10 (Symmetric lower bound). The theorem also admits a symmetric lower bound

$$\mathrm{AUC}(h_{\mathrm{SCX}}) \;\geq\; \mathrm{AUC}_{\mathrm{base}} - \sqrt{\frac{\delta}{2}} \cdot \left(\frac{1}{\eta} + \frac{1}{1-\eta}\right),$$

which follows from the same TV argument applied in the opposite direction. When $\delta = 0$, the SCX AUC is *exactly* equal to the loss baseline AUC – SCX can be neither better nor worse.

Remark 4.11 (Interpretation of the η dependence). When η is small (rare noise), the AUC and PR-AUC bounds become loose. This reflects the intrinsic difficulty of detecting rare noise: SCX requires stronger features (smaller δ) to reliably surpass the loss baseline. For $\eta = 0.5$ the factor $1/\eta + 1/(1-\eta) = 4$, the most favourable case; for $\eta \to 0$ or $\eta \to 1$ the factor diverges.

4.5 Corollaries

Corollary 1: Complete Failure ($\delta = 0$)

If $\phi(X)$ is independent of S (i.e. ϕ is completely uninformative), then

$$\mathrm{AUC}(h_{\mathrm{SCX}}) = \mathrm{AUC}_{\mathrm{base}}, \qquad \mathrm{F1}(h_{\mathrm{SCX}}) = \mathrm{F1}_{\mathrm{base}}.$$

If, in addition, the noise is loss-uninformative (e.g. uniform label noise), the loss baseline itself degenerates to random guessing:

$$\mathrm{AUC}_{\mathrm{base}} = 0.5, \qquad \mathrm{F1}_{\mathrm{base}} = \frac{2\eta}{1+\eta}.$$

Corollary 2: Loss-Uninformative Noise

Under uniform label noise, the loss distribution is independent of the noise indicator:

$$P(\ell \mid Z = 1) = P(\ell \mid Z = 0).$$

The loss baseline then reduces to the random baseline, and Theorem 2 gives

$$\mathrm{AUC}(h_{\mathrm{SCX}}) \leq 0.5 + \sqrt{\frac{\delta}{2}} \cdot \left(\frac{1}{\eta} + \frac{1}{1-\eta}\right), \tag{39}$$

$$\mathrm{PRAUC}(h_{\mathrm{SCX}}) \leq \eta + \sqrt{\frac{\delta}{2}} \cdot \left(\frac{1}{\eta} + \frac{1}{1-\eta}\right), \tag{40}$$

$$\mathrm{F1}(h_{\mathrm{SCX}}) \leq \frac{2\eta}{1+\eta} + C_F\sqrt{\frac{\delta}{2}}. \tag{41}$$

When $\delta = 0$, the PR-AUC is bounded by η and the F1 by $2\eta/(1+\eta)$.

Corollary 3: Minimum Information Threshold

Suppose we require the SCX detection F1 to exceed the loss baseline by at least $\Delta > 0$. The minimum mutual information necessary is

$$\delta_{\min} \;\geq\; 2\left(\frac{\Delta}{C_F}\right)^2. \tag{42}$$

For AUC improvements,

$$\delta_{\min} \;\geq\; 2\big(\Delta_{\mathrm{AUC}} \cdot \eta(1-\eta)\big)^2. \tag{43}$$

Example 4.12. To achieve an F1 improvement of $\Delta = 0.05$ with $C_F = 2$,

$$\delta_{\min} \;\geq\; 2\left(\frac{0.05}{2}\right)^2 = 2 \cdot 0.000625 = 1.25 \times 10^{-3} \text{ nats.}$$

To achieve $\Delta = 0.1$ with $C_F = 2$,

$$\delta_{\min} \;\geq\; 2\left(\frac{0.1}{2}\right)^2 = 5 \times 10^{-3} \text{ nats.}$$

Since δ is mutual information in nats and $\delta \leq \log K_{\mathcal{S}}$, these thresholds are modest: $\delta = 10^{-3}$ corresponds to a likelihood ratio of $\exp(10^{-3}) \approx 1.001$ – very little state information is needed for a measurable improvement.

Corollary 4: Effect of State Count $K_{\mathcal{S}}$

As the number of true states $K_{\mathcal{S}}$ grows, maintaining the same $\delta = I(\phi; S)$ requires progressively stronger features. The normalised weakness $\varepsilon_\phi = \delta / \log K_{\mathcal{S}}$ determines SCX effectiveness:

- $\varepsilon_\phi \approx 1$: features carry almost no state information; SCX degenerates.
- $\varepsilon_\phi \approx 0$: features carry full state information; SCX can be effective.
- $\varepsilon_\phi \approx 0.5$: state discovery is about half accurate; noise detection is partially effective.

Blindly increasing the number of states $K_{\mathcal{S}}$ is therefore counterproductive without a corresponding increase in feature information δ.

4.6 Practical Diagnostics

The following diagnostics operationalise Theorem 2 for practitioners.

Diagnostic 1: Within-state consistency convergence. If all estimated states have near-equal consensus scores $C(\hat{s})$, the features may be too weak:

$$\mathrm{Var}\big(\{C(\hat{s})\}_{\hat{s} \in \hat{\mathcal{S}}}\big) < \tau_C.$$

Diagnostic 2: Random baseline comparison. Compute the loss baseline AUC and PR-AUC on the training set. If they are close to the random baseline (AUC ≈ 0.5, PRAUC $\approx \eta$), the noise detection task is intrinsically difficult, and weak features further prevent SCX from surpassing this baseline.

Diagnostic 3: Supervised state matching (when available). If ground-truth state labels exist for a subset (e.g. simulated data or known batch sources), compute the adjusted Rand index (ARI) between the ϕ-based clustering and the true states:

$$\mathrm{ARI}(\hat{S}, S) < 0.1 \implies \text{features too weak.}$$

Diagnostic 4: Mutual information estimation. Estimate $I(\phi(X);Y)$ via a Kraskov k-NN estimator or MINE, and compare with $\log K_{\mathcal{S}}$:

$$\frac{\hat{I}(\phi(X);Y)}{\log K_{\mathcal{S}}} < 0.2 \implies \text{features may be too weak.}$$

Threshold guidelines. Based on Theorem 2 and Corollary 3, we recommend the operational thresholds in Table 2.

Table 2: Practical guidelines for feature strength assessment.

Condition	Recommendation
$\varepsilon_\phi < 0.2$	SCX likely effective
$0.2 \leq \varepsilon_\phi \leq 0.5$	SCX partially effective; check other metrics
$\varepsilon_\phi > 0.5$	SCX not recommended; strengthen features first
$\text{AUC}_{\text{base}} < 0.55$ and $\varepsilon_\phi > 0.5$	SCX noise detection will fail
$\text{AUC}_{\text{base}} > 0.7$	Loss baseline already strong; SCX may help marginally
$\eta < 0.1$ and $\varepsilon_\phi > 0.3$	Rare noise amplifies AUC bound; need stronger features
$0.3 \leq \eta \leq 0.7$	Moderate noise rate; AUC/PR-AUC bounds most favourable

Here $\varepsilon_\phi = I(\phi;S)/\log K_{\mathcal{S}}$ and $\eta = \mathbb{P}(Z=1)$ is the marginal noise rate.

References

[1] R. M. Fano, *Transmission of Information: A Statistical Theory of Communications.* MIT Press, 1961.

[2] T. M. Cover and J. A. Thomas, *Elements of Information Theory*, 2nd ed. Wiley, 2006.

[3] M. S. Pinsker, *Information and Information Stability of Random Variables and Processes.* Holden-Day, 1964.

5 Theorem 3: The Honest Person Theorem (The Honest Person Theorem)

5.1 Statement and Implications

The central discovery of this paper is that distinguishing label noise from intrinsic sample difficulty is mathematically impossible from observational data alone. Theorem 3 formalizes this impossibility.

Theorem 5.1 (The Honest Person Theorem (The Honest Person Theorem))**.** *For any $K \geq 2$ classification problem, any $M \geq 1$ experts, and any finite state space $\mathcal{S}$, there exist two data-generating processes $\mathcal{P}_{noise}$ and $\mathcal{P}_{hard}$ such that:*

*(i) Under $\mathcal{P}_{noise}$, label errors are caused by **label noise**—the observed label differs from the true label.*

(ii) *Under $\mathcal{P}_{hard}$, all observed labels equal the true labels, but some samples are intrinsically* ***difficult****—expert errors arise from label ambiguity, not from label errors.*

(iii) *The two processes produce identical observable joint distributions* ***if*** *the noise rate η_{err} in $\mathcal{P}_{noise}$ equals the ambiguity rate η_{amb} in $\mathcal{P}_{hard}$. (Note: only the sufficiency direction is proved—the constructed worlds are observationally equivalent when $\eta_{err} = \eta_{amb}$; necessity, i.e. whether observational equivalence forces equality of the two rates for* all possible *world pairs, is* ***not established****.) Under this symmetry condition $\eta_{err} = \eta_{amb} = \eta$:*

$$\mathcal{P}_{noise}\big(x,\, y,\, \{f_m(x)\}\big) = \mathcal{P}_{hard}\big(x,\, y,\, \{f_m(x)\}\big), \qquad \forall\, (x, y, \{f_m\}) \in \mathcal{X} \times \mathcal{Y} \times \mathcal{Y}^M.$$

(iv) *Consequently, for any algorithm $\mathcal{A}$ that maps n i.i.d. observations to per-sample noise flags, the expected error rate on the ambiguous subset satisfies*

$$\max\big(\mathrm{Error}_{\mathcal{P}_{noise}}(\mathcal{A}),\, \mathrm{Error}_{\mathcal{P}_{hard}}(\mathcal{A})\big) \geq \frac{\eta\rho}{2},$$

where η is the noise rate and ρ is the proportion of the ambiguous state. The bound is strictly positive whenever $\eta > 0$ and $\rho > 0$, proving that the two worlds cannot be perfectly distinguished.

Remark 5.2. Theorem 3 asserts that the two statements "this sample's label is wrong" and "this sample is so difficult that experts systematically err on it" are observationally equivalent. The SCX framework must introduce assumptions (A1)–(A6) to break this equivalence; Theorem 3 shows that these assumptions are not arbitrary technical conditions but rather the minimal structure required for identifiability. See §5.5 for the minimal sufficient set.

5.2 Proof for $K = 2$ (Binary Classification)

We prove Theorem 3 by explicit construction. Fix $K = 2$, label space $\mathcal{Y} = \{0, 1\}$, $M \geq 1$ experts, and two states $s_1, s_2 \in \mathcal{S}$. Let $\rho \in (0, 1)$ and $\eta \in (0, \frac{1}{2})$ be parameters to be specified.

5.2.1 World A: Noise-Driven Construction

In the noise world $\mathcal{P}_{\text{noise}}$:

- **State** s_1 ("clean but noisy"), with $\mathbb{P}(X \in s_1) = \rho$. All samples have true label $y^* \equiv 0$. Labels are flipped with probability η_{err}:

  $$\mathbb{P}(y \neq y^* \mid x \in s_1) = \eta_{\text{err}}.$$

  Expert accuracy on clean samples is $1 - \varepsilon_1$:

  $$\mathbb{P}(f_m(x) = y^* \mid x \in s_1,\ \text{clean}) = 1 - \varepsilon_1, \qquad \varepsilon_1 \in (0, \tfrac{1}{2}).$$

- **State** s_2 ("clean and easy"), with $\mathbb{P}(X \in s_2) = 1 - \rho$. True label $y^* \equiv 0$, no noise ($\eta = 0$). Expert accuracy is $1 - \varepsilon_2$ with $\varepsilon_1 < \varepsilon_2 \leq \frac{1}{2}$ (the state is genuinely harder for experts).

Expert predictions depend only on the true label y^* (through training on clean data) and are conditionally independent given x.

5.2.2 World B: Difficulty-Driven Construction

In the hardness world $\mathcal{P}_{\text{hard}}$:

- **State** s_1, with the same marginal probability ρ. All observed labels are correct: $y = y^*$ for every sample. However, the true label itself is random:

  $$\mathbb{P}(y^* = 0 \mid x \in s_1) = 1 - \eta_{\text{amb}}, \qquad \mathbb{P}(y^* = 1 \mid x \in s_1) = \eta_{\text{amb}}.$$

  Experts are biased toward class 0: for every m,

  $$\mathbb{P}(f_m(x) = 0 \mid x \in s_1,\, y^* = 0) = 1 - \varepsilon_1, \qquad \mathbb{P}(f_m(x) = 0 \mid x \in s_1,\, y^* = 1) = 1 - \varepsilon_1.$$

  Thus when $y^* = 0$ the expert accuracy is $1 - \varepsilon_1$ (good); when $y^* = 1$ the expert accuracy is ε_1 (systematically wrong).

- **State** s_2, identical to World A: $y = y^* \equiv 0$, expert accuracy $1 - \varepsilon_2$.

5.2.3 Observational Equivalence Verification

The observable joint distribution factorises as

$$\mathcal{P}(x, y, f_1, \ldots, f_M) = \mathcal{P}(x) \cdot \mathcal{P}(y \mid x) \cdot \prod_{m=1}^{M} \mathcal{P}(f_m \mid x),$$

using conditional independence of experts given x (Assumptions A1–A2).

Marginal $\mathcal{P}(x)$. In both worlds $\mathbb{P}(X \in s_1) = \rho$, $\mathbb{P}(X \in s_2) = 1 - \rho$. Identical.

Conditional label distribution $\mathcal{P}(y \mid x)$. In World A (state s_1):

$$\mathcal{P}_{\text{noise}}(y = 0 \mid s_1) = (1 - \eta_{\text{err}}) \cdot 1 + \eta_{\text{err}} \cdot 0 = 1 - \eta_{\text{err}},$$
$$\mathcal{P}_{\text{noise}}(y = 1 \mid s_1) = (1 - \eta_{\text{err}}) \cdot 0 + \eta_{\text{err}} \cdot 1 = \eta_{\text{err}}.$$

In World B (state s_1), $y = y^*$ directly:

$$\mathcal{P}_{\text{hard}}(y = 0 \mid s_1) = 1 - \eta_{\text{amb}}, \qquad \mathcal{P}_{\text{hard}}(y = 1 \mid s_1) = \eta_{\text{amb}}.$$

For state s_2, $\mathcal{P}(y = 0 \mid s_2) = 1$, $\mathcal{P}(y = 1 \mid s_2) = 0$ in both worlds. Hence $\mathcal{P}(y \mid x)$ is identical across worlds.

Conditional expert distribution $\mathcal{P}(f_m \mid x)$. In World A (state s_1), experts predict the true label $y^* \equiv 0$ with accuracy $1 - \varepsilon_1$:

$$\mathcal{P}_{\text{noise}}(f_m = 0 \mid s_1) = 1 - \varepsilon_1, \qquad \mathcal{P}_{\text{noise}}(f_m = 1 \mid s_1) = \varepsilon_1.$$

In World B (state s_1):

$$\begin{aligned}\mathcal{P}_{\text{hard}}(f_m = 0 \mid s_1) &= \mathbb{P}(y^* = 0 \mid s_1) \cdot \mathbb{P}(f_m = 0 \mid s_1, y^* = 0) \\ &\quad + \mathbb{P}(y^* = 1 \mid s_1) \cdot \mathbb{P}(f_m = 0 \mid s_1, y^* = 1) \\ &= (1 - \eta_{\text{amb}})(1 - \varepsilon_1) + \eta_{\text{amb}}(1 - \varepsilon_1) = 1 - \varepsilon_1, \\ \mathcal{P}_{\text{hard}}(f_m = 1 \mid s_1) &= (1 - \eta_{\text{amb}})\varepsilon_1 + \eta_{\text{amb}}\varepsilon_1 = \varepsilon_1.\end{aligned}$$

For state s_2, both worlds give $\mathcal{P}(f_m = 0 \mid s_2) = 1 - \varepsilon_2$, $\mathcal{P}(f_m = 1 \mid s_2) = \varepsilon_2$.

Joint identity. Combining the three components,

$$\mathcal{P}_{\text{noise}}(x, y, \{f_m\}) = \mathcal{P}_{\text{hard}}(x, y, \{f_m\}), \qquad \forall\,(x, y, \{f_m\}),$$

which proves part (iii) of Theorem 3. □

5.2.4 Algorithm Impossibility

Let $\mathcal{A}$ be any algorithm that uses n i.i.d. observations to output per-sample noise flags $\{z_i\}_{i=1}^n$, where $z_i = 1$ means "label is noisy." Since the observable distribution is identical in both worlds, the algorithm's output distribution is also identical:

$$\mathcal{L}_{\mathcal{P}_{\text{noise}}}(\mathcal{A}(D_n)) = \mathcal{L}_{\mathcal{P}_{\text{hard}}}(\mathcal{A}(D_n)).$$

However, the *correct* noise flags differ between worlds. In World A, the truly noisy samples are those in state s_1 with $y = 1$ (a proportion $\rho\eta$ of the data). In World B, *no* samples are noisy; the same subset consists of correct but difficult samples. The algorithm therefore cannot simultaneously achieve zero error in both worlds.

Let a be the expected fraction of the ambiguous subset $\{x \in s_1,\ y = 1\}$ that $\mathcal{A}$ flags as noisy. Then:

$$\text{Error}_{\mathcal{P}_{\text{noise}}}(\mathcal{A}) = \rho\eta(1 - a), \qquad \text{Error}_{\mathcal{P}_{\text{hard}}}(\mathcal{A}) = \rho\eta a.$$

Hence

$$\max\big(\text{Error}_{\mathcal{P}_{\text{noise}}},\, \text{Error}_{\mathcal{P}_{\text{hard}}}\big) \geq \frac{\text{Error}_{\text{noise}} + \text{Error}_{\text{hard}}}{2} = \frac{\rho\eta}{2}.$$

This proves part (iv). □

5.3 Extension to $K > 2$ (Random Expert Construction)

For $K > 2$, the binary construction does not generalise directly because the expert marginals in the two worlds no longer match automatically. We instead use a different construction where experts in the hardness world are *fully random* (conditionally independent of the true label).

World A (Noise). State s_1: $y^* \equiv 0$. Noise flips the label uniformly over the remaining $K - 1$ classes with probability η:

$$\mathbb{P}(y = c \mid \text{noise}, x \in s_1) = \frac{1}{K - 1}, \quad c \neq 0.$$

Expert f_m predicts 0 with probability $1-\varepsilon_1$, and each other class with probability $\varepsilon_1/(K-1)$. Thus

$$\mathcal{P}_{\text{noise}}(y = 0 \mid s_1) = 1 - \eta, \quad \mathcal{P}_{\text{noise}}(y = c \mid s_1) = \frac{\eta}{K - 1}\ (c \neq 0),$$

$$\mathcal{P}_{\text{noise}}(f_m = 0 \mid s_1) = 1 - \varepsilon_1, \quad \mathcal{P}_{\text{noise}}(f_m = c \mid s_1) = \frac{\varepsilon_1}{K - 1}\ (c \neq 0).$$

By Assumption A4, noise and expert predictions are independent given the state, so $y \perp f_m \mid s_1$ in World A.

World B (Hardness). State s_1: all labels are correct ($y = y^*$). The true label distribution matches the World A observed distribution:

$$\mathbb{P}(y^* = 0 \mid s_1) = 1 - \eta, \quad \mathbb{P}(y^* = c \mid s_1) = \frac{\eta}{K-1} \ (c \neq 0).$$

Critical construction: experts are fully random, independent of the true label:

$$\mathbb{P}(f_m = 0 \mid s_1) = 1 - \varepsilon_1, \quad \mathbb{P}(f_m = c \mid s_1) = \frac{\varepsilon_1}{K-1} \ (c \neq 0),$$

with $f_m \perp y^* \mid s_1$ and $f_m \perp y \mid s_1$ (since $y = y^*$).

Equivalence. The marginal distributions of y, f_m, and the joint $(y, \{f_m\})$ are identical in both worlds because both y and $\{f_m\}$ are independent products of the same marginal distributions. Hence for any $K \geq 2$:

$$\mathcal{P}_{\text{noise}}(x, y, \{f_m\}) = \mathcal{P}_{\text{hard}}(x, y, \{f_m\}).$$

The construction is extreme—experts in World B are equivalent to random guessers—but it suffices as an existence proof: there exists at least one "hardness" interpretation observationally indistinguishable from the "noise" interpretation. □

5.4 Corollaries: Which Assumptions Break Unidentifiability

Theorem 3 shows that without additional assumptions, noise and difficulty are indistinguishable. Each corollary below identifies a *sufficient condition* that, if known to hold, breaks the equivalence.

5.4.1 Anchor Labels (Corollary 1)

Corollary 5.3 (Anchor Labels Break Unidentifiability). *If for a known subset $\mathcal{X}_{anchor} \subset \mathcal{X}$ the researcher has access to the true labels y^* (e.g. via human review), then noise and difficulty are distinguishable on $\mathcal{X}_{anchor}$. If $\mathcal{X}_{anchor}$ covers all states, the problem is globally identifiable.*

Sketch. On anchor points, compare the observed label y with the true label y^*. If $y \neq y^*$, the sample is necessarily noisy. If $y = y^*$ but $f_m(x) \neq y$, the sample is necessarily difficult (the expert errs on a correctly labelled sample). Extrapolating expert error rates from anchor points to non-anchor points breaks global unidentifiability. □

This corresponds to the SCX cold-start protocol, where a small number of human-reviewed samples serve as anchors for initialising state-level noise rate estimates.

5.4.2 Disjoint Training Data (Corollary 2)

Corollary 5.4 (Known Disjoint Training Data Break Unidentifiability). *If M experts are trained on M disjoint data subsets (Assumption A1) and the researcher knows the training set composition, then noise and difficulty are distinguishable.*

Sketch. Disjoint training implies expert errors are conditionally independent given x. On noisy samples, *all* experts have elevated error rates simultaneously (because they all predict the true label y^*, while the observed label y differs). On difficult samples, expert errors are independent. The average error rate $\bar{e}(x) = \frac{1}{M}\sum_m \mathbf{1}\{f_m(x) \neq y\}$ satisfies

$$\mathbb{E}[\bar{e}(x) \mid \text{noise}] = 1 - \frac{1 - \mathbb{E}[\bar{e}(x) \mid \text{clean}]}{K-1}$$

(Theorem 1, Lemma 1). This relation cannot hold simultaneously in a pure hardness world, providing a statistical test. □

5.4.3 Input-Independent Noise (Corollary 3)

Corollary 5.5 (Input-Independent Noise Rate Breaks Unidentifiability)**.** *If the label noise rate $\eta(x) = \mathbb{P}(Y \neq Y^* \mid X = x)$ is constant across x (Assumption A4), and the expert error rates differ between states, then noise and difficulty are distinguishable.*

Sketch. When $\eta(x) \equiv \eta$ is constant, the correlation pattern between the observed label y and expert errors $\{f_m\}$ differs fundamentally between the two worlds. In the noise world, $\mathbb{E}[\bar{e}(x)]$ is constant across states (marginalising over the label-flip event). In the hardness world, $\bar{e}(x)$ varies with the state-dependent true label distribution. Testing for input dependence of $\bar{e}(x)$ conditioned on y distinguishes the two. □

5.4.4 State Homogeneity (Corollary 4)

Corollary 5.6 (State Homogeneity Breaks Unidentifiability)**.** *If within each state s the expected expert error rate on clean samples is constant (Assumption A5), and states differ in this rate, then noise and difficulty are distinguishable.*

Sketch. In the hardness world construction (World B), state s_1 contains two subpopulations—samples with $y^* = 0$ (expert accuracy $1 - \varepsilon_1$) and samples with $y^* = 1$ (expert accuracy ε_1). This violates state homogeneity, which requires a single per-state error rate. Therefore, if Assumption A5 holds, World B's construction is invalid, and the unidentifiability is broken. □

5.4.5 Balanced Error Distribution (Corollary 5)

Corollary 5.7 (Balanced Error Distribution Breaks Unidentifiability)**.** *If expert errors are uniformly distributed across all error classes (Assumption A6 with $C_{bal} = 1$), then noise and difficulty are distinguishable.*

Sketch. In World B, experts are biased toward class 0: when $y^* = 1$, they predict class 0 with probability $1 - \varepsilon_1$. This creates a highly imbalanced error pattern (errors are almost exclusively in the $1 \to 0$ direction), violating the balanced error assumption. If Assumption A6 is known to hold, World B is excluded. □

5.5 Minimal Sufficient Assumption Set

Theorem 3's construction relies on specific conditions. To exclude the construction—and thereby break unidentifiability—at least one of the following combinations of assumptions is necessary and sufficient:

$$
\begin{aligned}
\mathcal{A}^{(1)}_{\min} &= \{A1,\ A4,\ A5\},\\
\mathcal{A}^{(2)}_{\min} &= \{A1,\ A4,\ A6\},\\
\mathcal{A}^{(3)}_{\min} &= \{A5,\ A6\} \quad \text{with } |\mathcal{S}| \geq 2.
\end{aligned}
$$

The SCX framework operates under $\mathcal{A}^{(1)}_{\min}$ (the standard set A1–A6 provides additional technical benefits for finite-sample bounds). Each assumption plays a distinct role:

- **A1** (disjoint training): ensures expert independence, enabling Corollary 2's statistical test.
- **A4** (uniform noise): removes the possibility of input-dependent noise mimicking difficulty (Corollary 3).
- **A5** (state homogeneity): prevents a single state from containing both easy and difficult subpopulations (Corollary 4).
- **A6** (balanced errors): prevents expert bias from creating pseudo-difficulty patterns (Corollary 5).

If *all* of these assumptions are violated, Theorem 3 applies fully: no algorithm can distinguish label noise from sample difficulty.

5.6 Corollary: The Everyone Equal Theorem (Everyone Equal Theorem)

Corollary 5.8 (The Everyone Equal Theorem (Everyone Equal Theorem))**.** *Let any observer O claim to distinguish label noise from sample difficulty using any algorithm $\mathcal{A}$. Then either*

(i) the structural assumptions relied upon by $\mathcal{A}$ are explicitly stated and independently verifiable by any other observer, or

(ii) the claim is mathematically empty—no amount of observational data can validate it, and any observer may legitimately reject it without further argument.

Furthermore, if (i) holds, the assumptions are verifiable by any observer without privileged access to O's data, model, or expertise. All observers are placed on equal epistemic footing.

Proof. Theorem 3 establishes $P_{W_A}(X,Y) = P_{W_B}(X,Y)$ for all observers, i.e., the joint distribution of inputs, labels, and multi-expert predictions is identical in the noise world and the difficulty world. Consequently, for any discrimination rule d,

$$
\max_{W \in \{W_A, W_B\}} P_W\big(d(X,Y) \neq W\big) \geq \frac{1}{2},
$$

as proved via Le Cam's lemma in the main text. This lower bound is observer-independent: it holds for *any* algorithm $\mathcal{A}$ deployed by *any* observer O, because the observational data do not encode observer identity.

Therefore, if O claims to achieve discrimination performance better than chance, the claim must rely on assumptions that break the observational equivalence $P_{W_A} = P_{W_B}$. These assumptions—the minimal sufficient set $\mathcal{A}_{\min}$ derived above—must be stated. Once stated, they are verifiable by any observer: checking disjoint training (A1), uniform noise (A4), state homogeneity (A5), and balanced errors (A6) requires only access to the training procedure and data distribution, not to O's internal state.

Hence no observer occupies a privileged epistemic position. The theorem places all observers—regardless of institutional authority, computational resources, or social status—on identical mathematical ground. □

Remark 5.9 (No privileged observer). The name "The Everyone Equal Theorem" (Everyone Equal Theorem) reflects the corollary's operational meaning: the mathematical structure of data quality assessment does not permit any person, algorithm, institution, or nation to claim inherent authority over the distinction between noise and difficulty. Authority must be earned through transparently stated, independently verifiable assumptions. This is the formal basis for the SCX framework's commitment to epistemic equality.

5.7 Conjecture: The Good Person Convergence Conjecture (Good Person Convergence)

Conjecture 5.10 (The Good Person Convergence Conjecture (Good Person Convergence Conjecture))**.** *Consider a community of N agents interacting under the SCX framework: each agent may produce data, claim data quality, or audit others' claims. Let the following conditions hold:*

(C1) ***The Honest Person Condition****: every data-quality claim must state its structural assumptions (Theorem 3);*

(C2) ***The Everyone Equal Condition****: no agent occupies a privileged epistemic position (Corollary 5.8);*

(C3) ***The No-Pretension Condition****: audit outputs are unconditionally verifiable by any third party (non-proliferation equilibrium, Theorem NPE).*

Then, as $t \to \infty$:

(i) The fraction of agents making unverifiable claims converges to zero—"evil-doers bow": dishonest behaviour becomes detectable and economically irrational;

(ii) The expected payoff of honest behaviour exceeds that of dishonest behaviour for all agents—"the honest stand tall": honesty is the strictly dominant strategy;

(iii) The community converges to an equilibrium in which every agent simultaneously satisfies the three conditions, and no agent can improve their position by deviating.

Remark 5.11 (Status of this claim). **This is a conjecture, not a theorem.** The original manuscript presented a 13-line "proof sketch" that merely cited Theorems 3, NPE, and Corollary 5.8 without establishing a formal logical chain from premises to conclusion. A rigorous proof would require:

(a) A formal agent model (type space, strategy space, payoff functions);

(b) A dynamic process description (discrete/continuous time, update rules);

(c) A rigorous definition of convergence (topological space, convergence mode);

(d) A Lyapunov or potential function establishing the stability of the all-honest equilibrium;

(e) Verification that the three conditions (C1–C3) are not circular—each is itself a consequence of other SCX theorems, so their joint satisfaction must be proved consistent.

Until such a proof is provided, the "Good Person Convergence" claim remains a compelling research programme rather than an established result. The intuition—that SCX's mathematical structure makes honesty strategically dominant—is provocative and merits further investigation, but it has not been formally demonstrated.

5.8 Epistemic Formalization System (E1-E5 + K Operator)

We formalize the classical epistemological question—"when does a community know a claim?"—as a probabilistic decision problem within the SCX framework. The key insight: *knowledge is a claim that survives multi-expert audit with provable statistical guarantees.*

5.8.1 Axioms for the Verification Framework

Definition 5.12 (Verification framework)**.** Let $\mathcal{S}$ be a claim with truth value $T(\mathcal{S}) \in \{0,1\}$ (unknown to verifiers). Let $\mathcal{C} = \{v_1, \ldots, v_M\}$ be M verifiers. Each v_m produces a noisy binary judgment $J_m(\mathcal{S}) \in \{0,1\}$ via access only to observational data $(X, Y, \{f_m\}_{m=1}^{M})$ as defined in the SCX framework.

Assumption 5.13 (Epistemic axioms E1–E5)**. E1 Assumption transparency:** the claim $\mathcal{S}$ is accompanied by an explicit assumption set $\mathcal{A}_{\mathcal{S}}$ (enforced by Theorem 3, The Honest Person Theorem). If $\mathcal{A}_{\mathcal{S}}$ is empty or unstated, $\mathcal{S}$ is rejected without evaluation.

E2 Assumption soundness: under $\mathcal{A}_{\mathcal{S}}$, verifier judgments satisfy $\mathbb{P}(J_m = T(\mathcal{S}) \mid \mathcal{A}_{\mathcal{S}}) = p_m > 1/2$. That is, each verifier is better than chance when assumptions hold. The verifier errors are conditionally independent given $\mathcal{A}_{\mathcal{S}}$ and $T(\mathcal{S})$.

E3 State-homogeneous verifier quality: within each state $s \in \mathcal{S}$, $\mathbb{P}(J_m \neq T(\mathcal{S}) \mid s) \leq \varepsilon_s < 1/2$, with $\Delta_s = 1/2 - \varepsilon_s > 0$.

E4 Observational indistinguishability: verifiers have access only to $(X, Y, \{f_m\})$—the same data as any SCX observer. They possess no oracle access to $T(\mathcal{S})$.

E5 Counterexample logging: the Yajie CEC $\mathcal{E}_t$ accumulates all adjudicated judgments, so that any counterexample detected at time t is permanently recorded.

5.8.2 The Knowledge Operator

Definition 5.14 (SCX-knowledge operator $\mathbf{K}$)**.** For claim $\mathcal{S}$, assumption set $\mathcal{A}_{\mathcal{S}}$, community $\mathcal{C}$ of size M, and state $s \in \mathcal{S}$, define the **verification score**:

$$V_M(\mathcal{S}, s) = \frac{1}{M} \sum_{m=1}^{M} J_m(\mathcal{S}; s)$$

where $J_m(\mathcal{S}; s)$ is verifier m's judgment on samples in state s. Define the **decision threshold** $\theta_s \in (0, 1)$.

A **counterexample** is a triple $(x, y, \{f_m(x)\})$ such that $\mathcal{A}_\mathcal{S}$ holds but $V_M(\mathcal{S}, s(x)) < \theta_{s(x)}$. Let $\mathcal{E}_t$ be the set of counterexamples accumulated up to time t. A counterexample $e \in \mathcal{E}_t$ is **resolved** if either (a) it has been adjudicated and found consistent with $\mathcal{A}_\mathcal{S}$ under additional scrutiny, or (b) the claim is amended to accommodate it. Otherwise it is **unresolved**.

The **SCX-knowledge operator** is defined per-state:

$$\mathbf{K}_t(\mathcal{S}, \mathcal{C}, s) = \begin{cases} 1 & \text{if } V_M(\mathcal{S}, s) \geq \theta_s \text{ and no unresolved counterexample in } \mathcal{E}_t \text{ involves state } s, \\ 0 & \text{otherwise.} \end{cases}$$

The claim-level knowledge operator aggregates state-level judgments: $\mathbf{K}_t(\mathcal{S}, \mathcal{C}) = \mathbf{1}\{\mathbf{K}_t(\mathcal{S}, \mathcal{C}, s) = 1 \text{ for all } s \in (ho)\}$.

5.8.3 Rigorous Guarantee

Theorem 5.15 (Epistemic Formalization System — E1-E5 + K Operator)**.** *Under the* ***SCX Axiom System (SCX Axiom System)*** *E1–E5 and the SCX framework conditions A1–A6 (Theorem 1), for any claim $\mathcal{S}$ with stated assumptions $\mathcal{A}_\mathcal{S}$ and any community $\mathcal{C}$ of M independent verifiers, the knowledge operator satisfies:*

$$\boxed{\mathbb{P}\big(\mathbf{K}(\mathcal{S}, \mathcal{C}) \neq T(\mathcal{S}) \mid \mathcal{A}_\mathcal{S}\big) \leq \sum_{s \in \mathcal{S}} \rho_s \exp\big(-2M\Delta_s^2\big)}$$

where ρ_s is the fraction of verification samples in state s, and $\Delta_s = \theta_s - (1 - p_s)$ is the separation between the decision threshold and the expected error rate. The error probability decays exponentially in the number of verifiers M.

Furthermore, this bound is tight: no knowledge operator operating under E1–E5 can achieve a smaller error exponent without additional structural assumptions, by Theorem 4 (The Extreme Precision Theorem).

Proof. The proof is a direct application of Theorem 1's concentration argument to the verification setting. For each state s, the verifier judgments $\{J_m(\mathcal{S}; s)\}_{m=1}^M$ are M independent Bernoulli trials with success probability p_s (E2, E3). If $T(\mathcal{S}) = 1$ (the claim is true), then $p_s \geq 1 - \varepsilon_s$ by E3. The verification score V_M is a sample mean of M independent Bernoulli variables. By Hoeffding's inequality and the state-conditioned decomposition of Theorem 1 (equation S1.7):

$$\mathbb{P}\big(V_M \leq \theta_s \mid T(\mathcal{S}) = 1\big) \leq \exp\big(-2M(\theta_s - p_s)^2\big) = \exp\big(-2M\Delta_s^2\big).$$

The same bound holds symmetrically when $T(\mathcal{S}) = 0$, with $\Delta_s = p_s - (1 - \theta_s)$.

The verifier errors are conditionally independent given the state (E3 strengthened by A2' from Theorem 1). The optional stopping argument from Theorem 1's proof extends to the sequential audit setting, ensuring that the bound does not degrade with multiple looks at the data.

For the tightness claim, Theorem 4 (The Extreme Precision Theorem) establishes that the exponential rate $2M\Delta_s^2$ is minimax-optimal: any decision rule operating on M verifier judgments incurs an error exponent at least this large. Hence the SCX-knowledge operator achieves the fundamental limit. □ □

5.8.4 Properties of K

Corollary 5.16 (Knowledge operator properties). *The SCX-knowledge operator* $\mathbf{K}$ *satisfies:*

(i) ***Non-triviality:*** $\mathbf{K}(\mathcal{S}) \neq 1$ *for any* $\mathcal{S}$ *with* $\mathcal{A}_{\mathcal{S}} = \emptyset$*, because E1 rejects unstated assumptions;*

(ii) ***Exponential reliability:*** $\mathbb{P}(\mathbf{K}$ *correct*$) \to 1$ *exponentially in* M*;*

(iii) ***Monotonicity under audit:*** *if* $\mathbf{K}_t(\mathcal{S}) = 1$*, then* $\mathbf{K}_{t+1}(\mathcal{S}) = 1$ *unless a counterexample emerges (E5), in which case* $\mathbf{K}$ *flips to* 0 *and stays there;*

(iv) ***Gettier immunity (operational):*** *Gettier (1963) cases arise when a claim is justified but true only accidentally—the justification chain contains an unstated contingency. Under E1, any such contingency must appear in* $\mathcal{A}_{\mathcal{S}}$*, where it becomes either (a) verified (the claim holds under it) or (b) falsified (the claim fails). This addresses the* operational *Gettier problem—cases where the contingency is expressible as a verifiable condition. It does* not *claim to resolve* all metaphysical *Gettier cases (e.g., those involving non-verifiable possible-worlds semantics);*

(v) ***Epistemic equality:*** $\mathbf{K}$ *depends only on* $\mathcal{A}_{\mathcal{S}}$ *and* $\{J_m\}$*, never on the identity, credentials, or institutional affiliation of the verifiers (Corollary 5.8).*

5.8.5 Connection to Classical Epistemology

Remark 5.17 (What we did to Plato). Plato (c. 369 BCE) defined knowledge as justified true belief. Gettier (1963) found counterexamples where justification and truth coincide accidentally due to unstated contingencies. The subsequent literature attempted to *redefine* justification.

The SCX approach does not redefine—it *operationalizes.* Knowledge is not a metaphysical state of an individual mind; it is the mathematical status of a claim that has survived multi-expert audit under stated assumptions with provable statistical guarantees. The bound $\exp(-2M\Delta_s^2)$ replaces philosophical debate with a quantitative threshold: "how many independent verifiers do we need?" becomes a solvable equation.

5.9 Connections to Prior Art

5.9.1 Measurement Error Models

In classical measurement error models [1, 2], the observed variable Y is modelled as $Y = Y^* + U$ where U is measurement error. Without validation data or instrumental variables, the distributions of Y^* and U are not separately identifiable from the marginal distribution of Y. Theorem 3 extends this classical result from continuous measurement error to the classification setting with multi-expert predictions. The key novelty is that even with the additional information $\{f_m(X)\}$—which one might expect to resolve the ambiguity—the unidentifiability persists.

5.9.2 Label Noise Transition Matrix Unidentifiability

The label noise literature [3, 4] established that the noise transition matrix $T_{ij} = \mathbb{P}(\tilde{Y} = j \mid Y^* = i)$ is not identifiable from $(\tilde{Y}, X)$ alone without anchor points or diagonal-dominance assumptions. Theorem 3 generalises this to the multi-expert setting, showing that even with M expert predictions, the ambiguity between noise and intrinsic difficulty remains.

5.9.3 Dawid-Skene Unidentifiability

The Dawid-Skene model [5] treats the true label as a latent variable and each annotator's confusion matrix as a parameter. Standard identifiability issues include label-switching symmetries. Theorem 3 addresses a different and more fundamental ambiguity: even after resolving label-switching, one cannot distinguish whether an annotator's error is due to noise in the label or difficulty of the sample. The two unidentifiability problems are orthogonal: both must be resolved for reliable label quality assessment.

5.9.4 Causal Structure Unidentifiability

From a causal perspective, Theorem 3 can be understood as a Markov equivalence class problem [6]. Two causal directed acyclic graphs generate identical observational distributions:

- **Noise graph:** $S \to Y^* \to Y \leftarrow$ Noise, $X \to \{f_m\}$, where Noise is an independent exogenous variable.
- **Difficulty graph:** $S \to Y^*(= Y)$, $X \to \{f_m\}$, where state S affects both the true label distribution and the expert error rates.

These two graphs are observationally indistinguishable, forming a Markov equivalence class of size 2.

References for Section S3

References

[1] Fuller, W. A. *Measurement Error Models.* Wiley (1987).

[2] Carroll, R. J., Ruppert, D., Stefanski, L. A. & Crainiceanu, C. M. *Measurement Error in Nonlinear Models: A Modern Perspective.* Chapman and Hall/CRC (2006).

[3] Menon, A. K., van Rooyen, B., Ong, C. S. & Williamson, R. C. Learning from corrupted binary labels via class-probability estimation. *Proc. ICML* (2015).

[4] Patrini, G., Rozza, A., Menon, A. K., Nock, R. & Qu, L. Making deep neural networks robust to label noise: A loss correction approach. *Proc. CVPR* (2017).

[5] Dawid, A. P. & Skene, A. M. Maximum likelihood estimation of observer error-rates using the EM algorithm. *Applied Statistics* **28**(1), 20–28 (1979).

[6] Pearl, J. *Causality: Models, Reasoning, and Inference.* Cambridge University Press, 2nd ed. (2009).

[7] Bahadur, R. R. & Rao, R. R. On deviations of the sample mean. *Annals of Mathematical Statistics* **31**(4), 1015–1027 (1960).

6 Theorem 4': Exact Constant Minimax Optimality

6.1 Theorem 4: Exact Constant Minimax Optimality (The Extreme Precision Theorem)

Theorem 1 establishes that the SCX noise detector achieves an exponential convergence rate $2M\Delta^2$. Theorem 4' refines this result in three ways: it replaces the Hoeffding rate $2\Delta^2$ with the exact Chernoff information κ, it establishes the *exact constant* multiplying the exponential, and it proves that no algorithm can improve upon this constant.

Throughout this section we work within a single state s satisfying Assumptions (A1)–(A6). Let

$$p_0 = \mu_s, \qquad p_1 = 1 - C_{\mathrm{bal}} \cdot \frac{\mu_s}{K-1},$$

where $K = K_{\mathcal{Y}}$ is the number of classes and $0 < p_0 < p_1 < 1$. The expert error indicators $e_m = \mathbf{1}\{f_m(x) \neq y\}$ are i.i.d. $\mathrm{Bern}(p_0)$ under the clean hypothesis H_0 and i.i.d. $\mathrm{Bern}(p_1)$ under the noise hypothesis H_1.

Define the following quantities:

$$\kappa = C\big(\mathrm{Bern}(p_0), \mathrm{Bern}(p_1)\big) \qquad \text{Chernoff information,}$$

$$\theta^* = \frac{\log \frac{1-p_0}{1-p_1}}{\log \frac{p_1(1-p_0)}{p_0(1-p_1)}} \qquad \text{Chernoff point (unique root of } \mathrm{KL}(\theta \| p_0) = \mathrm{KL}($$

$$\lambda_0^* = \log \frac{\theta^*(1-p_0)}{p_0(1-\theta^*)} > 0, \qquad \lambda_1^* = \log \frac{\theta^*(1-p_1)}{p_1(1-\theta^*)} < 0 \quad \text{Saddlepoints at } \theta^*,$$

$$D^* = \lambda_0^* + |\lambda_1^*| = \log \frac{p_1(1-p_0)}{p_0(1-p_1)} > 0,$$

$$s = \frac{|\lambda_1^*|}{D^*} \in (0,1).$$

Theorem 6.1 (Exact Constant Minimax Optimality of SCX)**.** *Under Assumptions (A1)–(A6), for any state s with $p_0 < p_1$, the following hold.*

(a) ***SCX achievability with adaptive threshold.*** *The SCX detector using threshold*

$$\theta^\dagger = \theta^* + \frac{1}{M} \frac{\log \frac{1-\eta}{\eta}}{D^*} + O\left(\frac{1}{M^2}\right)$$

achieves

$$\lim_{M \to \infty} e^{M\kappa} \sqrt{2\pi M} \left(1 - \mathrm{F1}_{SCX}(\theta^\dagger)\right) = \frac{C_{\min}}{\eta},$$

where the canonical minimax constant is

$$\boxed{C_{\min} = \frac{\eta}{2} \left(\frac{1-\eta}{\eta}\right)^s \frac{1/\lambda_0^* + 1/|\lambda_1^*|}{\sqrt{\theta^*(1-\theta^*)}}}.$$

(b) ***Minimax lower bound.*** *For* any *noise detection algorithm $\mathcal{A}$,*

$$\liminf_{M \to \infty} e^{M\kappa} \sqrt{2\pi M} \left(1 - \mathrm{F1}_{\mathcal{A}}\right) \geq \frac{C_{\min}}{\eta}.$$

(c) ***Constant optimality.*** *Since SCX with the adaptive threshold $\theta^\dagger$ attains the limit with constant $C_{\min}$, it is* exact constant minimax optimal*: no algorithm can achieve a smaller asymptotic error constant.*

(d) ***Multi-state generalisation.*** *For S states with proportions ρ_s, Chernoff information κ_s, and per-state constants C_s,*

$$\kappa_{global} = \min_{1\le s\le S} \kappa_s, \qquad C_{global} = \sum_{s:\,\kappa_s=\kappa_{global}} \rho_s C_s,$$

and $\liminf_{M\to\infty} e^{M\kappa_{global}}\sqrt{2\pi M}\,(1-\mathrm{F1}_{\mathcal{A}}) \ge C_{global}/\eta$.

The remainder of this section provides the complete proof, organised into lemmas that are of independent interest.

6.2 Preliminaries: Chernoff Information and Bahadur-Rao

6.2.1 Chernoff Information

For two distributions P_0, P_1 with densities $p_0(\cdot), p_1(\cdot)$, the Chernoff information is

$$C(P_0,P_1) = -\min_{\lambda\in[0,1]} \log \mathbb{E}_{P_0}\left[\left(\frac{dP_1}{dP_0}\right)^{\lambda}\right].$$

For $\mathrm{Bern}(p_0)$ vs. $\mathrm{Bern}(p_1)$, we have

$$\mathbb{E}_{P_0}\left[\left(\frac{dP_1}{dP_0}\right)^{\lambda}\right] = p_0^{1-\lambda}p_1^{\lambda} + (1-p_0)^{1-\lambda}(1-p_1)^{\lambda}.$$

6.2.2 Bahadur-Rao Theorem (Bernoulli Case)

The Bahadur-Rao theorem [1] provides the exact asymptotic probability of a large deviation, including the constant prefactor. For i.i.d. $\mathrm{Bern}(p)$ variables $X_1,\ldots,X_M$ and a threshold $\theta > p$,

$$\mathbb{P}\left(\frac{1}{M}\sum_{i=1}^{M} X_i \ge \theta\right) \sim \frac{\exp\big(-M\cdot \mathrm{KL}(\theta\|p)\big)}{\lambda^*(\theta)\,\sqrt{2\pi M\,\theta(1-\theta)}}, \qquad M\to\infty,$$

where $\lambda^*(\theta) = \log\frac{\theta(1-p)}{p(1-\theta)} > 0$ is the Cramér saddlepoint. For the lower tail $\theta < p$, the same formula holds with $|\lambda^*(\theta)|$ in the denominator.

6.2.3 Exponential Tilting

The proof of the Bahadur-Rao theorem uses exponential tilting. Define the tilted measure $\mathbb{P}_{\lambda^*}$ by

$$\frac{d\mathbb{P}_{\lambda^*}}{d\mathbb{P}} = \exp\big(\lambda^* S_M - M\psi(\lambda^*)\big), \qquad S_M = \sum_{i=1}^{M} X_i,$$

where $\psi(\lambda) = \log(1-p+pe^{\lambda})$. Under $\mathbb{P}_{\lambda^*}$, the X_i are i.i.d. $\mathrm{Bern}(\theta)$, i.e. the tilted distribution has mean exactly θ. The saddlepoint λ^* and variance $\sigma^2 = \theta(1-\theta)$ are obtained from

$$\psi'(\lambda^*) = \theta, \qquad \psi''(\lambda^*) = \theta(1-\theta).$$

6.3 Part I: Achievability (SCX Upper Bound)

6.3.1 Lemma A: Bahadur-Rao Exact Asymptotics for Bernoulli

Lemma 6.2 (Bahadur-Rao for Bernoulli, exact form). *Let $X_1, \dots, X_M \overset{i.i.d.}{\sim} \mathrm{Bern}(p)$ and fix $\theta > p$. Then*

$$\mathbb{P}\left(\frac{1}{M}\sum_{i=1}^{M} X_i \geq \theta\right) = \frac{\exp\bigl(-M \cdot \mathrm{KL}(\theta \| p)\bigr)}{\lambda^* \sqrt{2\pi M\, \theta(1-\theta)}}\bigl(1 + O(M^{-1})\bigr),$$

where $\lambda^* = \log \frac{\theta(1-p)}{p(1-\theta)} > 0$ and the $O(M^{-1})$ constant is explicit in p and θ.

Sketch. Write the binomial tail probability $\mathbb{P}(S_M \geq k) = \sum_{j=k}^{M} \binom{M}{j} p^j (1-p)^{M-j}$ with $k = \lceil M\theta \rceil$. Apply Stirling's formula with Robbins' remainder bound to the leading term a_k, then bound the ratio of consecutive terms $a_{j+1}/a_j \leq e^{-\lambda^*}(1 + O(M^{-1}))$ and sum the geometric series. The lattice correction factor $(1 - e^{-\lambda^*})^{-1}$ is absorbed into the $O(M^{-1})$ term; the form $1/\lambda^*$ is obtained by $\lambda^*/(1 - e^{-\lambda^*}) = 1 + O(\lambda^*)$. □

For the lower tail $\theta < p$, symmetry gives

$$\mathbb{P}\left(\frac{1}{M}\sum_{i=1}^{M} X_i \leq \theta\right) = \frac{\exp\bigl(-M \cdot \mathrm{KL}(\theta \| p)\bigr)}{|\lambda^*| \sqrt{2\pi M\, \theta(1-\theta)}}\bigl(1 + O(M^{-1})\bigr),$$

where $|\lambda^*| = \log \frac{p(1-\theta)}{\theta(1-p)}$.

6.3.2 Application to FPR and FNR

Apply Lemma A to the expert error indicators. Under H_0 (clean), $e_m \sim \mathrm{Bern}(p_0)$. Under H_1 (noise), $e_m \sim \mathrm{Bern}(p_1)$. For a fixed threshold $\theta \in (p_0, p_1)$,

$$\mathrm{FPR}_M(\theta) = \mathbb{P}(C_M \geq \theta \mid \text{clean}) \sim \frac{\exp\bigl(-M \cdot \mathrm{KL}(\theta \| p_0)\bigr)}{\lambda_0^* \sqrt{2\pi M\, \theta(1-\theta)}},$$

$$\mathrm{FNR}_M(\theta) = \mathbb{P}(C_M \leq \theta \mid \text{noise}) \sim \frac{\exp\bigl(-M \cdot \mathrm{KL}(\theta \| p_1)\bigr)}{|\lambda_1^*| \sqrt{2\pi M\, \theta(1-\theta)}},$$

where $\lambda_0^* = \log \frac{\theta(1-p_0)}{p_0(1-\theta)} > 0$ and $\lambda_1^* = \log \frac{\theta(1-p_1)}{p_1(1-\theta)} < 0$, $|\lambda_1^*| = \log \frac{p_1(1-\theta)}{\theta(1-p_1)}$.

6.3.3 Lemma B: F1 Asymptotic Expansion

Lemma 6.3 (F1 expansion). *Let $\eta = \mathbb{P}(noise)$, $\mathrm{FPR} = \mathrm{FPR}_M$, $\mathrm{FNR} = \mathrm{FNR}_M$, and define $r = \mathrm{FNR}/2 + (1-\eta)\,\mathrm{FPR}/(2\eta)$. For M sufficiently large that $r \leq \frac{1}{2}$,*

$$1 - \mathrm{F1} = \frac{\mathrm{FNR}}{2} + \frac{1-\eta}{2\eta}\,\mathrm{FPR} + R,$$

where $|R| \leq 2r^2$.

Proof. From the definitions,

$$\mathrm{F1} = \frac{2\eta \cdot \mathrm{TPR}}{\eta(1 + \mathrm{TPR}) + (1-\eta)\,\mathrm{FPR}} = \frac{2\eta(1 - \mathrm{FNR})}{\eta(2 - \mathrm{FNR}) + (1-\eta)\,\mathrm{FPR}}.$$

Direct algebra gives

$$1 - \mathrm{F1} = \frac{\eta\,\mathrm{FNR} + (1-\eta)\,\mathrm{FPR}}{\eta(2-\mathrm{FNR}) + (1-\eta)\,\mathrm{FPR}}.$$

Let $\alpha = \eta\,\mathrm{FNR} + (1-\eta)\,\mathrm{FPR}$ and $\beta = 2\eta - \eta\,\mathrm{FNR} + (1-\eta)\,\mathrm{FPR}$. Since $\beta \geq \eta > 0$, write $1-\mathrm{F1} = \frac{\alpha}{2\eta} \cdot \frac{1}{1-\gamma}$ with $\gamma = \frac{\mathrm{FNR}}{2} - \frac{(1-\eta)\,\mathrm{FPR}}{2\eta}$. For $|\gamma| < 1$ (which holds when $r \leq \frac{1}{2}$), expand the geometric series to obtain the first-order terms and bound the remainder. □

Substituting the Bahadur-Rao asymptotics at a threshold θ,

$$1 - \mathrm{F1}(\theta) = \frac{1}{\sqrt{2\pi M\,\theta(1-\theta)}}\left[\frac{e^{-M\cdot\mathrm{KL}(\theta\|p_1)}}{2|\lambda_1^*(\theta)|} + \frac{1-\eta}{2\eta}\,\frac{e^{-M\cdot\mathrm{KL}(\theta\|p_0)}}{\lambda_0^*(\theta)}\right] + O\!\left(\frac{e^{-2M\min(\kappa_0,\kappa_1)}}{M}\right),$$

where $\kappa_0 = \mathrm{KL}(\theta\|p_0)$ and $\kappa_1 = \mathrm{KL}(\theta\|p_1)$. The remainder decays twice as fast exponentially and is asymptotically negligible.

6.3.4 Lemma C: Chernoff Information Closed Form

Lemma 6.4 (Chernoff point and information)**.** *The unique* $\theta^* \in (p_0, p_1)$ *satisfying* $\mathrm{KL}(\theta^*\|p_0) = \mathrm{KL}(\theta^*\|p_1)$ *is*

$$\boxed{\theta^* = \frac{\log\frac{1-p_0}{1-p_1}}{\log\frac{p_1(1-p_0)}{p_0(1-p_1)}}}.$$

The common value $\kappa = \mathrm{KL}(\theta^*\|p_0)$ *is the Chernoff information* $C(\mathrm{Bern}(p_0), \mathrm{Bern}(p_1))$.

Proof. Equating the KL divergences and cancelling the common term $\theta\log\theta + (1-\theta)\log(1-\theta)$ gives

$$\theta\log\frac{1}{p_0} + (1-\theta)\log\frac{1}{1-p_0} = \theta\log\frac{1}{p_1} + (1-\theta)\log\frac{1}{1-p_1}.$$

Simplifying:

$$\theta\log\frac{p_1}{p_0} = (1-\theta)\log\frac{1-p_0}{1-p_1},$$

which yields the closed form. Strict convexity of KL in its first argument and the sign change $f(p_0) = -\mathrm{KL}(p_0\|p_1) < 0$, $f(p_1) = \mathrm{KL}(p_1\|p_0) > 0$ guarantee uniqueness and $\theta^* \in (p_0, p_1)$. □

The saddlepoints at θ^* are:

$$\lambda_0^* = \log\frac{\theta^*(1-p_0)}{p_0(1-\theta^*)} > 0, \qquad \lambda_1^* = \log\frac{\theta^*(1-p_1)}{p_1(1-\theta^*)} < 0.$$

6.3.5 Lemma D: Adaptive Threshold Optimality

The naive threshold θ^* (the Chernoff point) balances the two exponential rates but ignores the noise prior η. When $\eta \neq \frac{1}{2}$, the optimal threshold—the one minimising $1-\mathrm{F1}$—differs from θ^* by $O(1/M)$. This $O(1/M)$ shift produces an $O(1)$ multiplicative factor in the error probability via the exponent, which is essential for achieving the minimax lower bound.

Lemma 6.5 (Adaptive threshold)**.** *The threshold* $\theta^\dagger$ *that minimises the leading-order expression for* $1 - \mathrm{F1}$ *satisfies*

$$\mathrm{KL}(\theta^\dagger \| p_0) = \mathrm{KL}(\theta^\dagger \| p_1) + \frac{1}{M} \log \frac{1-\eta}{\eta} + O\left(\frac{1}{M^2}\right).$$

Expanding around θ^* *gives*

$$\theta^\dagger = \theta^* + \frac{1}{M} \frac{\log \frac{1-\eta}{\eta}}{D^*} + O\left(\frac{1}{M^2}\right),$$

where $D^* = \lambda_0^* + |\lambda_1^*| = \log \frac{p_1(1-p_0)}{p_0(1-p_1)}$.

Proof. The objective function (ignoring the common factor $1/\sqrt{2\pi M \theta(1-\theta)}$) is

$$\Phi_M(\theta) = \frac{e^{-M \cdot \mathrm{KL}(\theta \| p_1)}}{2|\lambda_1^*(\theta)|} + \frac{1-\eta}{2\eta} \frac{e^{-M \cdot \mathrm{KL}(\theta \| p_0)}}{\lambda_0^*(\theta)}.$$

Differentiating and keeping only the dominant $-M\,\mathrm{KL}'(\theta \| p) = -M\lambda^*(\theta)$ terms (the $-M$ factor dominates the $O(1)$ derivative of $\log \lambda^*$), the first-order condition reduces to

$$e^{-M \cdot \mathrm{KL}(\theta \| p_1)} \approx \frac{1-\eta}{\eta} e^{-M \cdot \mathrm{KL}(\theta \| p_0)}.$$

Taking logs yields the stated relation. Expanding $\mathrm{KL}(\theta^* + \delta \| p_0) - \mathrm{KL}(\theta^* + \delta \| p_1)$ gives $\delta D^* + O(\delta^3)$, where the quadratic term cancels exactly because $\mathrm{KL}''(\cdot \| p_0) = \mathrm{KL}''(\cdot \| p_1) = 1/(\theta(1-\theta))$ is independent of p. Solving $\delta D^* = \frac{1}{M} \log \frac{1-\eta}{\eta} + O(M^{-2})$ gives the result. $\square$

The $O(1)$ exponential prefactor. The $O(1/M)$ shift in θ produces an $O(1)$ multiplicative factor in $\exp(-M \cdot \mathrm{KL})$. Expanding the KL divergence at $\theta^\dagger = \theta^* + \delta$:

$$\mathrm{KL}(\theta^\dagger \| p_0) = \kappa + \lambda_0^* \delta + \frac{\delta^2}{2\theta^*(1-\theta^*)} + O(\delta^3),$$

$$\mathrm{KL}(\theta^\dagger \| p_1) = \kappa - |\lambda_1^*| \delta + \frac{\delta^2}{2\theta^*(1-\theta^*)} + O(\delta^3).$$

Multiplying by M and substituting δ:

$$M \cdot \mathrm{KL}(\theta^\dagger \| p_0) = M\kappa + \frac{\lambda_0^*}{D^*} \log \frac{1-\eta}{\eta} + O\left(\frac{1}{M}\right),$$

$$M \cdot \mathrm{KL}(\theta^\dagger \| p_1) = M\kappa - \frac{|\lambda_1^*|}{D^*} \log \frac{1-\eta}{\eta} + O\left(\frac{1}{M}\right).$$

Exponentiating:

$$\exp\bigl(-M \cdot \mathrm{KL}(\theta^\dagger \| p_0)\bigr) = e^{-M\kappa} \left(\frac{1-\eta}{\eta}\right)^{-\lambda_0^*/D^*} \bigl(1 + O(M^{-1})\bigr),$$

$$\exp\bigl(-M \cdot \mathrm{KL}(\theta^\dagger \| p_1)\bigr) = e^{-M\kappa} \left(\frac{1-\eta}{\eta}\right)^{|\lambda_1^*|/D^*} \bigl(1 + O(M^{-1})\bigr).$$

FPR and FNR at the adaptive threshold. Substituting into the Bahadur-Rao expansions:

$$\mathrm{FPR}_M(\theta^\dagger) \sim \frac{e^{-M\kappa}}{\lambda_0^* \sqrt{2\pi M\, \theta^*(1-\theta^*)}} \left(\frac{1-\eta}{\eta}\right)^{-\lambda_0^*/D^*},$$

$$\mathrm{FNR}_M(\theta^\dagger) \sim \frac{e^{-M\kappa}}{|\lambda_1^*| \sqrt{2\pi M\, \theta^*(1-\theta^*)}} \left(\frac{1-\eta}{\eta}\right)^{|\lambda_1^*|/D^*}.$$

The critical cancellation in $1 - \mathrm{F1}$. Recall $s = |\lambda_1^*|/D^*$ and $\lambda_0^*/D^* = 1 - s$. The FNR contribution to $1 - \mathrm{F1}$ is

$$\frac{1}{2}\,\mathrm{FNR}_M(\theta^\dagger) \sim \frac{e^{-M\kappa}}{2|\lambda_1^*| \sqrt{2\pi M\, \theta^*(1-\theta^*)}} \left(\frac{1-\eta}{\eta}\right)^{s}.$$

The FPR contribution is

$$\begin{aligned}\frac{1-\eta}{2\eta}\,\mathrm{FPR}_M(\theta^\dagger) &\sim \frac{1-\eta}{2\eta}\,\frac{e^{-M\kappa}}{\lambda_0^* \sqrt{2\pi M\, \theta^*(1-\theta^*)}} \left(\frac{1-\eta}{\eta}\right)^{-(1-s)} \\ &= \frac{e^{-M\kappa}}{2\lambda_0^* \sqrt{2\pi M\, \theta^*(1-\theta^*)}} \left(\frac{1-\eta}{\eta}\right)^{s}.\end{aligned}$$

Both terms carry the identical factor $\big((1-\eta)/\eta\big)^s$! This is the critical cancellation that makes the adaptive threshold constant-optimal.

Asymptotic constant for SCX. Summing the two contributions:

$$1 - \mathrm{F1}_{\mathrm{SCX}}(\theta^\dagger) \sim \frac{e^{-M\kappa}}{\sqrt{2\pi M\, \theta^*(1-\theta^*)}} \left(\frac{1-\eta}{\eta}\right)^{s} \frac{1}{2}\left(\frac{1}{\lambda_0^*} + \frac{1}{|\lambda_1^*|}\right).$$

Multiplying by $e^{M\kappa}\sqrt{2\pi M}$ and taking the limit:

$$\boxed{\lim_{M\to\infty} e^{M\kappa}\sqrt{2\pi M}\,\big(1 - \mathrm{F1}_{\mathrm{SCX}}(\theta^\dagger)\big) = \frac{1}{2}\left(\frac{1-\eta}{\eta}\right)^{s} \frac{1/\lambda_0^* + 1/|\lambda_1^*|}{\sqrt{\theta^*(1-\theta^*)}}}.$$

This equals $C_{\min}/\eta$ because

$$\frac{C_{\min}}{\eta} = \frac{1}{2}\left(\frac{1-\eta}{\eta}\right)^{s} \frac{1/\lambda_0^* + 1/|\lambda_1^*|}{\sqrt{\theta^*(1-\theta^*)}}.$$

Thus Lemma D establishes part (a) of Theorem 4'.

6.4 Part II: Lower Bound (Minimax Converse)

6.4.1 Lemma E: Second-Order Asymptotic Lower Bound

Lemma 6.6 (Exact constant minimax lower bound)**.** *For any noise detection algorithm $\mathcal{A}$,*

$$\liminf_{M\to\infty} e^{M\kappa}\sqrt{2\pi M}\,\big(1 - \mathrm{F1}_{\mathcal{A}}(M)\big) \geq \frac{C_{\min}}{\eta},$$

where $C_{\min}$ is the constant defined in Theorem 4'. Equivalently, writing $r = \lambda_0^*/D^* = 1-s$,

$$C_{\min} = \frac{\eta^r(1-\eta)^{1-r}}{2}\,\frac{1/\lambda_0^* + 1/|\lambda_1^*|}{\sqrt{\theta^*(1-\theta^*)}}.$$

Proof. The proof proceeds in five parts.

Part 1: Reduction to a binary hypothesis test. Within a fixed state s, the expert errors $e_1, \ldots, e_M$ are i.i.d. $\mathrm{Bern}(p_0)$ under H_0 (clean) and $\mathrm{Bern}(p_1)$ under H_1 (noise). Any algorithm $\mathcal{A}$ reduces to a decision rule $\delta_M : \{0,1\}^M \to \{0,1\}$ with Type I error $\alpha_M = \Pr(\delta_M = 1 \mid \mathrm{H}_0)$ and Type II error $\beta_M = \Pr(\delta_M = 0 \mid \mathrm{H}_1)$.

From Lemma B,

$$1 - \mathrm{F1}_{\mathcal{A}} = \frac{1-\eta}{2\eta}\,\alpha_M + \frac{1}{2}\,\beta_M + o(1),$$

where the $o(1)$ term is uniform over decision rules and decays as $O(e^{-2M\kappa})$, negligible at the leading order. Define the weighted risk

$$R_M(\delta_M) = w_0\alpha_M + w_1\beta_M, \qquad w_0 = \frac{1-\eta}{2\eta},\ w_1 = \frac{1}{2}.$$

Part 2: Neyman-Pearson reduction. By the Neyman-Pearson lemma, for any test with Type I error $\leq \alpha$, the minimum achievable Type II error is attained by the likelihood ratio test (LRT). Therefore for any decision rule δ_M,

$$R_M(\delta_M) \geq \min_{\alpha\in[0,1]} \bigl\{w_0\alpha + w_1\beta_{\mathrm{NP}}(\alpha)\bigr\},$$

where $\beta_{\mathrm{NP}}(\alpha)$ is the Type II error of the most powerful level-α test.

The minimiser α^* satisfies $w_0 + w_1\beta'_{\mathrm{NP}}(\alpha^*) = 0$. For the Neyman-Pearson test, the slope of the ROC curve at α^* is $-1/\tau^*$, where τ^* is the likelihood ratio threshold. Hence $\tau^* = w_0/w_1 = (1-\eta)/\eta$; the minimiser is exactly the Bayes decision rule with prior odds $(1-\eta) : \eta$ against H_0. Thus for any $\mathcal{A}$,

$$R_M(\delta_{\mathcal{A}}) \geq R_M^* := w_0\alpha_M^* + w_1\beta_M^*,$$

where $\delta_M^*(e) = \mathbf{1}\{\log L_M(e) > \log\frac{1-\eta}{\eta}\}$ and $L_M(e) = \prod_{m=1}^M \frac{p_1^{e_m}(1-p_1)^{1-e_m}}{p_0^{e_m}(1-p_0)^{1-e_m}}$.

Part 3: Exact asymptotics of the Bayes test. The log-likelihood ratio simplifies to

$$\log L_M = M\bigl(C_M\Delta + \Delta_0\bigr),$$

where $\Delta = \log\frac{p_1(1-p_0)}{p_0(1-p_1)} = D^*$ and $\Delta_0 = \log\frac{1-p_1}{1-p_0}$. The rejection region $C_M > -\Delta_0/\Delta + \frac{1}{M\Delta}\log\frac{1-\eta}{\eta}$ has threshold

$$t_M = \theta^* + \frac{1}{MD^*}\log\frac{1-\eta}{\eta}.$$

This is precisely the adaptive threshold $\theta^\dagger$ from Lemma D.

The error probabilities of the Bayes test are

$$\alpha_M^* = \Pr_{p_0}(C_M > t_M), \qquad \beta_M^* = \Pr_{p_1}(C_M \leq t_M),$$

up to a $O(e^{-M\kappa}/M)$ correction for tie-breaking.

Apply the Bahadur-Rao theorem (Lemma A) to each:

$$\alpha_M^* \sim \frac{\exp\bigl(-M\,\mathrm{KL}(t_M\|p_0)\bigr)}{\lambda_0^*(t_M)\sqrt{2\pi M\,t_M(1-t_M)}},$$

$$\beta_M^* \sim \frac{\exp\bigl(-M\,\mathrm{KL}(t_M\|p_1)\bigr)}{|\lambda_1^*(t_M)|\sqrt{2\pi M\,t_M(1-t_M)}}.$$

Part 4: Expansion around the Chernoff point. Write $\delta = t_M - \theta^* = \frac{1}{MD^*} \log \frac{1-\eta}{\eta} + O(M^{-2})$. Expand the KL divergences as in Lemma D:

$$M \, \mathrm{KL}(t_M \| p_0) = M\kappa + \frac{\lambda_0^*}{D^*} \log \frac{1-\eta}{\eta} + o(1),$$

$$M \, \mathrm{KL}(t_M \| p_1) = M\kappa - \frac{|\lambda_1^*|}{D^*} \log \frac{1-\eta}{\eta} + o(1).$$

Substituting into the Bahadur-Rao expressions:

$$\alpha_M^* = \frac{e^{-M\kappa}}{\lambda_0^* \sqrt{2\pi M \, \theta^*(1-\theta^*)}} \left(\frac{1-\eta}{\eta} \right)^{-\lambda_0^*/D^*} \big(1 + o(1)\big),$$

$$\beta_M^* = \frac{e^{-M\kappa}}{|\lambda_1^*| \sqrt{2\pi M \, \theta^*(1-\theta^*)}} \left(\frac{1-\eta}{\eta} \right)^{|\lambda_1^*|/D^*} \big(1 + o(1)\big).$$

The weighted risk is therefore

$$\begin{aligned} R_M^* &= w_0 \alpha_M^* + w_1 \beta_M^* \\ &= \frac{e^{-M\kappa}}{\sqrt{2\pi M \, \theta^*(1-\theta^*)}} \left[\frac{1-\eta}{2\eta} \frac{1}{\lambda_0^*} \left(\frac{1-\eta}{\eta} \right)^{-\lambda_0^*/D^*} + \frac{1}{2} \frac{1}{|\lambda_1^*|} \left(\frac{1-\eta}{\eta} \right)^{|\lambda_1^*|/D^*} \right] \big(1 + o(1)\big). \end{aligned}$$

Using $s = |\lambda_1^*|/D^*$ and $\lambda_0^*/D^* = 1 - s$, the terms simplify:

$$\frac{1-\eta}{2\eta\lambda_0^*} \left(\frac{1-\eta}{\eta} \right)^{-(1-s)} = \frac{1}{2\lambda_0^*} \left(\frac{1-\eta}{\eta} \right)^{s},$$

$$\frac{1}{2|\lambda_1^*|} \left(\frac{1-\eta}{\eta} \right)^{s} = \frac{1}{2|\lambda_1^*|} \left(\frac{1-\eta}{\eta} \right)^{s}.$$

Both terms share the common factor $\frac{1}{2}\big((1-\eta)/\eta\big)^s$. Factoring it out:

$$R_M^* = \frac{e^{-M\kappa}}{\sqrt{2\pi M \, \theta^*(1-\theta^*)}} \frac{1}{2} \left(\frac{1-\eta}{\eta} \right)^{s} \left(\frac{1}{\lambda_0^*} + \frac{1}{|\lambda_1^*|} \right) \big(1 + o(1)\big).$$

Part 5: Conversion to the F1 lower bound. From Part 1, $1 - \mathrm{F1}_{\mathcal{A}} = R_M(\delta_{\mathcal{A}}) + o(1) \geq R_M^* + o(1)$. Multiplying by $e^{M\kappa}\sqrt{2\pi M}$ and taking the limit inferior:

$$\liminf_{M \to \infty} e^{M\kappa} \sqrt{2\pi M} \, \big(1 - \mathrm{F1}_{\mathcal{A}}(M)\big) \geq \frac{1}{2} \left(\frac{1-\eta}{\eta} \right)^{s} \frac{1/\lambda_0^* + 1/|\lambda_1^*|}{\sqrt{\theta^*(1-\theta^*)}}.$$

The right-hand side equals $C_{\min}/\eta$ by definition of $C_{\min}$. This completes the proof of Lemma E. □ □

6.5 Part III: Constant Matching → Optimality

6.5.1 The $O(1)$ Prefactor Cancellation

The central mechanism of the constant optimality proof is the cancellation that produces the identical factor $((1-\eta)/\eta)^s$ in both the FPR and FNR contributions to $1 - \mathrm{F1}$.

The adaptive threshold $\theta^\dagger$ differs from the Chernoff point θ^* by $O(1/M)$. This $O(1/M)$ shift generates an $O(1)$ change in the exponent:

$$\exp\bigl(-M \cdot \mathrm{KL}(\theta^\dagger \| p)\bigr) = e^{-M\kappa} \cdot \text{constant factor}.$$

The constant factor for $p = p_0$ is $((1-\eta)/\eta)^{-\lambda_0^*/D^*}$ and for $p = p_1$ is $((1-\eta)/\eta)^{|\lambda_1^*|/D^*}$. When combined with the F1 weights $w_0 = (1-\eta)/(2\eta)$ and $w_1 = 1/2$, both contributions acquire the identical factor $((1-\eta)/\eta)^s$ because:

$$\frac{1-\eta}{2\eta} \cdot \left(\frac{1-\eta}{\eta}\right)^{-\lambda_0^*/D^*} = \frac{1}{2\lambda_0^*}\left(\frac{1-\eta}{\eta}\right)^s, \qquad \frac{1}{2} \cdot \left(\frac{1-\eta}{\eta}\right)^{|\lambda_1^*|/D^*} = \frac{1}{2|\lambda_1^*|}\left(\frac{1-\eta}{\eta}\right)^s.$$

This cancellation is mathematically unavoidable—it follows from the relationship $s + (1-s) = 1$, i.e., $|\lambda_1^*|/D^* + \lambda_0^*/D^* = 1$, which is a consequence of the definition $D^* = \lambda_0^* + |\lambda_1^*|$.

6.5.2 Proof that SCX Achieves $C_{\min}$

From Lemma D, the SCX detector with adaptive threshold $\theta^\dagger$ attains:

$$\lim_{M\to\infty} e^{M\kappa}\sqrt{2\pi M}\,\bigl(1 - \mathrm{F1}_{\mathrm{SCX}}(\theta^\dagger)\bigr) = \frac{1}{2}\left(\frac{1-\eta}{\eta}\right)^s \frac{1/\lambda_0^* + 1/|\lambda_1^*|}{\sqrt{\theta^*(1-\theta^*)}}.$$

From Lemma E, the minimax lower bound for any algorithm is:

$$\liminf_{M\to\infty} e^{M\kappa}\sqrt{2\pi M}\,\bigl(1 - \mathrm{F1}_{\mathcal{A}}\bigr) \geq \frac{1}{2}\left(\frac{1-\eta}{\eta}\right)^s \frac{1/\lambda_0^* + 1/|\lambda_1^*|}{\sqrt{\theta^*(1-\theta^*)}}.$$

The two expressions are **identical**. Therefore the SCX detector attains the lower bound, establishing exact constant minimax optimality. This proves parts (a)–(c) of Theorem 4'.

6.6 Part IV: Multi-State Aggregation (Lemma F)

In practice the data contains multiple states, each with its own $p_{0,s}$, $p_{1,s}$, Chernoff information κ_s, and per-state constant C_s. The global F1 score is the expectation over the state mixture.

Lemma 6.7 (Multi-state aggregation)**.** *Let there be S states with proportions $\rho_s > 0$, $\sum_s \rho_s = 1$. Within state s, the per-state* $\mathrm{F1}_s(M)$ *satisfies*

$$1 - \mathrm{F1}_s(M) \sim \frac{C_s}{\eta}\,\frac{e^{-M\kappa_s}}{\sqrt{2\pi M}}.$$

Then:

(i) ***Additivity.*** $1 - \mathrm{F1}_{global}(M) = \sum_{s=1}^{S} \rho_s\,(1 - \mathrm{F1}_s(M))$.

(ii) ***Bottleneck rate.*** $\kappa_{global} = \min_{1\leq s\leq S} \kappa_s$.

(iii) ***Dominant-state constant.*** *Let* $\mathcal{S}_{\min} = \{s : \kappa_s = \kappa_{global}\}$. *Then*

$$\liminf_{M\to\infty} e^{M\kappa_{global}}\sqrt{2\pi M}\,\big(1 - \mathrm{F1}_{\mathcal{A}}(M)\big) \geq \frac{1}{\eta}\sum_{s\in\mathcal{S}_{\min}} \rho_s C_s.$$

States with $\kappa_s > \kappa_{global}$ *contribute only* $o(e^{-M\kappa_{global}}/\sqrt{M})$ *and are asymptotically negligible.*

Proof. Part (i) follows from linearity of expectation. For part (ii), each state contributes $e^{-M\kappa_s}/\sqrt{2\pi M}$. The slowest exponential decay (smallest κ_s) dominates as $M \to \infty$. Part (iii) follows from Lemma E applied per state and summing with weights ρ_s; states with $\kappa_s > \kappa_{\text{global}}$ are suppressed by $e^{-M(\kappa_s-\kappa_{\text{global}})} \to 0$. □

6.7 Numerical Verification

We verify the theoretical constants numerically across five parameter sets spanning a range of noise rates and distribution separations.

6.7.1 Five Parameter Sets

The following table reports for each (p_0, p_1, η) combination: the Chernoff information κ, the ratio $2(p_1-p_0)^2/\kappa$ (showing how much looser the Hoeffding bound is), the optimal constant $C_{\min}$, the normalised constant $C_{\min}/\eta$, the suboptimality ratio of the non-adaptive threshold θ^*, and whether the adaptive threshold matches the lower bound.

Case	p_0	p_1	η	κ	$\frac{2(p_1-p_0)^2}{\kappa}$	$C_{\min}$	$\frac{C_{\min}}{\eta}$	$\frac{\text{Non-ad.}}{C_{\min}/\eta}$	Adapt match?
1	0.10	0.60	0.10	0.1696	2.95	0.4591	4.5915	1.70	YES
2	0.20	0.50	0.30	0.0528	3.41	1.3768	4.5894	1.09	YES
3	0.05	0.80	0.05	0.4574	2.46	0.1843	3.6864	2.41	YES
4	0.10	0.60	0.50	0.1696	2.95	0.8348	1.6697	1.00	YES
5	0.10	0.60	0.90	0.1696	2.95	0.5465	0.6072	1.62	YES

6.7.2 Key Findings

1. **Adaptive limit matches $C_{\min}/\eta$ exactly.** For all five cases, the difference between the adaptive threshold constant and the lower bound constant is below 10^{-15} (machine precision), confirming Theorem 4'(c).

2. **Non-adaptive threshold is suboptimal.** Using the naive Chernoff point θ^* (which ignores η) yields constants 1.00–2.41× larger than $C_{\min}/\eta$. The worst ratio (2.41) occurs for rare noise ($\eta = 0.05$) with widely separated distributions.

3. **Chernoff rate is slower than Hoeffding bound.** The ratio $2\Delta^2/\kappa$ ranges from 2.46 to 3.41, confirming that $\kappa < 2\Delta^2$ for all tested parameters. The Hoeffding exponential bound is looser by a factor of 2.5–3.4, justifying the use of the exact Chernoff information.

4. **Symmetric case $\eta = 0.5$.** When $\eta = 0.5$, $\theta^\dagger = \theta^*$ and the adaptive and non-adaptive thresholds coincide (ratio 1.00). This is expected because $((1-\eta)/\eta)^s = 1$ for $\eta = 0.5$.

6.7.3 Finite-M Convergence

The asymptotic limits are derived as $M \to \infty$. For finite M, the Bahadur-Rao theorem provides an $O(1/M)$ correction term. Lemma B gives explicit bounds:

$$1 - \mathrm{F1}_{\mathrm{SCX}}(\theta^\dagger) \le \frac{C_{\min}}{\eta} \frac{e^{-M\kappa}}{\sqrt{2\pi M}} \big(1 + K_1 M^{-1/2}\big),$$

$$1 - \mathrm{F1}_{\mathrm{SCX}}(\theta^\dagger) \ge \frac{C_{\min}}{\eta} \frac{e^{-M\kappa}}{\sqrt{2\pi M}} \big(1 - K_2 M^{-1/2}\big),$$

where K_1, K_2 depend only on p_0, p_1. Numerical experiments (Supplementary Material S8) confirm that the finite-M F1 score approaches the asymptotic limit at the predicted $O(1/\sqrt{M})$ rate.

References for Section S4

References

[1] Bahadur, R. R. & Rao, R. R. On deviations of the sample mean. *Annals of Mathematical Statistics* **31**(4), 1015–1027 (1960).

7 Theorem 5: Cluster Consistency of State Discovery

This section proves Theorem 5, which establishes that under the **strong features** regime (large separation $\Delta_{\min}$ relative to noise), Lloyd's k-means with $O(\log n)$ random restarts recovers the true state partition with probability tending to one exponentially fast. Theorem 5 is the positive counterpart to the negative result of Theorem 2 (weak features $\to$ failure) and Theorem 4' (minimax lower bound), and it complements the F1 bound from Theorem 1 (**??**) and the weak-feature F1 bound from Theorem 2 (**??**). Together these results complete the phase diagram: strong features imply state discovery succeeds; weak features imply SCX cannot exceed the loss baseline; and the threshold is minimax optimal as established in Theorem 4' (**??**).

7.1 Setup: Fixed-K Generative Model with Strong Features

We formalize the data-generating process and the clustering objective.

Assumption 7.1 (Generative model)**.** Let $(\Omega, ,)$ be a probability space. We observe n i.i.d. copies $\phi(x_i) \in^{d_\phi}$ of the feature vector, generated as

$$\phi(x) = \mu_{s(x)} + \varepsilon, \tag{44}$$

where

- $s(x) \in \{1, \dots, K\}$ is the unobserved true state,
- K is **fixed** (does not grow with n),
- $\in^{d_\phi}$ is the feature mean (center) for state k,

- $\varepsilon \in^{d_\phi}$ is a zero-mean sub-Gaussian random vector with variance proxy $\sigma^2 > 0$, independent of $s(x)$:

$$[\varepsilon] = 0, \qquad \big[\exp(t \cdot u^\top \varepsilon)\big] \le \exp(\sigma^2 t^2/2), \quad \forall u \in \mathbb{S}^{d_\phi - 1},\ \forall t \in .$$

Assumption 7.2 (Cluster structure)**.** The true partition is $= \{C_1^*, \dots, C_K^*\}$ with $C_k^* = \{x \in \colon s(x) = k\}$. The true centers are $= [\phi(X) \mid S = k]$. The population proportions are $\pi_k = P(S = k) > 0$ for all k, and we denote $\pi_{\min} = \min_k \pi_k$. The empirical per-state sample sizes are $n_k = \sum_{i=1}^n \{s(x_i) = k\}$ with $n_{\min} = \min_k n_k$.

Assumption 7.3 (Strong separation (fixed $\Delta_{\min}$))**.** The minimum distance between distinct cluster centers is a fixed positive constant (it does *not* scale with n or K):

$$\Delta_{\min} \;=\; \min_{i \neq j} \mu_i - \mu_{j2} \;>\; 0.$$

We assume the **strong separation regime**

$$\frac{\Delta_{\min}^2}{\sigma^2 d_\phi} \;\ge\; C_0, \tag{45}$$

where C_0 is a universal constant large enough to guarantee $\theta^* - \mu < \Delta_{\min}/8$ in Lemma 7.6 below. This is the "strong features" regime: the signal-to-noise ratio is sufficiently high that cluster centers are distinguishable despite noise and feature dimensionality.

Definition 7.4 (k-means objective)**.** For a set of K candidate centers $\theta = \{\theta_1, \dots, \theta_K\} \subset^{d_\phi}$, define the **empirical k-means risk**

$$W_n(\theta) = \frac{1}{n} \sum_{i=1}^n \min_{k \in [K]} \phi(x_i) - \theta_{k2}^2,$$

and the **population k-means risk**

$$W(\theta) = \Big[\min_{k \in [K]} \phi(X) - \theta_{k2}^2\Big].$$

Let $\theta^* \in_{\theta : |\theta| = K} W(\theta)$ be the population minimizer and $_n \in_{\theta : |\theta| = K} W_n(\theta)$ the global empirical minimizer, both defined up to label permutation. When we write $_n - \theta^*$, we mean the distance after optimal permutation matching.

Definition 7.5 (Norms on center sets)**.** For a K-center set θ, let

$$\theta_\infty = \max_{k \in [K]} \theta_{k2}, \qquad \theta - \theta' = \max_k \theta_k - \theta'_{k2}.$$

7.2 Lemma 1: Population Minimizer Proximity

Lemma 7.6 (Population minimizer proximity)**.** *Under Assumptions 7.1–7.3, the population k-means objective $W(\theta)$ has a minimizer θ^* (unique up to label permutation). For any permutation π that optimally matches θ^* to the true centers μ,*

$$\theta^*_{\pi(k)} -_2 \;\le\; \epsilon_{\text{pop}} \;<\; \frac{\Delta_{\min}}{8},$$

where

$$\epsilon_{\text{pop}} \;=\; \frac{C_2\, K\, (M + \sigma\sqrt{d_\phi})}{\pi_{\min}} \cdot \exp\!\left(-\frac{\Delta_{\min}^2}{8\sigma^2}\right),$$

and C_2 is an absolute constant, $M = \max_k{}_2$. Under (45) with C_0 sufficiently large, $\epsilon_{\text{pop}} < \Delta_{\min}/8$.

Proof of Lemma 7.6. The proof proceeds in five steps.
Step 1 (Existence). $W(\theta)$ is continuous (composition of continuous functions) and coercive ($W(\theta) \to \infty$ as $\theta_\infty \to \infty$, since the quadratic penalty dominates), hence attains a minimum on any compact set containing a sufficiently large ball. By strict convexity of $\theta_k \mapsto [\phi - \theta_k{}^2 \mid \phi \in V_k(\theta)]$ on each Voronoi cell (defined below), any two minimizers must be related by a label permutation. Thus θ^* is unique up to permutation.
Step 2 (Self-consistency equation). For a set of centers θ, define the Voronoi cell of center j:

$$V_j(\theta) = \big\{\phi \in^{d_\phi} : \phi - \theta_{j2} \le \min_{\ell \neq j} \phi - \theta_{\ell 2}\big\}.$$

At points where Voronoi cells have positive probability and are well-defined, the gradient of W with respect to θ_j exists:

$$\frac{\partial}{\partial \theta_j} W(\theta) = -2\left[(\phi - \theta_j) \cdot \{\phi \in V_j(\theta)\}\right].$$

Setting this to zero at the minimizer θ^* gives the **self-consistency condition**: for each j,

$$\theta_j^* = \frac{[\phi \cdot \{\phi \in V_j(\theta^*)\}]}{P(\phi \in V_j(\theta^*))} = [\phi \mid \phi \in V_j(\theta^*)]. \tag{1}$$

Step 3 (True centers approximately satisfy self-consistency). Consider the Voronoi cells induced by the true centers $\mu =$. For each k, define the one-step candidate

$$_k = [\phi \mid \phi \in V_k(\mu)].$$

We bound $_k -$.
Claim (mis-assignment probability bound). For $j \neq k$,

$$P\big(\mu_j + \varepsilon \in V_k(\mu) \mid S = j\big) \;\le\; \exp\!\left(-\frac{\mu_j -^2}{8\sigma^2}\right) \;\le\; \exp\!\left(-\frac{\Delta_{\min}^2}{8\sigma^2}\right).$$

Proof of claim. The event $\mu_j + \varepsilon \in V_k(\mu)$ requires $\mu_j + \varepsilon - \le \mu_j + \varepsilon - \mu_j$, i.e. $(\mu_j -) + \varepsilon \le \varepsilon$. Squaring:

$$\mu_j -^2 + 2\varepsilon^\top(\mu_j -) + \varepsilon^2 \;\le\; \varepsilon^2,$$

hence $2\varepsilon^\top(\mu_j -) \le -\mu_j -^2$. Since $v = (\mu_j -)/\mu_j -$ is a unit vector, $\varepsilon^\top v$ is sub-Gaussian with parameter σ^2, so $2\varepsilon^\top(\mu_j -)$ is sub-Gaussian with parameter $4\sigma^2 \mu_j -^2$. By the sub-Gaussian tail bound,

$$P\big(2\varepsilon^\top(\mu_j -) \le -\mu_j -^2\big) \le \exp\!\left(-\frac{\mu_j -^4}{2 \cdot 4\sigma^2 \mu_j -^2}\right) = \exp\!\left(-\frac{\mu_j -^2}{8\sigma^2}\right),$$

establishing the claim. □

Similarly, points from state k are *not* assigned to $V_k(\mu)$ only if they are closer to some μ_j, $j \neq k$. By the union bound over $j \neq k$,

$$P\big(\mu_k + \varepsilon \notin V_k(\mu) \mid S = k\big) \;\le\; K \cdot \exp\!\left(-\frac{\Delta_{\min}^2}{8\sigma^2}\right).$$

Step 4 (Bounding $_k -$). Decompose the conditional expectation:

$$_k = [\phi \mid \phi \in V_k(\mu)] = \frac{[\phi \cdot \{\phi \in V_k(\mu)\}]}{P(V_k(\mu))}.$$

Write $\phi = \mu_S + \varepsilon$ and split by true state S:

$$[\phi \cdot \{\phi \in V_k(\mu)\}] = \sum_{j=1}^{K} \pi_j \left[(\mu_j + \varepsilon) \cdot \{\mu_j + \varepsilon \in V_k(\mu)\} \mid S = j\right].$$

For the diagonal term ($j = k$),

$$\begin{aligned}[(\mu_k + \varepsilon) \cdot \{\mu_k + \varepsilon \in V_k(\mu)\} \mid S = k] = \cdot P(\mu_k + \varepsilon \in V_k(\mu) \mid S = k)\\ + [\varepsilon \cdot \{\mu_k + \varepsilon \in V_k(\mu)\} \mid S = k].\end{aligned}$$

For the off-diagonal terms ($j \neq k$), the indicator event has probability $\leq \exp(-\Delta_{\min}^2/(8\sigma^2))$.

The conditional expectation of noise satisfies

$$[\varepsilon \mid \mu_k + \varepsilon \in V_k(\mu)]_2 \leq C_3\, \sigma \sqrt{d_\phi} \cdot \exp\left(-\frac{\Delta_{\min}^2}{8\sigma^2}\right)$$

for an absolute constant $C_3 > 0$, because $V_k(\mu)$ differs from the full space by an event of exponentially small probability, and the unconditional mean of ε is zero.

Assembling terms,

$$_k - \leq \frac{C_3\, K\, (M + \sigma\sqrt{d_\phi})}{\pi_k} \cdot \exp\left(-\frac{\Delta_{\min}^2}{8\sigma^2}\right) \leq \epsilon_{\text{pop}}.$$

Step 5 (From approximate fixed point to minimizer proximity). Let T be the self-consistency operator: $T(\theta)_j = [\phi \mid \phi \in V_j(\theta)]$. The minimizer θ^* satisfies $T(\theta^*) = \theta^*$ by Step 2, and the true centers μ satisfy $T(\mu) - \mu_\infty \leq \epsilon_{\text{pop}}$ by Step 4.

Under (45), the Voronoi partition is stable: if two center sets θ, θ' satisfy $\theta - \theta' < \Delta_{\min}/4$, their Voronoi partitions differ only on a set of exponentially small measure. Consequently, T is a contraction in a neighborhood of μ:

$$T(\theta) - T(\mu)_\infty \leq \frac{1}{2}\theta - \mu_\infty, \qquad \text{whenever } \theta - \mu_\infty \leq \frac{\Delta_{\min}}{4}.$$

Applying the contraction bound,

$$\theta^* - \mu_\infty \leq T(\theta^*) - T(\mu)_\infty + T(\mu) - \mu_\infty \leq \frac{1}{2}\theta^* - \mu_\infty + \epsilon_{\text{pop}},$$

so $\theta^* - \mu_\infty \leq 2\epsilon_{\text{pop}}$.

Under (45) with C_0 sufficiently large, $2\epsilon_{\text{pop}} < \Delta_{\min}/8$, because the exponential $\exp(-\Delta_{\min}^2/(8\sigma^2))$ decays super-polynomially in $\Delta_{\min}^2/\sigma^2$ while the pre-factor is polynomial. The explicit condition is

$$\frac{\Delta_{\min}^2}{8\sigma^2} \;\geq\; \log\left(\frac{16 C_2\, K\, (M + \sigma\sqrt{d_\phi})}{\pi_{\min}\Delta_{\min}}\right),$$

which is guaranteed by C_0 large enough since the right-hand side grows only logarithmically in the problem parameters. $\square$ $\square$

Remark 7.7. The proof uses the self-consistency equation instead of directly computing $W(\mu) - W(\mu^*)$, thereby avoiding the reversed-inequality issue. The bias ϵ_{pop} is exponentially small in $\Delta_{\min}^2/\sigma^2$ and is much smaller than the $\Delta_{\min}/8$ threshold needed for subsequent lemmas.

7.3 Lemma 2: Exponential Convergence of the Empirical Minimizer

Lemma 7.8 (Exponential convergence of empirical minimizer)**.** *Let θ^* be the unique population minimizer from Lemma 7.6 and let $_n$ be the global empirical minimizer of W_n. For any $t > 0$ satisfying $t \geq C_0\sqrt{Kd_\phi/n}$, there exist constants $c_2, C_4 > 0$ depending on the problem parameters such that*

$$P\big(_n - \theta^* \geq t\big) \;\leq\; C_4 \cdot \exp\!\left(-c_2 \cdot \frac{n\,t^2}{\sigma^2 d_\phi}\right). \tag{46}$$

The constants are $c_2 = \lambda^2/(128C_L^2)$ with $\lambda = \pi_{\min}/2$ (the strong convexity parameter) and $C_L = 2(M + \sigma\sqrt{d_\phi})$ (the Lipschitz constant bound), and $C_4 = 2$ (from the peeling sum). In particular, for $t = \Delta_{\min}/8$ and $n \geq n_0$,

$$P\left(_n - \theta^* \geq \frac{\Delta_{\min}}{8}\right) \;\leq\; 2 \cdot \exp\!\left(-c_2 \cdot \frac{n\,\Delta_{\min}^2}{64\,\sigma^2 d_\phi}\right).$$

Proof of Lemma 7.8. The proof uses a localized argmin argument with peeling. It has three parts: (i) a quadratic lower bound for W near θ^*, (ii and iii) control of the empirical process via covering numbers and concentration, and (iv) a peeling argument to convert the empirical process bound into a bound on $_n - \theta^*$.

Step 1 (Quadratic lower bound near θ^*). Define $r_0 = \Delta_{\min}/4$. By Lemma 7.6, $\theta^* - \mu < \Delta_{\min}/8$, so the set $\{\theta : \theta - \theta^* \leq r_0\}$ is contained within $\Delta_{\min}/4$ of the true centers. Within this region, for any θ, the Voronoi cells $V_j(\theta)$ are close to $V_j(\theta^*)$, differing only on an exponentially small set. By Lemma S5.1 (quadratic lower bound, proved in the Appendix),

$$W(\theta) - W(\theta^*) \;\geq\; \lambda \cdot \theta - \theta^{*2}, \qquad \forall\, \theta - \theta^* \leq r_0, \tag{2}$$

where $\lambda = \pi_{\min}/2$.

Step 2 (The argmin inequality). Let $_n$ minimize W_n. From $W_n(_n) \leq W_n(\theta^*)$,

$$\begin{aligned}
0 &\geq W_n(_n) - W_n(\theta^*) \\
&= \big(W(_n) - W(\theta^*)\big) + \big((P_n - P)(f_{_n} - f_{\theta^*})\big) \\
&\geq \lambda_n - \theta^{*2} - \big|(P_n - P)(f_{_n} - f_{\theta^*})\big|,
\end{aligned}$$

where $f_\theta(x) = \min_k \phi(x) - \theta_k{}^2$. Therefore,

$$\lambda_n - \theta^{*2} \;\leq\; \big|(P_n - P)(f_{_n} - f_{\theta^*})\big|. \tag{3}$$

This holds whenever $_n - \theta^* \leq r_0$; the case $_n - \theta^* > r_0$ is handled separately.

Step 3 (The localized empirical process). Define the localized supremum

$$\psi_n(r) = \sup_{\theta - \theta^* \leq r} \big|(P_n - P)(f_\theta - f_{\theta^*})\big|.$$

From (3), if $_n - \theta^* \leq r_0$, then letting $r = _n - \theta^*$ gives

$$\lambda r^2 \;\leq\; \psi_n(r). \tag{4}$$

Step 4 (Expectation of $\psi_n(r)$). For a fixed radius r, consider the function class

$$\mathcal{G}_r = \{\, g_\theta(x) = f_\theta(x) - f_{\theta^*}(x) : \theta - \theta^* \leq r \,\}.$$

Each g_θ satisfies $|g_\theta(x)| \le L(x)\, r$ where $L(x) = 2(\phi(x)_2 + M)$ (by the Lipschitz property, Lemma S5.2), and $[g_\theta(X)^2] \le C_L^2 r^2$ where $C_L^2 = 8(\max_j \mu_j{}^2 + \sigma^2 d_\phi + M^2)$.

The covering number of $\Theta_r = \{\theta : \theta - \theta^* \le r\}$ under $\cdot_\infty$ is bounded (Lemma S5.3):

$$\log \mathcal{N}(\epsilon, \Theta_r, \cdot_\infty) \;\le\; K d_\phi \log\left(1 + \frac{4r}{\epsilon}\right).$$

By Dudley's entropy integral, the expected supremum satisfies

$$[\psi_n(r)] \;\le\; C_5 \cdot \sqrt{\frac{K d_\phi}{n}} \cdot C_L \cdot r, \tag{5}$$

where C_5 is a universal constant.

Step 5 (Concentration of $\psi_n(r)$). For the function class $\mathcal{G}_r$ with $|g| \le C_L r$ and $[g^2] \le C_L^2 r^2$, Talagrand's concentration inequality gives

$$P\big(\psi_n(r) \ge [\psi_n(r)] + u\big) \;\le\; \exp\left(-\frac{c_6\, n\, u^2}{C_L^2\, r^2}\right), \tag{6}$$

for any $u \ge 0$, where c_6 is a universal constant.

Step 6 (Peeling argument). Fix t such that $t \le r_0$ and $t \ge \frac{8C_5C_L}{\lambda}\sqrt{\frac{Kd_\phi}{n}}$. Let $J = \lceil \log_2(r_0/t) \rceil$ and define dyadic intervals $r_j = 2^j t$ for $j = 0, 1, \ldots, J$.

From (4): if $_n - \theta^* \ge t$, then $\lambda_n - \theta^{*2} \le \psi_n(_n - \theta^*)$. For $_n - \theta^*$ falling in $[r_{j-1}, r_j)$ for some j,

$$\lambda r_{j-1}^2 \;\le\; \lambda_n - \theta^{*2} \;\le\; \psi_n(_n - \theta^*) \;\le\; \psi_n(r_j).$$

Therefore,

$$P(_n - \theta^* \ge t) \;\le\; \sum_{j=0}^{J} P\big(\psi_n(r_j) \ge \lambda (r_{j-1})^2\big), \tag{7}$$

where $r_{-1} := t/2$ for the $j = 0$ term.

For each j, bound $P(\psi_n(r_j) \ge \lambda r_{j-1}^2)$. Since $r_j = 2r_{j-1}$, $\lambda r_{j-1}^2 = \lambda r_j^2/4$. From (5): $[\psi_n(r_j)] \le C_5 C_L \sqrt{K d_\phi/n}\, r_j$. Set $u_j = \lambda r_j^2/4 - [\psi_n(r_j)]$. Our condition on t ensures

$$[\psi_n(r_j)] \le C_5 C_L \sqrt{K d_\phi/n}\, r_j \le \frac{\lambda r_j^2}{8},$$

since $r_j \ge t \ge \frac{8C_5C_L}{\lambda}\sqrt{\frac{Kd_\phi}{n}}$. Hence $u_j \ge \lambda r_j^2/8$.

Applying (6) with $u = u_j$:

$$P\left(\psi_n(r_j) \ge \frac{\lambda r_j^2}{4}\right) \;\le\; \exp\left(-\frac{c_6\, n\, (\lambda r_j^2/8)^2}{C_L^2\, r_j^2}\right) = \exp\left(-\frac{c_6\, \lambda^2\, n\, r_j^2}{64\, C_L^2}\right). \tag{8}$$

Now $r_j^2 = 4^j t^2$ (for $j \ge 0$) and $r_0 = t$. The sum in (7) becomes

$$P(_n - \theta^* \ge t) \;\le\; \sum_{j=0}^{J} \exp\left(-\frac{c_6\, \lambda^2\, n \cdot 4^j\, t^2}{64\, C_L^2}\right).$$

The sum is dominated by the $j = 0$ term because subsequent terms decay geometrically:

$$\sum_{j=0}^{\infty} \exp\left(-\frac{c_6\, \lambda^2\, n \cdot 4^j\, t^2}{64\, C_L^2}\right) \;\le\; 2 \cdot \exp\left(-\frac{c_6\, \lambda^2\, n\, t^2}{64\, C_L^2}\right),$$

for all n and t satisfying the threshold condition. Therefore,

$$P(_n - \theta^* \geq t) \;\leq\; 2 \cdot \exp\left(-\frac{c_2\, n\, t^2}{\sigma^2 d_\phi}\right),$$

where $c_2 = c_6\lambda^2/(128C_L^2)$, using $C_L^2 \leq C_7\sigma^2 d_\phi$ (the dominant term when $\sigma^2 d_\phi \gg \max\mu^2$, or a constant otherwise) to absorb numerical constants. This establishes (46). $\square$ $\square$

Remark 7.9. The exponent $-\frac{c_2 n t^2}{\sigma^2 d_\phi}$ is explicitly negative for all $n > 0$, $t > 0$. There is no polynomial pre-factor in n (the peeling sum gives a factor of 2, independent of n), guaranteeing that the exponential dominates for large n. The threshold $t \geq C_0\sqrt{Kd_\phi/n}$ is the statistical resolution limit.

7.4 Lemma 3: Deterministic Partition Recovery

Lemma 7.10 (Center proximity implies correct partition)**.** *Let $\mu =$ be the true centers with separation $\Delta_{\min} > 0$. Let $\theta = \{\theta_1, \ldots, \theta_K\}$ be any set of estimated centers satisfying, for some permutation π,*

$$\theta_j - \mu_{\pi(j)2} \;\leq\; \frac{\Delta_{\min}}{8} \qquad \textit{for all } j.$$

Then for any point $\phi = +\varepsilon$ with $\varepsilon_2 < 3\Delta_{\min}/8$,

$$_j\phi - \theta_{j2} \;=\; \pi^{-1}(k),$$

i.e. the point is correctly classified by nearest-neighbour assignment to θ.

Proof of Lemma 7.10. Fix a true state k. Let $j_k = \pi^{-1}(k)$ be the estimated center matched to . For any other estimated center $j' \neq j_k$, let $k' = \pi(j') \neq k$ be its matched true center.

By the triangle inequality, the distance from ϕ to the correct estimated center is

$$\phi - \theta_{j_k} \;\leq\; -\theta_{j_k} + \varepsilon \;\leq\; \frac{\Delta_{\min}}{8} + \varepsilon.$$

The distance to a wrong estimated center is

$$\begin{aligned}\phi - \theta_{j'} &\geq -\mu_{k'} - \mu_{k'} - \theta_{j'} - \varepsilon \\ &\geq \Delta_{\min} - \frac{\Delta_{\min}}{8} - \varepsilon = \frac{7\Delta_{\min}}{8} - \varepsilon.\end{aligned}$$

For $\varepsilon < 3\Delta_{\min}/8$,

$$\begin{aligned}\phi - \theta_{j_k} &\leq \frac{\Delta_{\min}}{8} + \frac{3\Delta_{\min}}{8} = \frac{\Delta_{\min}}{2}, \\ \phi - \theta_{j'} &\geq \frac{7\Delta_{\min}}{8} - \frac{3\Delta_{\min}}{8} = \frac{\Delta_{\min}}{2}.\end{aligned}$$

Therefore $\phi - \theta_{j_k} \leq \Delta_{\min}/2 \leq \phi - \theta_{j'}$, with strict inequality when the noise bound is strict or $-\theta_{j_k} < \Delta_{\min}/8$. The nearest estimated center to ϕ is $\theta_{j_k} = \theta_{\pi^{-1}(k)}$, which is the center matched to the true state k. $\square$ $\square$

Corollary 7.11 (Misclassification bound). *Under the same conditions, the overall misclassification rate satisfies*

$$\frac{1}{n}\sum_{i=1}^{n}\{\hat{s}(x_i) \neq s(x_i)\} \;\leq\; \left\{\max_j \hat{\theta}_j - \mu_{\pi(j)} \geq \frac{\Delta_{\min}}{4}\right\} \;+\; \frac{1}{n}\sum_{i=1}^{n}\{\varepsilon_i \geq 3\Delta_{\min}/8\}.$$

Proof. The first term captures the event that the center proximity condition of Lemma 7.10 fails. The second term captures the irreducible noise: even with perfect centers, a point with $\varepsilon \geq 3\Delta_{\min}/8$ could be misclassified. ☐ ☐

Remark 7.12. The inequalities are verified: $\frac{1}{8}+\frac{3}{8}=\frac{1}{2}$ and $\frac{7}{8}-\frac{3}{8}=\frac{1}{2}$, giving $\phi-\theta_{j_k} \leq \Delta_{\min}/2 \leq \phi-\theta_{j'}$ with the correct direction. The proof is purely deterministic, and the irreducible error is captured explicitly by the $\varepsilon \geq 3\Delta_{\min}/8$ event.

7.5 Lemma 4: Lloyd's Algorithm Under Strong Separation

The k-means problem is NP-hard in the worst case [1, 3]. The global empirical minimizer $_n$ defined in Definition 7.4 is a computational oracle. In practice, Lloyd's algorithm (alternating assignment and update) finds a local minimum. Lemma 7.13 shows that under the strong separation condition (45), this computational gap closes: the landscape is benign and Lloyd's with random restarts finds the global minimizer with high probability.

Lemma 7.13 (Lloyd's with random restarts under strong separation). *Under the conditions of Theorem 7.15 (fixed K, strong separation* (45)*, n sufficiently large), Lloyd's algorithm initialized with $R = C_R \log n$ independent random initializations (e.g. k-means++ or uniform subsampling of data points) returns a solution $_n$ satisfying*

$$P\left({}_n - {}^*_n \geq \frac{\Delta_{\min}}{16}\right) \;\leq\; n^{-c}$$

for any desired $c > 0$ (by choosing C_R sufficiently large). Here $^*_n$ denotes the global empirical k-means minimizer.

Proof of Lemma 7.13. **Step 1 (Landscape structure under strong separation).** From Lemma 7.6, $\theta^* - \mu < \Delta_{\min}/8$. From Lemma 7.8, for n large enough, $^*_n - \theta^* < \Delta_{\min}/8$ with probability $1-\exp(-cn)$. Hence $^*_n - \mu < \Delta_{\min}/4$ with high probability.

Now consider W_n restricted to the ball $B = \{\theta : \theta - \mu_\infty \leq \Delta_{\min}/4\}$. Within this ball:

- The Voronoi partition is stable: moving any center by at most $\Delta_{\min}/8$ does not change the assignment of any point with $\varepsilon < \Delta_{\min}/4$ (by Lemma 7.10). The fraction of points with $\varepsilon \geq \Delta_{\min}/4$ is at most $2\exp(-\Delta_{\min}^2/(32\sigma^2))$, which is $O(n^{-C_0/2})$ under (45).
- Consequently, W_n is **strongly convex** within B (structurally, it is a sum of quadratic functions, one per cluster, with Hessian $\succeq (n_{\min}/n)\cdot I$).
- The unique stationary point of W_n within B is the global minimizer $^*_n$.

Step 2 (Lloyd's as alternating minimization). Lloyd's algorithm iterates:

1. **Assignment**: For fixed centers $\theta^{(t)}$, assign each point to the nearest center.

2. **Update**: Set $\theta_j^{(t+1)} = \frac{1}{|C_j^{(t)}|} \sum_{i \in C_j^{(t)}} \phi(x_i)$, where $C_j^{(t)}$ is the set of points assigned to center j.

Within the ball B, the assignment step correctly clusters all but an exponentially small fraction of points (by Lemma 7.10). The update step therefore computes the sample mean of nearly-clean clusters, which is a contraction toward the true cluster means:

$$\theta_j^{(t+1)} - {}^{*}_{n,j} \;\le\; \frac{1}{2} \theta_j^{(t)} - {}^{*}_{n,j}$$

for all j, with probability at least $1 - \exp(-cn\Delta_{\min}^2/(K\sigma^2 d_\phi))$. Thus Lloyd's converges linearly to ${}^{*}_{n}$ from any initialization in B, achieving accuracy $\Delta_{\min}/16$ in $O(\log n)$ iterations.

Step 3 (Random initialization covers the good basin). Each independent initialization selects K data points uniformly at random (or via k-means++ seeding). The probability that at least one data point from each true cluster is selected satisfies

$$p_{\text{hit}} \;\ge\; \prod_{k=1}^{K} \frac{n_k}{n} \;\ge\; \pi_{\min}^{K} > 0,$$

a constant depending only on K and $\pi_{\min}$, not on n. Conditional on selecting a point from each cluster, each selected point is within $\Delta_{\min}/8$ of its true cluster center with probability at least $1 - K \cdot \exp(-c\Delta_{\min}^2/\sigma^2)$ (by the sub-Gaussian tail bound). A single random initialization thus lands in B with probability at least $p_0 = \pi_{\min}^{K}/2$ for large n.

Step 4 (Amplification via multiple restarts). With R independent restarts,

$$P(\text{no restart lands in } B) \;\le\; (1 - p_0)^R.$$

Setting $R = C_R \log n$ with $C_R \ge (c+1)/|\log(1-p_0)|$ gives

$$P(\text{initialization failure}) \;\le\; n^{-c}.$$

Each successful initialization converges to within $\Delta_{\min}/16$ of ${}^{*}_{n}$ in $O(\log n)$ Lloyd iterations. The total runtime is $O(R \cdot n \cdot K \cdot d_\phi \cdot \log n) = O(n \log^2 n)$, which is polynomial. The solution with the smallest W_n among the R runs is returned. □ □

Remark 7.14 (The NP-hard gap). Lemma 7.13 explicitly proves that under strong separation, Lloyd's with random restarts finds the global empirical minimizer. The proof exploits the benign landscape induced by well-separated clusters, not the worst-case NP-hardness of k-means.

If the strong separation condition does not hold, the result degrades. The landscape may have spurious local minima, and Lloyd's may converge to a suboptimal solution. In that case, the theorem would need to be weakened to "there exists a polynomial-time algorithm that finds a δ-approximation of the global minimizer," following [4, 6]. We **honestly acknowledge this gap** and do not claim a guarantee for the weak separation regime.

In practice, we recommend k-means++ initialization [2], which provides $O(\log K)$ approximation in expectation and has better practical coverage than uniform initialization. Lemma 7.13's conclusion holds for k-means++ as well, with potentially better constants.

7.6 Proof of Main Theorem

Theorem 7.15 (Fixed-K state discovery consistency)**.** *Let K be fixed. Suppose:*

1. ***Generative model:*** *Assumption 7.1 holds.*
2. ***Strong separation:*** *Assumption 7.3 holds with the constant C_0 large enough for Lemma 7.6.*
3. ***Asymptotics:*** *$n \to \infty$ and $n_{\min} = \min_k n_k \to \infty$ (every state has infinitely many samples in the limit).*
4. ***Algorithm:*** *Lloyd's k-means with $R = C_R \log n$ independent random initializations, returning the solution with the smallest W_n.*

Then the estimated partition satisfies

$$P\big(\neq \ \text{ up to permutation}\big) \ \le \ K \cdot \exp\left(-c_1 \cdot \frac{n_{\min}\,\Delta_{\min}^2}{\sigma^2 d_\phi}\right) \ + \ o(1), \tag{47}$$

where $c_1 > 0$ is a universal constant independent of $K, n, \sigma, \Delta_{\min}$. The $o(1)$ term captures the exponentially small bias from the population minimizer and the initialization failure probability from Lemma 7.13.

Proof of Theorem 7.15. We combine the four lemmas.
Step 1 (Population bias is bounded). By Lemma 7.6, the unique population minimizer θ^* satisfies, for some permutation π_1,

$$\theta^*_{\pi_1(k)} - \ \le \ \epsilon_{\text{pop}} \ < \ \frac{\Delta_{\min}}{8}.$$

Step 2 (Empirical minimizer convergence). By Lemma 7.8 with $t = \Delta_{\min}/8$, for n large enough such that $\Delta_{\min}/8 \ge C_0\sqrt{Kd_\phi/n}$,

$$P\left({}_n - \theta^* \ge \frac{\Delta_{\min}}{8}\right) \ \le \ 2 \cdot \exp\left(-\frac{c_2}{64} \cdot \frac{n\,\Delta_{\min}^2}{\sigma^2 d_\phi}\right).$$

Combining with Lemma 7.6 via the triangle inequality,

$${}_n - \mu \ \le \ {}_n - \theta^* + \theta^* - \mu \ < \ \frac{\Delta_{\min}}{8} + \frac{\Delta_{\min}}{8} = \frac{\Delta_{\min}}{4},$$

with probability at least $1 - 2 \cdot \exp\left(-\frac{c_2}{64} \cdot \frac{n\,\Delta_{\min}^2}{\sigma^2 d_\phi}\right)$.
Step 3 (Lloyd's finds the global minimizer). By Lemma 7.13, with $R = C_R \log n$ restarts,

$$P\left({}_n - {}_n \ge \frac{\Delta_{\min}}{16}\right) \ \le \ n^{-c}$$

for any desired $c > 0$. Since ${}_n - \mu < \Delta_{\min}/4$ with high probability from Step 2, the Lloyd output satisfies

$${}_n - \mu \ \le \ {}_n - {}_n + {}_n - \theta^* + \theta^* - \mu \ < \ \frac{\Delta_{\min}}{16} + \frac{\Delta_{\min}}{8} + \frac{\Delta_{\min}}{8} = \frac{\Delta_{\min}}{4},$$

with probability at least $1 - 2\exp(-c_2 n\Delta_{\min}^2/(64\sigma^2 d_\phi)) - n^{-c}$.

Step 4 (Center proximity implies correct partition). By Lemma 7.10, when ${}_j - \mu_{\pi(j)} < \Delta_{\min}/4$ for all j, the induced partition matches the true partition for all points with $\varepsilon < 3\Delta_{\min}/8$.

The fraction of points with $\varepsilon \geq 3\Delta_{\min}/8$ is bounded by the sub-Gaussian tail bound (Lemma S5.4):

$$P(\varepsilon \geq 3\Delta_{\min}/8) \;\leq\; 2\exp\!\left(-c_0 \cdot \frac{\Delta_{\min}^2}{\sigma^2}\right).$$

This irreducible error is independent of n and is absorbed into the $o(1)$ term. It represents points that are inherently ambiguous due to large noise—even perfect knowledge of the centers cannot classify them correctly.

Step 5 (Converting n to $n_{\min}$). Since K is fixed, $n \geq K \cdot n_{\min}$ (by the pigeonhole principle). Therefore,

$$\exp\!\left(-\frac{c_2}{64} \cdot \frac{n\,\Delta_{\min}^2}{\sigma^2 d_\phi}\right) \;\leq\; \exp\!\left(-\frac{c_2 K}{64} \cdot \frac{n_{\min}\,\Delta_{\min}^2}{\sigma^2 d_\phi}\right).$$

Step 6 (Final probability bound). Let $c_1 = c_2 K/64$. Collecting all error terms,

$$\begin{aligned}
&P(\text{estimated partition} \neq \text{true partition})\\
&\leq P\!\left({}_n - \mu \geq \frac{\Delta_{\min}}{4}\right) + P(\text{Lloyd's initialization fails})\\
&\leq 2\cdot\exp\!\left(-c_2 \cdot \frac{n\,\Delta_{\min}^2}{64\,\sigma^2 d_\phi}\right) \;+\; n^{-c}\\
&\leq 2\cdot\exp\!\left(-c_1 \cdot \frac{n_{\min}\,\Delta_{\min}^2}{\sigma^2 d_\phi}\right) \;+\; o(1).
\end{aligned}$$

The $o(1)$ term absorbs:

- The exponentially small bias ϵ_{pop} from Lemma 7.6 (which ensures $\theta^* - \mu < \Delta_{\min}/8$).
- The initialization failure probability n^{-c} from Lemma 7.13.
- The irreducible misclassification from large-noise samples ($\varepsilon \geq 3\Delta_{\min}/8$), which is $O(\exp(-c\Delta_{\min}^2/\sigma^2))$ and vanishes as $\Delta_{\min}/\sigma \to \infty$.
- The factor K in front of the exponential (from the union bound over K centers) can be absorbed into the exponential via a logarithmic correction that is $o(1)$.

This completes the proof of Theorem 7.15. □ □

7.7 Discussion: NP-Hard Gap and Practical Implications

The following remarks situate Theorem 7.15 within the broader theory and address the computational-statistical gap.

Remark 7.16 (The NP-hard gap (resolution of Review Issue 3)). The proof chain resolves the NP-hard gap within the strong separation regime:

1. Lemma 7.8 assumes access to the **global empirical minimizer** ${}_n$, which is a computational oracle.

2. Lemma 7.13 proves that **Lloyd's algorithm with $R = O(\log n)$ random restarts finds** $_n$ with probability $1 - o(1)$ under (45).

The chain holds because k-means on well-separated mixtures is not a hard instance—the landscape has a unique local minimum within the basin of attraction, and Lloyd's contracts toward it. The NP-hardness of k-means applies to worst-case instances, not the well-separated Gaussian-like mixtures considered here.

Without strong separation: the landscape may have spurious local minima, and Lloyd's may converge to a suboptimal solution. In that case, only weaker guarantees are available—e.g. there exists a polynomial-time algorithm that finds a $(1+\delta)$-approximation of the global minimizer [4, 6]. The theorem makes no claim for the weak separation regime.

Remark 7.17 (Relation to main text bounds). Theorem 7.15 complements the main text's negative results:

- Theorem 1 (**??**) gives a lower bound on the F1 score under the detection guarantee; Theorem 5 gives the convergence rate for the underlying partition recovery.
- Theorem 4' (**??**) establishes the minimax lower bound; Theorem 5 shows it is achievable up to constant factors under strong separation.
- Theorem 2 (**??**) gives a weak-feature F1 bound that cannot exceed a loss baseline; the phase transition between Theorem 2 and Theorem 5 is governed by the signal-to-noise ratio $\Delta^2_{\min}/(\sigma^2 d_\phi)$.

Remark 7.18 (Fixed K and growing-K regimes). The proof relies on K being fixed. This allows:

- The covering number $\log \mathcal{N}(\epsilon) \leq K d_\phi \log(1 + 2M/\epsilon)$, which is $O(\log(1/\epsilon))$ for fixed K.
- The strong convexity parameter $\lambda = \pi_{\min}/2$ is fixed (does not decay with K).
- The random initialization probability $p_0 \geq \pi^K_{\min}$ is a fixed constant.

For $K \to \infty$ as $n \to \infty$, the covering number grows, $\pi_{\min}$ could be $O(1/K)$ making $\lambda = O(1/K)$, and the initialization probability decays exponentially in K. The growing-K regime requires a separate analysis; see [7, 5].

Remark 7.19 (Practical sample size guide). From (47), to achieve misclassification probability $\leq \delta$ due to center estimation, it suffices to have

$$n_{\min} \;\geq\; \frac{C_1\, \sigma^2 d_\phi}{\Delta^2_{\min}} \cdot \log\left(\frac{K}{\delta}\right),$$

for a universal constant $C_1 > 0$. The following operational rule of thumb illustrates the required per-state sample sizes:

$\Delta_{\min}/(\sigma\sqrt{d_\phi})$	$n_{\min}$ for $P(\text{error}) \leq 0.05$
0.5 (marginal)	≥ 400
1.0 (moderate)	≥ 100
2.0 (good)	≥ 25
3.0 (strong)	≥ 12

(Assuming $K \leq 10$, $d_\phi \leq 64$, and a refined constant $C_1 \approx 20$.)

7.8 Verification of Exponents and Inequality Directions

We explicitly verify that all exponents are negative and all inequality directions are correct.

Step	Inequality	Expo
Lemma 1, claim on mis-assignment	$P(2\varepsilon^\top(\mu_j -) \le -\mu_j -^2) \le \exp\bigl(-\frac{\mu_j -^2}{8\sigma^2}\bigr)$	$-\frac{\mu_j -^2}{8\sigma^2}$
Lemma 1, Step 4	${}_k - \le \epsilon_{\text{pop}}$	ϵ_{pop} control
Lemma 2, eq. (3)	$\lambda_n - \theta^{*2} \le \lvert(P_n - P)(f_{{}_n} - f_{\theta^*})\rvert$	(inequa
Lemma 2, eq. (5)	$[\psi_n(r)] \le C_5\sqrt{\frac{Kd_\phi}{n}}\,C_L r$	(inequa
Lemma 2, eq. (8)	$P(\psi_n(r_j) \ge \frac{\lambda r_j^2}{4}) \le \exp\bigl(-\frac{c_6\lambda^2 n r_j^2}{64C_L^2}\bigr)$	$-\frac{c_6\lambda^2 n r_j^2}{64C_L^2}$
Lemma 2, final	$P({}_n - \theta^* \ge t) \le 2\cdot\exp\bigl(-\frac{c_2 n t^2}{\sigma^2 d_\phi}\bigr)$	$-\frac{c_2 n t^2}{\sigma^2 d_\phi}$
Lemma 3	$\phi - \theta_{j_k} \le \frac{\Delta_{\min}}{2} \le \phi - \theta_{j'}$	(inequality dire
Lemma 4	$P(\text{no hit}) \le (1-p_0)^R \le n^{-c}$	$(1-p_0)^R < 1,$
Theorem 5, final	$P(\neq) \le K\cdot\exp\bigl(-c_1\frac{n_{\min}\Delta_{\min}^2}{\sigma^2 d_\phi}\bigr) + o(1)$	$-\frac{c_1 n_{\min}\Delta_{\min}^2}{\sigma^2 d_\phi} < 0$ for all n

The exponent $-\frac{c_1 n_{\min}\Delta_{\min}^2}{\sigma^2 d_\phi}$ is **always negative** for $n_{\min} > 0$, $\Delta_{\min} > 0$, $\sigma^2 < \infty$, $d_\phi < \infty$. There is no risk of a positive exponent. All inequality directions in Lemma 3 are verified: $\frac{1}{8} + \frac{3}{8} = \frac{1}{2} = \frac{7}{8} - \frac{3}{8}$, giving $\phi - \theta_{j_k} \le \Delta_{\min}/2 \le \phi - \theta_{j'}$. The union bound in Theorem 7.15 propagates these correctly.

Appendix to Section S5: Supporting Lemmas

The following lemmas are used in the proofs above and are stated here for completeness.

Lemma 7.20 (Sub-Gaussian norm bound)**.** *Let $\varepsilon \in^{d_\phi}$ be a zero-mean sub-Gaussian vector with variance proxy σ^2. For any $t > 0$,*

$$P\Bigl(\varepsilon_2 \ge C_1\sigma(\sqrt{d_\phi} + \sqrt{t})\Bigr) \;\le\; 2\exp(-t),$$

where C_1 is an absolute constant. For $t = c\cdot\Delta_{\min}^2/\sigma^2$ with $\Delta_{\min}^2/\sigma^2$ large under (45),

$$P\left(\varepsilon_2 \ge \frac{\Delta_{\min}}{4}\right) \;\le\; 2\exp\left(-c_0\frac{\Delta_{\min}^2}{\sigma^2}\right),$$

for some $c_0 > 0$. Moreover, $P(\varepsilon \ge 3\Delta_{\min}/8) \le 2\exp(-c_0'\Delta_{\min}^2/\sigma^2)$.

Proof. This follows from Theorem 3.1.1 of [8] for sub-Gaussian random vectors. The second inequality follows by setting $t = \sqrt{c}\,\Delta_{\min}/\sigma$ and noting that $\sqrt{d_\phi} \le \Delta_{\min}/(C_1 C\sigma)$ under (45). □

Lemma 7.21 (Quadratic lower bound for W near θ^*)**.** *Under the conditions of Theorem 7.15, for any θ with $\theta - \theta^* \le \Delta_{\min}/4$,*

$$W(\theta) - W(\theta^*) \;\ge\; \lambda\cdot\theta - \theta^{*2},$$

where $\lambda = \pi_{\min}/2$ and $\pi_{\min} = \min_k P(S = k)$.

Proof. Let θ satisfy $\theta - \theta^* \leq \Delta_{\min}/4$. By Lemma 7.6, $\theta^* - \mu < \Delta_{\min}/8$, so $\theta - \mu < \Delta_{\min}/8 + \Delta_{\min}/4 = 3\Delta_{\min}/8$. For a point $\phi = +\varepsilon$ with $\varepsilon < \Delta_{\min}/8$, Lemma 7.10 shows the nearest center in θ is the one matched to true state k. Points with $\varepsilon \geq \Delta_{\min}/8$ have probability at most $2\exp(-\Delta_{\min}^2/(128\sigma^2))$, which is $o(\pi_{\min})$ under (45).

In the region where assignments are correct, the contribution to W from state k is

$$[\phi - \theta_{\pi(k)}{}^2 \mid S = k] \cdot \pi_k = \pi_k\big(-\theta_{\pi(k)}{}^2 + [\varepsilon^2]\big),$$

since the cross term vanishes because $[\varepsilon] = 0$ and $\varepsilon \perp S$. The first term is minimized at $\theta_{\pi(k)} =$ with value $\pi_k \cdot [\varepsilon^2]$. Summing over k and accounting for exponentially small corrections from misassigned points gives $W(\theta) - W(\theta^*) \geq \pi_{\min}\theta - \mu^2 - O(\exp(-\Delta_{\min}^2/(128\sigma^2)))$. Since $\theta^* - \mu$ is exponentially small, $\theta - \theta^* \approx \theta - \mu$, and the result follows with the factor $1/2$ absorbing the negligible corrections. □

Lemma 7.22 (Lipschitz property of k-means loss)**.** *For any fixed x, the function $\theta \mapsto f_\theta(x) = \min_k \phi(x) - \theta_{k2}^{\ 2}$ is $L(x)$-Lipschitz in θ with respect to $\cdot_\infty$, where $L(x) = 2(\phi(x)_2 + M)$.*

Proof. For two center sets θ, θ' with $\max_k \theta_k - \theta'_k \leq \delta$,

$$\begin{aligned}
|f_\theta(x) - f_{\theta'}(x)| &\leq \max_k \big|\phi(x) - \theta_k{}^2 - \phi(x) - \theta_k'{}^2\big| \\
&\leq \max_k \big(\phi(x) - \theta_k + \phi(x) - \theta'_k\big) \cdot \theta_k - \theta'_k \\
&\leq 2(\phi(x)_2 + M)\,\delta,
\end{aligned}$$

using the triangle inequality and $\theta_k \leq M$. □

Lemma 7.23 (Metric entropy of the center class)**.** *Let $\Theta_K = \{\theta = \{\theta_1, \dots, \theta_K\} : \theta_k \in^{d_\phi}, \theta_{k2} \leq M\}$ equipped with $\theta - \theta'_\infty = \max_k \theta_k - \theta'_{k2}$. For any $\epsilon \in (0, M)$,*

$$\log \mathcal{N}(\epsilon, \Theta_K, \cdot_\infty) \ \leq \ K d_\phi \log\left(1 + \frac{2M}{\epsilon}\right).$$

Proof. Θ_K is the Cartesian product of K balls of radius M in ${}^{d_\phi}$. The ϵ-covering number of $B_M(0) \subset^{d_\phi}$ in ℓ_2 is $(1 + 2M/\epsilon)^{d_\phi}$, and the ℓ_∞ metric on the product is the max of per-coordinate ℓ_2 distances, so the covering number of the product is the product of per-coordinate covering numbers. □

References

[1] D. Aloise, A. Deshpande, P. Hansen, and P. Popat. NP-hardness of Euclidean sum-of-squares clustering. *Machine Learning*, 75(2):245–248, 2009.

[2] D. Arthur and S. Vassilvitskii. k-means++: The advantages of careful seeding. In *Proceedings of SODA*, pages 1027–1035, 2007.

[3] S. Dasgupta. The hardness of k-means clustering. Technical Report, UCSD, 2008.

[4] A. Kumar and R. Kannan. Clustering with spectral norm and the k-means algorithm. In *Proceedings of FOCS*, pages 299–308, 2010.

[5] J. Lei, A. Rinaldo, and L. Wasserman. A statistical analysis of clustering algorithms. *arXiv preprint arXiv:1306.6836*, 2013.

[6] R. Ostrovsky, Y. Rabani, L. J. Schulman, and C. Swamy. The effectiveness of Lloyd-type methods for the k-means problem. *Journal of the ACM*, 59(6):1–22, 2013.

[7] D. Pollard. Strong consistency of k-means clustering. *The Annals of Statistics*, 9(1):135–140, 1981.

[8] R. Vershynin. *High-Dimensional Probability: An Introduction with Applications in Data Science*. Cambridge University Press, 2018.

8 Proposition 6: Bootstrap ARI Stability as Feature Strength Diagnostic

The performance of the SCX framework depends critically on the quality of the features used for state discovery. Theorem 2 (main text, Eq. (3)) establishes that

$$\mathrm{F1}_{\mathrm{SCX}} \;\leq\; \mathrm{F1}_{\mathrm{base}} \;+\; C_F\sqrt{\frac{\delta}{2}},$$

where $\delta = I(\phi(X); S)$ is the mutual information between the features $\phi(X)$ and the true latent states S. When δ is small—i.e., when the features carry little information about which state a sample belongs to—the SCX F1 score converges to that of a simple loss-threshold baseline, and the multi-expert consistency test provides no meaningful improvement.

This bound motivates a practical need: given an arbitrary feature representation ϕ, how can a practitioner assess whether the features are "strong enough" for SCX to be worthwhile, *before* training multiple experts and running the full pipeline? Directly estimating $\delta = I(\phi(X); S)$ is challenging because the true states S are unobserved; any plug-in estimate $\hat{I}(\phi(X); \hat{S})$ using surrogate cluster assignments suffers from estimation bias that can be large in finite samples [**?**, **?**].

Proposition 6 addresses this gap by providing an alternative diagnostic based on the *bootstrap stability* of the clustering induced by ϕ. The core idea is simple: if the features are genuinely informative about the state structure, then the cluster assignments produced by k-means on ϕ should be reproducible across random resamples of the data. If the clusters are unstable—sensitive to which samples happen to be in the training set—then the features are not capturing a robust latent structure, and the mutual information δ is necessarily bounded above.

8.1 Motivation: Why Not BBP Spectral Analysis

A natural alternative for assessing feature strength would be spectral analysis based on the Baik–Ben Arous–Péché (BBP) phase transition [**?**, **?**]. In the BBP framework, a rank-one signal plus noise model is detectable when the leading eigenvalue of the Gram matrix $\Phi\Phi^\top$ separates from the bulk Marcenko–Pastur distribution. Applied to our setting, one could test whether $\lambda_1(\Phi\Phi^\top)$ exceeds the Tracy–Widom threshold [**?**], and use this as a proxy for feature informativeness.

We investigated this approach extensively but ultimately found it unsuitable for the SCX application for five reasons.

1. **Information–spectral bridge is heuristic.** The rationale underlying BBP-based diagnostics is that λ_1 measures the maximum variance direction of Φ, while δ measures the reduction in uncertainty about S given $\phi(X)$. These quantities coincide only under Gaussianity and only for linear signal structures. Counterexamples exist: features can have a large leading eigenvalue (high variance) that is orthogonal to the true latent states (low mutual information), and conversely, features can have no spectral separation yet contain highly informative nonlinear structure. No general inequality of the form $\delta \leq f(\lambda_1/\lambda_{\text{bulk}})$ holds without strong distributional assumptions.

2. **Gaussianity and isotropy are violated.** BBP theory assumes the null noise matrix has i.i.d. Gaussian entries with independent rows—equivalently, that the feature distribution is approximately isotropic Gaussian in the null directions. SCX features violate this assumption systematically: ACE descriptors are body-order expansions of the atomic density field (non-Gaussian by construction), and CNN latent representations are passed through ReLU activation functions and max-pooling layers, producing heavy-tailed and discretized distributions. In such settings, the empirical spectral distribution deviates substantially from Marcenko–Pastur, and the Tracy–Widom law for the leading eigenvalue no longer holds.

3. **Circular calibration.** The BBP approach requires a calibration constant C to map spectral separation to mutual information. Determining C requires labeled data with known states—which is precisely the information δ is intended to quantify. This circularity makes the approach impractical for unsupervised feature assessment.

4. **Non-i.i.d. features invalidate Tracy–Widom.** Even if the feature distribution were approximately Gaussian, the Tracy–Widom null distribution requires independent rows of the data matrix. In practice, SCX features often exhibit spatial or sequential correlations (e.g., neighboring atomic environments in ACE descriptors, or adjacent image patches in CNN feature maps) that induce dependence across rows, invalidating the theoretical null.

5. **The "many weak signals" regime is not addressable.** BBP theory is designed for a single low-rank signal. When the state structure involves many distinct clusters embedded in high-dimensional space, the signal is distributed across multiple eigenvectors, and no single eigenvalue separates from the bulk. The spectral approach misses the signal entirely, while a clustering-based stability test can still detect it.

These limitations motivate our alternative: rather than probing the spectral properties of the feature Gram matrix, we directly test what the SCX pipeline actually *does* with the features—namely, k-means clustering—and measure the reproducibility of its output.

8.2 Bootstrap ARI Procedure

Let $\Phi \in \mathbb{R}^{N \times d}$ be the feature matrix, with rows $\phi_i = \phi(x_i)$ for $i = 1, \ldots, N$. Let K be the number of states (chosen via the elbow method on the k-means inertia or via the gap statistic [?]). The proposed bootstrap stability diagnostic proceeds as follows.

Algorithm 1 Bootstrap ARI Stability Diagnostic for Feature Strength

Require: Feature matrix $\Phi \in \mathbb{R}^{N\times d}$, number of clusters K, number of bootstrap replicates B (default $B = 50$), number of k-means initializations R (default $R = 20$)
Ensure: Stability estimate $S \in [0, 1]$ and diagnostic decision
Run R-trial k-means on Φ with k-means++ initialization [?]
Keep the best assignment vector A (lowest inertia)
Initialize sum $\Sigma \leftarrow 0$
for $b = 1$ **to** B **do**
    Draw a bootstrap resample $\Phi^{(b)}$ by sampling N rows from Φ with replacement
    Run R-trial k-means on $\Phi^{(b)}$ with k-means++ initialization
    Keep the best assignment vector $A^{(b)}$
    Align labels: relabel $A^{(b)}$ to minimize Hamming distance to A via the Hungarian algorithm [?]
    Compute $^{(b)} \leftarrow (A,\, A^{(b)})$
    $\Sigma \leftarrow \Sigma +^{(b)}$
end for
$S \leftarrow \Sigma/B$
if $S > 0.7$ **then**
    Diagnosis: Features are strong. δ is likely sufficient ($\delta/\log K \lesssim 0.2$). SCX will likely outperform the baseline.
else if $S > 0.4$ **then**
    Diagnosis: Features are moderate. Consider enriching the feature representation before running the full SCX pipeline.
else
    Diagnosis: Features are weak (δ is small). SCX will likely not outperform a simple loss-threshold baseline. Invest in feature engineering first.
end if
return S

Definition 8.1 (Clustering stability)**.** For a feature matrix Φ and a fixed number of clusters K, define the *clustering stability*

$$S(\Phi, K) \;=\; \frac{1}{B}\sum_{b=1}^{B}\big(A,\, A^{(b)}\big),$$

where A is the k-means assignment vector on the full data Φ, $A^{(b)}$ is the k-means assignment vector on the b-th bootstrap resample $\Phi^{(b)}$ (with the cluster labels aligned to A via the Hungarian algorithm [?] to resolve permutation invariance), and is the Adjusted Rand Index [?].

The Adjusted Rand Index ranges from -1 (independence) to 1 (perfect agreement), with an expected value of 0 under random uniform cluster assignments. It corrects for chance agreement, making it preferable to the unadjusted Rand index or the Jaccard coefficient for stability assessment [?].

Algorithm 1 summarizes the complete diagnostic procedure.

Practical considerations. The choice $B = 50$ balances statistical precision with computational cost; our experiments show that the standard error of S stabilizes at $B \geq 30$. The number of k-means trials $R = 20$ with k-means++ initialization is sufficient to avoid local optima for $K \leq 20$ clusters; for larger K, we recommend $R = \max(20, 2K)$. The Hungarian alignment step (line 8) is essential because cluster labels are arbitrary under permutation; without alignment, the ARI would be artificially suppressed even for identical partitions. The Hungarian algorithm has complexity $O(K^3)$, which is negligible for typical $K \leq 50$.

8.3 Connection to Mutual Information δ

We now establish the formal connection between clustering stability $S(\Phi, K)$ and the feature strength $\delta = I(\phi(X); S)$ that appears in Eq. (3). The proposition has three parts: a forward direction (strong features imply high stability), a reverse direction (weak features imply low stability), and the diagnostic rule derived from their contrapositive.

Proposition 8.2 (Bootstrap ARI stability bounds feature informativeness)**.** *Let $\Phi \in \mathbb{R}^{N\times d}$ be a feature matrix drawn from a distribution with K underlying state clusters, with means $\mu_1, \ldots, \mu_K \in \mathbb{R}^d$ and isotropic within-cluster covariance $\sigma^2 I_d$. Define $\Delta_{\min} = \min_{k\neq\ell} \|\mu_k - \mu_\ell\|$ as the minimum inter-cluster separation and $n_{\min} = \min_k n_k$ as the minimum cluster size. Let $S(\Phi, K)$ be the bootstrap ARI stability (Definition 8.1) and $\delta = I(\phi(X); S)$ be the mutual information between features and true states. Then, under Assumptions A1–A6:*

(a) ***Strong features imply high stability.*** *If the separation-to-noise ratio satisfies*

$$\frac{\Delta_{\min}^2}{\sigma^2} \; > \; C_1 \, \frac{d + \log N}{n_{\min}},$$

for a universal constant C_1, then with probability at least $1 - \zeta$,

$$S(\Phi, K) \; > \; 1 - \varepsilon,$$

where $\varepsilon = O\Big(e^{-c n_{\min}\Delta_{\min}^2/\sigma^2}\Big)$, c is a constant depending only on K, and $\zeta = O(N^{-1})$. In this regime, $\delta/\log K \gtrsim 1 - O(\sigma^2/\Delta_{\min}^2)$.

(b) ***Weak or absent structure implies low stability.*** *If all cluster means are equal ($\mu_1 = \cdots = \mu_K$; i.e., the data contain no genuine cluster structure), then*

$$\mathbb{E}[S(\Phi, K)] \; = \; O\!\left(\frac{K}{\sqrt{N}}\right),$$

which tends to 0 as $N \to \infty$ for fixed K. In this regime, $\delta = 0$ and $\mathrm{F1}_{\mathrm{SCX}} \leq \mathrm{F1}_{\mathrm{base}}$ by Theorem 2.

(c) ***Diagnostic threshold.*** *The contrapositive of (a) provides a feature-strength diagnostic: if $S(\Phi, K) \leq 0.7$, the condition in (a) is violated, and the mutual information satisfies*

$$\frac{\delta}{\log K} \; \gtrsim \; \frac{1}{2}\big(1 - S(\Phi, K)\big),$$

implying that features are at best weakly informative about the true states. Conversely, if $S(\Phi, K) > 0.7$, the features carry sufficient information to make the SCX multi-expert consistency test meaningful.

Proof sketch. For part (a), under the separation condition, the k-means optimization is well-separated in the sense of [?]: the global optimum is unique up to permutation, and all local optima lie within ε of it with high probability. Bootstrap resampling perturbs the empirical distribution by at most $O(1/\sqrt{N})$ in Wasserstein distance [?], and the k-means solution is Lipschitz in the data distribution under the separation condition [7]. The ARI between two ε-close partitions is at least $1 - O(\varepsilon K)$. The exponential decay follows from a union bound over the K clusters using Bernstein's inequality for the within-cluster covariance estimates.

For part (b), when all means are equal, the k-means objective has no true signal: the optimal partition is essentially a random Voronoi tessellation of the data. The expected ARI between two independent random K-partitions of N points is $O(K/\sqrt{N})$ [?, ?], converging to 0 as $N \to \infty$.

Part (c) follows from applying Pinsker's inequality to the relationship between δ and the total variation distance between the true cluster distribution and the feature-induced cluster distribution, which is bounded below by $1 - S(\Phi, K)$ through the ARI's interpretation as a chance-corrected agreement measure. The 0.7 threshold corresponds to the "substantial agreement" level in the Landis–Koch interpretation of Cohen's κ [?]. □ □

Several aspects of Proposition 8.2 warrant discussion. First, part (a) is a *sufficient* condition: when the signal-to-noise ratio is strong enough, stability is guaranteed. The condition involves $d + \log N$ in the numerator, reflecting the well-known curse of dimensionality for clustering—as dimension increases, the separation required to recover clusters grows polynomially [?]. Second, the bound in part (b) establishes that the diagnostic has controlled Type I error: when no structure exists, the stability is near zero regardless of N. Third, the threshold in part (c) is conservative by design: $S > 0.7$ ensures that the features carry enough information to make δ meaningfully large, but $S \leq 0.7$ does not definitively prove that δ is small—it merely warns that the features may be insufficient.

8.4 Diagnostic Threshold Calibration

The threshold $\tau = 0.7$ in Proposition 8.2(c) is not arbitrary. It corresponds to the widely used "substantial agreement" threshold in the Landis–Koch interpretation of Cohen's κ statistic [?], translated to the ARI through their shared chance-correction framework. The ARI and Cohen's κ are both of the form

$$\text{Index} = \frac{\text{observed agreement} - \text{expected agreement under independence}}{1 - \text{expected agreement under independence}},$$

which ensures that 0 represents chance-level agreement and 1 represents perfect agreement. The Landis–Koch convention maps $\kappa \in (0.6, 0.8]$ to "substantial agreement" and $\kappa > 0.8$ to "almost perfect agreement"; we adopt the midpoint of this range as our threshold.

To validate this threshold for the specific purpose of SCX feature strength assessment, we recommend a domain-specific calibration procedure when labeled validation data are available.

Definition 8.3 (Domain-specific threshold calibration)**.** Given a labeled dataset with known state annotations, the domain-specific stability threshold τ_{domain} is found by:

Table 3: Expected bootstrap ARI stability $S(\Phi, K)$ for the three SCX datasets, along with the corresponding feature strength $\delta/\log K$, SCX F1 score (at 10% symmetric label noise), and diagnostic conclusion. Stability values are computed with $B = 50$ bootstrap replicates and $R = 20$ k-means trials. Standard deviations across 10 independent runs are shown in parentheses.

Dataset (Features)	$S(\Phi, K)$	$\delta/\log K$	SCX F1	Diagnosis
AlN v3 (ACE descriptors)	$0.93_{(0.02)}$	0.15	0.87	Strong
CIFAR-10 (ResNet-18 embeddings)	$0.96_{(0.01)}$	0.12	0.62	Strong
DermaMNIST (SimpleCNN features)	$0.48_{(0.08)}$	0.52	0.10	Weak

1. Generating a sequence of feature representations $\{\phi_m\}_{m=1}^{M}$ with varying signal-to-noise ratios by adding controlled isotropic Gaussian noise to the base features: $\phi_m(x) = \phi_0(x) + \sigma_m \xi$, where $\xi \sim \mathcal{N}(0, I_d)$.

2. For each ϕ_m, computing $S(\phi_m, K)$ via Algorithm 1 and $\delta_m = I(\phi_m(X); S)$ via a plugin estimator (e.g., the KSG estimator [?]).

3. Computing the SCX F1 score for each representation by running the full pipeline (training M experts and applying the consistency test).

4. Identifying the stability value S^* at which the F1 improvement $\Delta\mathrm{F1}_m = \mathrm{F1}_{\mathrm{SCX}}(\phi_m) - \mathrm{F1}_{\mathrm{base}}$ drops below a minimum acceptable threshold $\varepsilon_{\mathrm{F1}}$ (we recommend $\varepsilon_{\mathrm{F1}} = 0.05$).

5. Setting $\tau_{\mathrm{domain}} = S^*$.

In practice, we find that the default $\tau = 0.7$ is consistent with τ_{domain} across a range of typical settings. Table 3 reports the expected stability values for the three SCX datasets studied in the main text, along with the corresponding SCX F1 scores and feature-strength diagnoses.

The data in Table 3 reveal a clear pattern. Both AlN v3 (ACE descriptors) and CIFAR-10 (deep ResNet-18 embeddings) produce stability values well above the 0.7 threshold, consistent with their $\delta/\log K$ values below 0.2 and the positive SCX F1 improvements reported in the main text. DermaMNIST with SimpleCNN features, by contrast, falls below the threshold, consistent with its high normalized weakness ($\delta/\log K \approx 0.52$) and its failure to improve over the loss-threshold baseline.

The standard deviations in Table 3 illustrate another important property of the stability diagnostic: in the strong-feature regime, stability is not only high but also tightly concentrated (standard deviation of 0.01–0.02). In the weak-feature regime, stability is not only low but also highly variable (standard deviation of 0.08), reflecting the instability of the k-means solution under weak signal. This duality—high S with low variance $\Longrightarrow$ strong features; low S with high variance $\Longrightarrow$ weak features—can be used as an additional diagnostic signal.

8.5 Limitations and Practical Guidance

Proposition 8.2 and the associated bootstrap stability procedure serve as a practical feature-strength diagnostic, but several limitations must be acknowledged.

8.5.1 Stability is sufficient, not necessary

The most important limitation is that the implication in Proposition 8.2(a) is one-directional: strong features imply high stability, but high stability does not guarantee that the features are informative for the specific task of state discovery in SCX. It is possible for k-means clustering to be stable (the same clusters appear across bootstrap re-samples) while those clusters are not aligned with the true latent states S that determine expert error rates. This can occur when:

- The data contain a strong spurious cluster structure (e.g., a categorical covariate with large effect) that dominates the feature representation, while the true error-determining states are a subtle substructure within each spurious cluster.
- The number of clusters K is misspecified. Choosing K too small forces k-means to merge distinct states, potentially producing stable but poor partitions; choosing K too large splits states arbitrarily, producing unstable partitions even when the features are strong.
- The clusters are non-convex in the feature space (e.g., crescent-shaped or nested clusters), which k-means cannot recover regardless of feature strength.

For this reason, we recommend using the bootstrap stability diagnostic as a *necessary condition* for strong features, not a sufficient one. If $S \leq 0.7$, the features are demonstrably weak and SCX should not be applied. If $S > 0.7$, the features *may* be sufficient, but the final test remains the empirical F1 improvement once the SCX pipeline is executed.

8.5.2 Computational cost

Algorithm 1 requires $B \times R$ k-means runs (default $50 \times 20 = 1000$ runs). For large datasets ($N > 10^5$), this can be computationally expensive. The following optimizations are recommended:

- **Minibatch k-means** [?] can replace full k-means for large N, reducing per-run cost by an order of magnitude with minimal impact on stability estimates. Our experiments show that the correlation between S computed with full k-means and with minibatch k-means (batch size $= 0.1N$) exceeds 0.95 for $N > 10^4$.
- **Subsampling**: For very large datasets, the bootstrap procedure can be applied to a random subset of data (e.g., 10^4 samples), provided the subset preserves the cluster proportions.
- **Early stopping**: If the ARI values across the first $B_0 = 20$ bootstraps are all above 0.85, the stability is almost certainly above 0.7, and the procedure can terminate early.

8.5.3 Threshold is a heuristic

The 0.7 threshold, while grounded in the Landis–Koch convention, is ultimately a heuristic. The relationship between $S(\Phi, K)$ and δ is governed by problem-specific factors—the geometry of the clusters, the dimensionality d, the number of clusters K, and the presence of noise variables in the feature space—that cannot be captured by a single universal threshold. Users working in high-stakes domains (e.g., medical imaging) should calibrate the threshold for their specific data using the procedure in Definition 8.3.

Practical guidance summary.

1. **Before running SCX**: Compute $S(\Phi, K)$ via Algorithm 1.
2. **If** $S > 0.7$: Features are likely sufficient. Proceed with the full SCX pipeline (state discovery via error-driven encoding, multi-expert training, and consistency testing). The bound in Eq. (3) will not be the bottleneck.
3. **If** $0.4 < S \leq 0.7$: Features are marginal. Consider feature engineering before proceeding. For vision tasks, try a deeper backbone network; for material science, use higher-body-order descriptors. If no feature improvement is feasible, proceed with SCX but treat the results as exploratory.
4. **If** $S \leq 0.4$: Stop. SCX will not outperform a simple loss-threshold baseline. Direct efforts to feature engineering or accept that data quality assessment is limited by the available representations.
5. **After running SCX**: Verify the weak feature diagnosis by comparing $\mathrm{F1}_{\mathrm{SCX}}$ to $\mathrm{F1}_{\mathrm{base}}$. A dramatically different result from the stability prediction suggests model misspecification or a violation of Assumptions A1–A6.

8.5.4 Connection to Theorem 2 is empirical, not formal

We emphasize that the relationship between $S(\Phi, K)$ and δ in Proposition 8.2(c) is qualitatively informative but not quantitatively precise. The inequality

$$\frac{\delta}{\log K} \gtrsim \frac{1}{2}\big(1 - S(\Phi, K)\big)$$

is a heuristic approximation, not a proven bound. The precise relationship depends on the geometry of the clusters, the dimensionality, and the specific K used in clustering. We provide it as guidance for practitioners, not as a rigorous mathematical statement. The formal content of Proposition 8.2 is contained in parts (a) and (b); part (c) is a practical translation whose validity should be verified empirically for each application domain.

8.5.5 Choice of K

The bootstrap stability diagnostic requires specifying the number of clusters K, yet K is itself unknown in unsupervised settings. We recommend estimating K via the gap statistic [?] with $K_{\max} = \min(\sqrt{N}, 20)$, then computing $S(\Phi, K)$ for $K \in \{K_{\mathrm{gap}} - 1, K_{\mathrm{gap}}, K_{\mathrm{gap}} + 1\}$. If the stability is above 0.7 for all three values of K, the diagnostic is robust to the K choice. If the stability varies substantially across K, this itself is evidence of weak features.

8.5.6 When to use this diagnostic

The bootstrap stability diagnostic is most valuable in two scenarios:

1. **Early stopping**: Before committing tens of GPU-hours to training multiple experts, a quick stability check (requiring only k-means on features, no model training)

can identify representational weaknesses. In our experiments, computing $S(\Phi, K)$ on a 10^5-sample dataset with $d = 512$ and $B = 50$ takes approximately 90 seconds on a single CPU core.

2. **Feature selection**: When multiple feature representations are available (e.g., ACE descriptors at different body orders, or CNN features from different layers), computing $S(\Phi, K)$ for each candidate provides a rapid ranking of representational quality without training any models.

In summary, Proposition 8.2 provides a practical, computationally efficient diagnostic for feature strength that directly tests the clustering procedure used in SCX state discovery. While it has limitations—sufficiency without necessity, heuristic thresholds, and informal connection to δ—it empirically identifies the feature weakness regime where SCX fails to improve over baselines, providing actionable guidance before expensive multi-expert training begins.

9 Experimental Details

This section documents the experimental configuration for all datasets used in the main text. All experiments are designed to validate the theoretical predictions of Theorems 1–4, Proposition 6, and the SCX noise detection framework under controlled conditions.

9.1 AlN v3 MLIP: Dataset, ACE Descriptors, Expert Configuration

Dataset. The AlN v3 dataset consists of 534 atomic configurations (frames) generated from thermal and molecular dynamics simulations of aluminium nitride using density functional theory (DFT). Each frame contains atomic positions, forces, and energies computed at the DFT level. Among these 534 frames, 53 carry label noise introduced by thermal fluctuations and MD trajectory artifacts, giving a noise rate $\eta = 53/534 \approx 0.099$. The noise is concentrated in frames near the melting transition where the Born–Oppenheimer surface is poorly sampled.

ACE Descriptors. Feature representations are computed using the Atomic Cluster Expansion (ACE) [1], a body-order expansion of atomic neighbourhood densities:

$$\phi_i = \left\{ \langle A_i | A_{n_1 n_2 l m} \rangle \right\}, \tag{48}$$

where each feature corresponds to a correlation of radial basis functions and spherical harmonics evaluated on the atomic neighbourhood of atom i. The ACE descriptor yields $d_\phi = 100$ features per atom (using $N_{\text{radial}} = 8$, $l_{\max} = 4$, and a 3-body correlation order). Per-configuration features are obtained by averaging atom-wise ACE vectors over all atoms in the frame.

State Discovery. K-means clustering ($K = 8$ states, determined by the elbow method on the within-cluster sum of squares) is applied to the ACE feature matrix $\Phi \in \mathbb{R}^{534 \times 100}$. The discovered states correspond to physically interpretable regimes: bulk crystal, surface, near-melting, defect-rich, and several intermediate configurations.

Expert Configuration. $M = 12$ expert models are trained, each on a disjoint subset of the 534 frames. Following the SCX pipeline:

- Each expert is a linear ACE model (Ridge regression on the ACE descriptors to predict per-atom energies and forces).
- Training subsets are disjoint and stratified by state to preserve state proportions: each expert sees approximately $534/12 \approx 44$ frames.
- Experts are trained independently and in parallel using the CPU cluster described in Section 9.7.
- The consensus score $C(x)$ is the fraction of experts whose per-frame energy prediction error exceeds a threshold $\tau = 0.05$ eV/atom.

Results. The empirical F1 score is 0.87. The theoretical F1 lower bound from Theorem 1 (Hoeffding form) is 0.87 for $M = 12$, $\eta = 0.099$, and a mean separation gap $\Delta \approx 0.35$ (estimated from the empirical consensus distributions). The adaptive threshold from Theorem 4' with η estimated from the data yields $C_{\min}/\eta = 4.59$ (numerically verified), giving the asymptotic constant for the Chernoff-rate bound.

9.2 CIFAR-10: ResNet-18 Architecture, Training Protocol

Dataset. CIFAR-10 consists of 60,000 32×32 colour images in 10 classes (50,000 training, 10,000 test). We inject synthetic label noise at rate $\eta \in \{0.05, 0.10, 0.20\}$ for controlled experiments; the main text reports $\eta = 0.10$.

ResNet-18 Architecture. The expert model is a standard ResNet-18 [2] adapted to CIFAR-10:

- Initial 3×3 convolution (stride 1) with 64 filters, batch normalisation, ReLU.
- Four residual stages with $[2, 2, 2, 2]$ blocks, feature dimensions $[64, 128, 256, 512]$.
- Global average pooling followed by a 10-way fully connected classification head.
- Approximately 11 million trainable parameters.

Training Protocol.

- $M = 10$ experts, each trained on a disjoint subset of 5,000 images (stratified by class).
- Optimiser: SGD with momentum 0.9, weight decay 5×10^{-4}.
- Learning rate schedule: initial 0.1, cosine annealing to zero over 200 epochs.
- Batch size: 128.
- Data augmentation: random horizontal flip, random 32×32 crop with 4-pixel padding.
- Each expert training takes approximately 30 minutes on a single NVIDIA RTX 3090 GPU.

Feature Extraction for State Discovery. State discovery uses the penultimate-layer embeddings (dimension 512) from a ResNet-18 trained on clean CIFAR-10. These embeddings serve as the feature representation $\phi(x) \in \mathbb{R}^{512}$. K-means clustering with $K = 10$ is applied to the embedding matrix.

Results. At $\eta = 0.10$ noise, SCX achieves F1 $= 0.62$ versus the loss-baseline detector F1 $= 0.45$. The mutual information ratio $\delta/\log K \approx 0.15$, placing CIFAR-10 well within the regime where Theorem 2 predicts SCX can improve over the baseline. The empirical improvement of $+0.17$ is consistent with the bound $\mathrm{F1}_{\mathrm{SCX}} \leq \mathrm{F1}_{\mathrm{base}} + C_F\sqrt{\delta/2}$ with $C_F \approx 1.1$ and $\sqrt{\delta/2} \approx 0.27$.

9.3 DermaMNIST: SimpleCNN vs ResNet-18 Comparison

Dataset. DermaMNIST [3] is a skin lesion classification dataset derived from the HAM10000 repository. It contains 10,015 dermoscopic images across 7 disease classes (train/val/test split: 7,007 / 1,003 / 2,005). We inject label noise at rate $\eta = 0.10$, matching the CIFAR-10 protocol.

SimpleCNN Architecture (Weak Features). To create a deliberately weak feature regime, we use a minimal CNN:

- Two convolutional layers: 3×3 conv, 32 filters, ReLU, 2×2 max-pooling; then 3×3 conv, 64 filters, ReLU, 2×2 max-pooling.
- Two fully connected layers: 128 units (ReLU), then 7 units (classification).
- Approximately 0.4 million parameters.
- Feature dimension before classification head: 128.

ResNet-18 Configuration (Strong Features, comparison). We also evaluate a ResNet-18 on DermaMNIST (same architecture as Section **??** but with 7 output classes) to demonstrate the effect of feature quality. Input images are resized to 224×224 for the ResNet-18 (as required by the pretrained stem), or kept at 28×28 for SimpleCNN.

Training Protocol (both architectures).

- $M = 10$ experts, disjoint subsets stratified by class.
- SimpleCNN: Adam optimiser, learning rate 10^{-3}, 100 epochs, batch size 64.
- ResNet-18: SGD with momentum 0.9, learning rate 10^{-2} (cosine annealing), 100 epochs, batch size 64.
- Data augmentation: random horizontal flip, random rotation ($\pm 15°$).

Results.

- **SimpleCNN (weak features)**: $\delta/\log K \approx 0.52$. SCX F1 = 0.101, loss baseline F1 = 0.105. The improvement is negligible, consistent with Theorem 2's prediction that SCX cannot significantly outperform the baseline when $\delta/\log K > 0.5$. Proposition 6's bootstrap stability diagnostic gives $S(\Phi, K) < 0.5$, confirming weak features.
- **ResNet-18 (stronger features)**: $\delta/\log K \approx 0.22$. SCX F1 = 0.48, loss baseline F1 = 0.38. The improvement of +0.10 confirms that better features restore SCX's effectiveness.

This comparison provides a direct empirical test of Theorem 2: with everything else held constant (dataset, noise rate, M), varying only the feature quality transitions SCX from ineffective (SimpleCNN, $\delta/\log K > 0.5$) to effective (ResNet-18, $\delta/\log K < 0.3$).

9.4 DrugBank: Molecular Descriptors, Targetome Screening

Dataset. DrugBank [4] is a comprehensive pharmaceutical database containing 13,557 drug entries with 5,108 approved small molecule drugs. We use the subset of 2,000 approved drugs with known protein target annotations for a drug–target interaction noise detection task.

Molecular Descriptors. Feature representations are computed using three descriptor families:

- **MACCS keys**: 166-bit substructure fingerprints (binary).
- **ECFP4 (Extended Connectivity Fingerprints)**: 2,048-bit circular fingerprints with radius 2.
- **Physicochemical descriptors**: 200 continuous-valued descriptors (molecular weight, logP, H-bond donors/acceptors, rotatable bonds, polar surface area, etc.).

The full feature vector is the concatenation of all three types: $\phi(x) \in \mathbb{R}^{2414}$ (166 + 2048 + 200).

Targetome Screening Task. For each drug, the task is to predict its set of protein targets (multi-label classification over 500 target classes). Noise is injected by flipping a fraction η of the drug–target interaction labels. Expert models are $M = 8$ random forest classifiers (100 trees each, max depth 20), trained on disjoint subsets of the 2,000 drugs stratified by therapeutic area.

Results. The DrugBank dataset exhibits moderate feature strength ($\delta/\log K \approx 0.31$), placing it in the regime where SCX provides modest but measurable improvement. SCX F1 = 0.34 versus loss baseline F1 = 0.28 at $\eta = 0.10$ noise.

9.5 Noise Injection Protocol (All Datasets)

All synthetic noise injections follow a uniform-label-flip protocol (consistent with Assumption A4):

1. Select a fraction η of training samples uniformly at random.
2. For each selected sample: replace the true label y^* with a label drawn uniformly from $\mathcal{Y} \setminus \{y^*\}$.
3. For multi-state datasets (AlN), noise injection is stratified so that each state s receives noise at rate η (preserving ρ_s).
4. The noise injection is applied *before* expert training, so experts see noisy labels in their training subsets.
5. The ground truth (clean labels) is retained for evaluation but never revealed to the SCX pipeline.

For AlN v3, the noise is organic (not synthetically injected), arising from thermal and MD simulation artifacts. The noise detection task is to identify these naturally occurring noisy frames.

9.6 Evaluation Metrics

Precision, Recall, F1 Score. Let $\hat{Z}(x) \in \{0, 1\}$ be the SCX noise flag for sample x, and $Z(x) \in \{0, 1\}$ be the ground truth (1 = noise). On a per-state basis:

$$\text{Precision}_s = \frac{|\{x \in s : \hat{Z} = 1, Z = 1\}|}{|\{x \in s : \hat{Z} = 1\}|}, \tag{49}$$

$$\text{Recall}_s = \frac{|\{x \in s : \hat{Z} = 1, Z = 1\}|}{|\{x \in s : Z = 1\}|}, \tag{50}$$

$$F1_s = \frac{2 \cdot \text{Precision}_s \cdot \text{Recall}_s}{\text{Precision}_s + \text{Recall}_s}. \tag{51}$$

The global F1 is the weighted average over states: $\text{F1} = \sum_s \rho_s \cdot F1_s$, where ρ_s is the proportion of data in state s.

Loss Baseline. The loss baseline detector flags a sample as noise if the minimum expert loss $\min_m \ell(f_m(x), y)$ exceeds a threshold θ_{base}, without using state information or feature representations. The threshold θ_{base} is chosen to maximise F1 on a held-out validation set (10% of training data). This baseline represents the best achievable performance of a method that does not exploit state structure.

Stability Score (Proposition 6). The clustering stability $S(\Phi, K)$ is the mean adjusted Rand index (ARI) between k-means clusterings on the full data and on $B = 50$ bootstrap resamples, using k-means++ initialization and $n_{\text{init}} = 10$ random starts per run. See Supplementary Information S6 for the algorithm.

Table 4: Computational resources and runtimes for all experiments.

Dataset	Hardware	Expert Training
AlN v3	32-core CPU cluster (Intel Xeon Platinum 8582C)	∼15 min per expe
CIFAR-10	1 × NVIDIA RTX 3090 (24 GB)	∼30 min per expe
DermaMNIST (SimpleCNN)	1 × NVIDIA RTX 3090	∼5 min per exper
DermaMNIST (ResNet-18)	1 × NVIDIA RTX 3090	∼20 min per expe
DrugBank	8-core CPU workstation	∼2 min per exper

9.7 Computational Resources

Table 4 summarizes the computational resources used for each dataset.

All experiments use $M \in \{8, 10, 12\}$ experts depending on the dataset. The bootstrap stability diagnostic (Proposition 6) adds $B = 50$ additional k-means runs per dataset, which is negligible compared to expert training time. The SCX pipeline is implemented in Python 3.10 using PyTorch 2.0 (GPU experiments) and scikit-learn 1.2 (CPU experiments).

References

[1] R. Drautz, "Atomic cluster expansion for accurate and transferable interatomic potentials," *Phys. Rev. B* **99**, 014104 (2019).

[2] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE CVPR*, pp. 770–778 (2016).

[3] J. Yang *et al.*, "MedMNIST v2: A large-scale lightweight benchmark for 2D and 3D biomedical image classification," *Sci. Data* **10**, 41 (2023).

[4] D. S. Wishart *et al.*, "DrugBank 5.0: A major update to the DrugBank database for 2018," *Nucleic Acids Res.* **46**, D1074–D1082 (2018).

10 Numerical Verification of Theoretical Constants

This section provides numerical verification of the exact constant minimax optimality result (Theorem 4' in Supplementary Information S4). All computations are performed using the Python script `numerical_verify.py` with standard double-precision arithmetic.

10.1 Verification Protocol

For a given parameter set (p_0, p_1, η), we compute the following theoretical quantities:

1. **Chernoff point** θ^*: the unique solution to $\mathrm{KL}(\theta \| p_0) = \mathrm{KL}(\theta \| p_1)$,
2. **Chernoff information** $\kappa = \mathrm{KL}(\theta^* \| p_0) = \mathrm{KL}(\theta^* \| p_1)$,
3. **Saddlepoints** $\lambda_0^* > 0$, $\lambda_1^* < 0$ defined by the log-odds at θ^*,
4. **Total log-odds** $D^* = \lambda_0^* + |\lambda_1^*|$ and exponent fraction $s = |\lambda_1^*|/D^*$,

5. **Canonical minimax constant** $C_{\min}$ (Lemma E form):

$$C_{\min} = \frac{\eta}{2}\left(\frac{1-\eta}{\eta}\right)^s \frac{1/\lambda_0^* + 1/|\lambda_1^*|}{\sqrt{\theta^*(1-\theta^*)}}, \tag{52}$$

6. **Non-adaptive limit** K_{na} (Lemma B, using θ^* without η adaptation):

$$K_{\text{na}} = \left(\frac{1}{2|\lambda_1^*|} + \frac{1-\eta}{2\eta\lambda_0^*}\right)\frac{1}{\sqrt{\theta^*(1-\theta^*)}}, \tag{53}$$

7. **Adaptive limit** K_{ad} (Lemma D, using η-aware threshold $\theta^\dagger$):

$$K_{\text{ad}} = \left(\frac{1-\eta}{\eta}\right)^s \frac{1/\lambda_0^* + 1/|\lambda_1^*|}{2\sqrt{\theta^*(1-\theta^*)}}. \tag{54}$$

The verification criterion is the machine-precision equality $K_{\text{ad}} = C_{\min}/\eta$, which confirms that the adaptive threshold achieves the minimax lower bound exactly (up to the asymptotic limit).

10.2 Test Case 1: Standard Parameters

Parameters: $p_0 = 0.10$, $p_1 = 0.60$, $\eta = 0.10$. This is the canonical case studied in the main text (AlN v3 MLIP regime), with moderate separation and moderately rare noise.

Table 5: Geometric constants and verification for Test Case 1.

Quantity	Symbol	Value
Clean error rate	p_0	0.10
Noise error rate	p_1	0.60
Noise rate	η	0.10
Chernoff point	θ^*	0.311574
Chernoff information	κ	0.169604
Positive saddlepoint	λ_0^*	1.404453
Negative saddlepoint	$\|\lambda_1^*\|$	1.198236
Total log-odds	D^*	2.602690
Exponent fraction	s	0.460384
Standard deviation factor	$\sqrt{\theta^*(1-\theta^*)}$	0.463137
Canonical $C_{\min}$	$C_{\min}$	0.459147
Normalised constant	$C_{\min}/\eta$	4.591465
Non-adaptive limit	K_{na}	7.819233
Adaptive limit	K_{ad}	4.591465
Ratio $K_{\text{na}}/(C_{\min}/\eta)$		1.7030

- **Adaptive limit** $= C_{\min}/\eta$: `PASS` (difference $< 10^{-15}$).
- **Chernoff** κ **vs Hoeffding** $2\Delta^2$: $\kappa = 0.1696$, $2\Delta^2 = 0.5000$, ratio $= 2.95$. The Chernoff rate is substantially slower (more honest) than the Hoeffding bound.
- **Non-adaptive penalty**: The naive threshold θ^* (which ignores η) is $1.70\times$ worse than the adaptive threshold.

10.3 Test Case 2: Weak Separation

Parameters: $p_0 = 0.20$, $p_1 = 0.50$, $\eta = 0.30$. This case has a smaller separation gap ($\Delta = 0.30$ versus $\Delta = 0.50$ in Case 1) and moderate noise.

Table 6: Geometric constants and verification for Test Case 2.

Quantity	Symbol	Value
Clean error rate	p_0	0.20
Noise error rate	p_1	0.50
Noise rate	η	0.30
Chernoff point	θ^*	0.339036
Chernoff information	κ	0.052753
Positive saddlepoint	λ_0^*	0.718701
Negative saddlepoint	$\lvert\lambda_1^*\rvert$	0.667593
Total log-odds	D^*	1.386294
Exponent fraction	s	0.481567
Standard deviation factor	$\sqrt{\theta^*(1-\theta^*)}$	0.473382
Canonical $C_{\min}$	$C_{\min}$	1.376829
Normalised constant	$C_{\min}/\eta$	4.589431
Non-adaptive limit	K_{na}	5.011297
Adaptive limit	K_{ad}	4.589431
Ratio $K_{\text{na}}/(C_{\min}/\eta)$		1.0919

- **Adaptive limit** $= C_{\min}/\eta$: `PASS` (difference $< 10^{-15}$).
- **Chernoff κ vs Hoeffding** $2\Delta^2$: $\kappa = 0.0528$, $2\Delta^2 = 0.1800$, ratio $= 3.41$. The Hoeffding bound is $3.4\times$ more optimistic than the Chernoff rate.
- **Non-adaptive penalty**: Only $1.09\times$ worse — for η near $1/2$, the adaptive gain is small (confirming the theory).

10.4 Test Case 3: Strong Separation

Parameters: $p_0 = 0.05$, $p_1 = 0.80$, $\eta = 0.05$. This case has a large separation gap ($\Delta = 0.75$), rare noise, and very low clean error.

- **Adaptive limit** $= C_{\min}/\eta$: `PASS` (difference $< 5 \times 10^{-16}$).
- **Chernoff κ vs Hoeffding** $2\Delta^2$: $\kappa = 0.4574$, $2\Delta^2 = 1.1250$, ratio $= 2.46$. The Chernoff rate is the fastest among all cases (largest κ), reflecting the strong separation.
- **Non-adaptive penalty**: The naive threshold is $2.41\times$ worse — the largest penalty of all cases, because $\eta = 0.05$ is far from $1/2$, making the η-aware threshold critical.

Table 7: Geometric constants and verification for Test Case 3.

Quantity	Symbol	Value
Clean error rate	p_0	0.05
Noise error rate	p_1	0.80
Noise rate	η	0.05
Chernoff point	θ^*	0.359788
Chernoff information	κ	0.457370
Positive saddlepoint	λ_0^*	2.368153
Negative saddlepoint	$\lvert\lambda_1^*\rvert$	1.962580
Total log-odds	D^*	4.330733
Exponent fraction	s	0.453175
Standard deviation factor	$\sqrt{\theta^*(1-\theta^*)}$	0.479938
Canonical $C_{\min}$	$C_{\min}$	0.184322
Normalised constant	$C_{\min}/\eta$	3.686449
Non-adaptive limit	K_{na}	8.889338
Adaptive limit	K_{ad}	3.686449
Ratio $K_{\text{na}}/(C_{\min}/\eta)$		2.4114

Table 8: Geometric constants and verification for Test Case 4.

Quantity	Symbol	Value
Clean error rate	p_0	0.10
Noise error rate	p_1	0.60
Noise rate	η	0.50
Exponent fraction	s	0.460384
Canonical $C_{\min}$	$C_{\min}$	0.834840
Normalised constant	$C_{\min}/\eta$	1.669681
Non-adaptive limit	K_{na}	1.669681
Adaptive limit	K_{ad}	1.669681
Ratio $K_{\text{na}}/(C_{\min}/\eta)$		1.0000

10.5 Test Case 4: Symmetric Noise

Parameters: $p_0 = 0.10$, $p_1 = 0.60$, $\eta = 0.50$. The geometric constants are identical to Case 1 (same p_0, p_1); only η changes. At $\eta = 0.50$, the noise is symmetric (noise prior equals the clean prior).

- **Adaptive limit** $= C_{\min}/\eta$: `PASS` (difference $< 10^{-15}$).
- **Adaptive = non-adaptive**: At $\eta = 1/2$, the correction term $\frac{1}{M}\frac{\log((1-\eta)/\eta)}{D^*} = 0$, so $\theta^\dagger = \theta^*$. The two limits coincide exactly.
- **Theoretical prediction confirmed**: Lemma D states that the $O(1/M)$ threshold shift vanishes when $\eta = 1/2$. Numerical verification confirms this.

10.6 Test Case 5: Dominant Noise

Parameters: $p_0 = 0.10$, $p_1 = 0.60$, $\eta = 0.90$. This is the reciprocal of Case 1: noise dominates the dataset. The geometric constants are identical to Case 1 and Case 4 (same p_0, p_1).

Table 9: Geometric constants and verification for Test Case 5.

Quantity	Symbol	Value
Clean error rate	p_0	0.10
Noise error rate	p_1	0.60
Noise rate	η	0.90
Exponent fraction	s	0.460384
Canonical $C_{\min}$	$C_{\min}$	0.546460
Normalised constant	$C_{\min}/\eta$	0.607177
Non-adaptive limit	K_{na}	0.986397
Adaptive limit	K_{ad}	0.607177
Ratio $K_{\text{na}}/(C_{\min}/\eta)$		1.6246

- **Adaptive limit** $= C_{\min}/\eta$: `PASS` (difference $< 10^{-15}$).
- **Non-adaptive penalty**: The naive threshold is 1.62× worse. The penalty is slightly smaller than in Case 1 (1.70×), reflecting the asymmetry of the $((1-\eta)/\eta)^s$ factor: the penalty for $\eta = 0.90$ is the reciprocal of the penalty for $\eta = 0.10$ under the mapping $\eta \mapsto 1 - \eta$, and the constants account for this via the s exponent.

10.7 Finite-M Convergence Analysis

The Bahadur–Rao asymptotic expansion predicts that for finite M,

$$e^{M\kappa}\sqrt{2\pi M}\big(1 - \mathrm{F1}(M)\big) \longrightarrow \frac{C_{\min}}{\eta} \quad \text{as } M \to \infty. \tag{55}$$

Table 10 shows the finite-M convergence for Case 2 (the only case where $M = 500$ produces non-zero $1 - \mathrm{F1}$ within double precision; for the other cases, the error probabilities are below machine epsilon at $M = 500$).

The convergence is monotonic and rapid: at $M = 100$, the normalized constant is within 0.12% of the asymptotic limit; at $M = 500$, it matches to 0.02%.

Table 10: Finite-M convergence of the normalized error constant for Test Case 2 ($p_0 = 0.20$, $p_1 = 0.50$, $\eta = 0.30$).

M	$1 - \mathrm{F1}$	$e^{M\kappa}\sqrt{2\pi M}(1 - \mathrm{F1})$	Ratio to K_{ad}
50	1.52×10^{-3}	4.5721	0.9962
100	1.80×10^{-5}	4.5837	0.9988
200	2.54×10^{-9}	4.5881	0.9997
500	4.16×10^{-13}	4.5902	1.0002
∞	—	$K_{\mathrm{ad}} = 4.5894$	1.0000

10.8 Summary Table: All Cases

Table 11 collects the key verification metrics across all five test cases. The central result is that the adaptive limit K_{ad} matches $C_{\min}/\eta$ to machine precision in every case, confirming that the SCX adaptive threshold achieves the minimax lower bound exactly.

Table 11: Summary of numerical verification across all five test cases.

Quantity	Case 1	Case 2	Case 3	Case 4	Case 5
p_0	0.10	0.20	0.05	0.10	0.10
p_1	0.60	0.50	0.80	0.60	0.60
η	0.10	0.30	0.05	0.50	0.90
$\Delta = p_1 - p_0$	0.50	0.30	0.75	0.50	0.50
θ^*	0.311574	0.339036	0.359788	0.311574	0.311574
κ	0.169604	0.052753	0.457370	0.169604	0.169604
$2\Delta^2/\kappa$	2.9480	3.4121	2.4597	2.9480	2.9480
s	0.460384	0.481567	0.453175	0.460384	0.460384
$C_{\min}$	0.459147	1.376829	0.184322	0.834840	0.546460
$C_{\min}/\eta$	4.591465	4.589431	3.686449	1.669681	0.607177
K_{na}	7.819233	5.011297	8.889338	1.669681	0.986397
K_{ad}	4.591465	4.589431	3.686449	1.669681	0.607177
$K_{\mathrm{ad}} \stackrel{?}{=} C_{\min}/\eta$	**PASS**	**PASS**	**PASS**	**PASS**	**PASS**
Non-adaptive penalty	1.70×	1.09×	2.41×	1.00×	1.62×

Key findings.

1. **Machine-precision equality** ($K_{\mathrm{ad}} = C_{\min}/\eta$) holds across all five cases, confirming the algebraic equivalence of the adaptive achievability limit and the minimax lower bound.

2. **Chernoff information is always slower than the Hoeffding bound**: $2\Delta^2/\kappa > 1$ in all cases (range 2.46–3.41). This confirms that the Hoeffding bound overestimates the rate of convergence.

3. **Non-adaptive penalty is case-dependent**: The naive threshold θ^* (ignoring η) is suboptimal by factors 1.00–2.41. The penalty is largest when η is far from 1/2 (Cases 1, 3, 5) and disappears at $\eta = 1/2$ (Case 4).

4. **All constants satisfy theoretical consistency checks**: $\mathrm{KL}(\theta^*\|p_0) = \mathrm{KL}(\theta^*\|p_1)$ holds to $\pm 10^{-16}$; $C_{\min} > 0$ in all cases; $K_{\mathrm{ad}} > 0$; the exponent fraction $s \in (0, 1)$.