

The Recursive Audit Theorem: Self-Reference of the Audit Sword, Convergence under $\rho < 1$, and the Fixed-Point Identity with Theorem 3

Anonymous Author(s)

June 29, 2026

Abstract

The SCX “Audit Sword” principle asserts that every use of the Ceric Score must be independently auditable. But this sword can — and must — be turned upon itself. If auditor $\mathcal{A}_1$ issues a noise estimate $\hat{\eta}_1$, a meta-auditor $\mathcal{A}_2$ can audit $\mathcal{A}_1$ ’s output, producing $\hat{\eta}_2$, and so on, ad infinitum. We formalize this **Recursive Audit Problem** and prove four results. (i) **Recursive Convergence** (Theorem 3.1): if each meta-audit layer’s additional noise σ_k^2 satisfies $\sigma_k^2 \leq \rho \sigma_{k-1}^2$ with $\rho < 1$ (the meta-auditor is strictly higher-fidelity than its object), the recursive sequence $\{\hat{\eta}_k\}$ converges almost surely to a fixed point η^* satisfying $\eta^* = f(\eta^*)$, with exponential rate when $\rho \leq 1/2$. This is **rigorous** from Banach’s fixed-point theorem and Doob’s martingale convergence theorem. (ii) **Infinite Regress Inevitability** (Theorem 3.3): if $\rho \geq 1$ (meta-auditor no more precise than the object auditor), the probability of divergence tends to 1 — the Audit Sword *breaks* under self-reference. **[Rigorous]** (Pólya’s theorem). (iii) **Fixed-Point Information-Theoretic Characterization** (Conjecture ??): η^* is conjectured to equal the audit heat death $H_{\min} = H(\varepsilon \mid S)$ from AE-Theorem — the convergence point of recursive auditing is exactly Theorem 3’s information barrier. **[Open]** (equivalence proof). (iv) **Gödel Analogy** (Open Problem ??): the $\rho \geq 1$ divergence may correspond to an **absolutely unauditable audit proposition** — a statement about audit quality that cannot be verified by any finite iteration of the Audit Sword. **[Open]**

1 Introduction

Self-reference is the crucible in which foundational claims are tested. Gödel’s incompleteness theorems showed that any sufficiently powerful formal system contains true statements it cannot prove. Turing’s halting problem showed that no algorithm can decide whether an arbitrary program terminates. The SCX framework makes an equally universal claim — the Audit Sword: every use of the Cercis Score can be independently audited — and must therefore face the same self-referential test.

The recursive audit problem is this: **Can the audit of an audit be audited? And if so, does this recursion converge, or does it spiral into infinite regress?**

Concretely:

- (1) Auditor $\mathcal{A}_0$ examines dataset $\mathcal{D}$ and outputs an estimated noise rate $\hat{\eta}_0$ with confidence interval $[\hat{\eta}_0 - \epsilon_0, \hat{\eta}_0 + \epsilon_0]$;
- (2) Meta-auditor $\mathcal{A}_1$ examines $\mathcal{A}_0$ ’s output (and the underlying data) and produces $\hat{\eta}_1$ with confidence $[\hat{\eta}_1 - \epsilon_1, \hat{\eta}_1 + \epsilon_1]$;
- (3) Meta-meta-auditor $\mathcal{A}_2$ examines $\mathcal{A}_1$ ’s output, producing $\hat{\eta}_2$;
- (4) ...
- (5) $\mathcal{A}_k$ produces $\hat{\eta}_k$.

The question is whether $\hat{\eta}_k$ converges as $k \rightarrow \infty$, and if so, to what value and under what conditions.

Contributions.

- (1) **Recursive Convergence** (Theorem 3.1): $\rho < 1 \Rightarrow \hat{\eta}_k \xrightarrow{a.s.} \eta^*$. [Rigorous]
- (2) **Infinite Regress** (Theorem 3.3): $\rho \geq 1 \Rightarrow P(\text{convergence}) \rightarrow 0$. [Rigorous]
- (3) **Fixed-Point Identity Conjecture** (Conjecture ??): $\eta^* = H_{\min}$. [Open]
- (4) **Gödel Analogy** (Open Problem ??): absolutely unauditable propositions. [Open]

2 Preliminaries

2.1 Recursive Audit Sequence

Definition 2.1 (Recursive Audit Sequence). *Let $\mathcal{A}_0$ be the base auditor operating on dataset $\mathcal{D}$. For $k \geq 1$, the k -th meta-auditor $\mathcal{A}_k$ takes as input:*

- (i) *The original dataset $\mathcal{D}$ (or its Cercis fingerprint);*
- (ii) *The audit trail $\{\hat{\eta}_j, \epsilon_j\}_{j=0}^{k-1}$ of all prior levels;*
- (iii) *Its own independent audit analysis.*

$\mathcal{A}_k$ outputs $\hat{\eta}_k$ with confidence half-width ϵ_k . The recursive sequence is $\{\hat{\eta}_k\}_{k=0}^\infty$.

Definition 2.2 (Meta-Audit Fidelity Decay Rate). *The **fidelity decay rate** ρ is the smallest constant such that the additional noise introduced by meta-auditor $\mathcal{A}_{k+1}$ (beyond that of $\mathcal{A}_k$) satisfies:*

$$\sigma_{k+1}^2 \leq \rho \cdot \sigma_k^2, \quad (1)$$

where $\sigma_k^2 = \mathbb{V}[\hat{\eta}_k \mid \eta^*]$ is the variance of the k -th level estimate around the true fixed point.

Intuitively, $\rho < 1$ means each meta-auditor is strictly more precise (higher-fidelity) than the auditor it audits. $\rho = 1$ means equal fidelity (same system auditing itself). $\rho > 1$ means the meta-auditor is less precise than its object.

2.2 SCX Recursive Audit Operator

Definition 2.3 (SCX Recursive Audit Operator). *The recursion is governed by the operator $f : [0, 1] \rightarrow [0, 1]$:*

$$\hat{\eta}_{k+1} = f(\hat{\eta}_k) + \xi_k, \quad (2)$$

where $\xi_k \sim (0, \sigma_k^2)$ is the meta-audit estimation noise and f encodes the systematic relationship between an audit estimate at level k and the corrected estimate produced by level $k + 1$.

3 Main Results

3.1 Recursive Convergence

Theorem 3.1 (Recursive Audit Convergence). *Assume the SCX recursive audit operator f is a contraction on $[0, 1]$: $|f(\eta) - f(\eta')| \leq \rho_f |\eta - \eta'|$ with*

$\rho_f < 1$, and that the meta-audit noise satisfies $\sigma_k^2 \leq \rho_\sigma \sigma_{k-1}^2$ with $\rho_\sigma < 1$. Let $\rho = \max(\rho_f, \rho_\sigma) < 1$. Then:

(i) $\hat{\eta}_k \xrightarrow{a.s.} \eta^*$ as $k \rightarrow \infty$, where $\eta^* = f(\eta^*)$ is the unique fixed point;

(ii) If $\rho \leq 1/2$, the convergence is exponential:

$$|\hat{\eta}_k - \eta^*| \leq \sigma_0 \cdot \rho^{k/2} \cdot O(\sqrt{\log k}), \quad (3)$$

holding with probability $\geq 1 - 1/k^2$ for all sufficiently large k .

Proof. Part (i): Contraction + Martingale. By Banach’s fixed-point theorem, a contraction on a complete metric space has a unique fixed point. The deterministic iteration $\eta_{k+1} = f(\eta_k)$ converges geometrically: $|\eta_k - \eta^*| \leq \rho_f^k |\eta_0 - \eta^*|$.

Adding the stochastic component, define the deviation $D_k = \hat{\eta}_k - \eta^*$. Under the noise decay condition, the sequence $\{D_k\}$ is a supermartingale: $\mathbb{E}[|D_{k+1}| \mid \mathcal{F}_k] \leq \rho |D_k|$. Since $|D_k|$ is non-negative, Doob’s martingale convergence theorem guarantees $D_k \xrightarrow{a.s.} 0$.

Part (ii): Exponential rate. When $\rho \leq 1/2$, the contraction dominates the noise. By the Hájek–Rényi inequality for martingale difference sequences, $P(\sup_{m \geq k} |D_m| > \varepsilon) \leq \sigma_0^2 \rho^k / \varepsilon^2$. Setting $\varepsilon = \sigma_0 \rho^{k/2} \sqrt{\log k}$ and applying the Borel–Cantelli lemma yields the almost-sure exponential rate. $\square$

Applications: Theorem 3.1 provides the **architectural constraint** for self-auditing SCX systems. The condition $\rho < 1$ means each meta-audit layer must be strictly more capable than the layer it audits. In practice, this implies a **hierarchical audit architecture**: $\mathcal{A}_0$ (base AI auditor, σ large) $\rightarrow \mathcal{A}_1$ (human-aided auditor, σ smaller) $\rightarrow \mathcal{A}_2$ (expert panel, σ smallest). Each escalation reduces the meta-audit noise relative to the layer below, guaranteeing $\rho < 1$ and hence convergence. The exponential rate for $\rho \leq 1/2$ — achievable when each layer is at least twice as precise as the previous — provides a concrete design target for audit hierarchy depth. In autonomous systems where self-audit is safety-critical (self-driving perception stacks), this theorem mandates that the meta-auditor use strictly higher-resolution sensors or strictly more data than the primary auditor — a “sensor hierarchy” that mirrors the audit hierarchy. The convergence guarantee also enables automated audit termination: stop the recursion when $|\hat{\eta}_k - \hat{\eta}_{k-1}| < \epsilon$, with the exponential rate ensuring this occurs within $O(\log(1/\epsilon))$ layers.

Remark 3.2. Theorem 3.1 is **rigorous**. It formalizes the condition under which the Audit Sword can safely point at itself: the meta-auditor must be strictly higher-fidelity than the auditor it evaluates. *[Rigorous]*

3.2 Infinite Regress Inevitability

Theorem 3.3 (Infinite Regress under $\rho \geq 1$). *If $\rho \geq 1$ (meta-audit noise does not strictly decay), then*

$$\lim_{k \rightarrow \infty} P(|\hat{\eta}_k - \eta^*| < \varepsilon) = 0, \quad (4)$$

for any $\varepsilon > 0$ and any putative fixed point η^ . The recursive audit sequence **almost surely diverges** — the Audit Sword breaks.*

Proof. When $\rho \geq 1$, the noise does not decay across levels. The sequence $\{D_k\}$ is a random walk with non-diminishing step variance (or increasing variance when $\rho > 1$). By Pólya’s recurrence theorem for random walks on $\mathbb{R}$, a one-dimensional random walk with i.i.d. increments of finite non-zero variance is recurrent (returns infinitely often to any neighborhood of the origin) but does *not* converge — it fluctuates without bound in the limit. Specifically, for $\rho = 1$ and independent increments, $\limsup_{k \rightarrow \infty} |\hat{\eta}_k| = \infty$ almost surely. For $\rho > 1$, the variance grows geometrically, making divergence even more rapid.

In the audit context, this means the confidence intervals $[\hat{\eta}_k - \epsilon_k, \hat{\eta}_k + \epsilon_k]$ expand without bound as $k \rightarrow \infty$, and no finite level of recursion produces a definitive estimate. $\square$

Applications: Theorem 3.3 provides the **prohibition on flat audit hierarchies**. The practical mandate is absolute: **never deploy a self-audit system where the auditor and meta-auditor have equal fidelity** ($\rho \approx 1$). A single AI system auditing itself — no matter how sophisticated — will produce an audit sequence that diverges rather than converges. This is the audit analogue of “a judge cannot judge their own case”: the information-theoretic structure of self-reference guarantees non-convergence. The critical phase transition at $\rho = 1$ provides a diagnostic: if two consecutive meta-audit layers produce estimates with equal confidence interval widths ($\epsilon_k \approx \epsilon_{k+1}$), the system is at the critical point and must be restructured. In regulatory contexts, this theorem provides the legal basis for mandatory independent audit: a financial institution’s internal audit cannot be the final arbiter of its own compliance — an external auditor with strictly higher fidelity ($\rho < 1$) is mathematically necessary for convergence to a reliable conclusion. The random-walk behavior at $\rho = 1$ means that, over time, the audit estimate wanders arbitrarily far from the truth, potentially certifying a severely mislabeled dataset as clean.

Remark 3.4. *The critical point $\rho = 1$ defines an **Audit Self-Reference Phase Transition**:*

- $\rho < 1$: *Audit-convergent phase* — recursion converges to fixed point;
- $\rho = 1$: *Critical point* — marginal divergence (random walk);
- $\rho > 1$: *Audit-divergent phase* — recursion explodes.

The practical implication is devastating for naive self-audit: if you use the same AI system to audit itself ($\rho \approx 1$), the recursion does not converge. **Equal-fidelity systems cannot audit equal-fidelity systems.** [Rigorous]

3.3 Fixed-Point Information-Theoretic Identity

Conjecture 3.5 (Fixed-Point — Audit Heat Death Equivalence). *The recursive audit fixed point η^* is identical to the audit heat death entropy $H_{\min} = H(\varepsilon \mid S)$ from AE-Theorem:*

$$\eta^* = H_{\min} = H(\varepsilon \mid S). \quad (5)$$

If true, this means: **the convergence point of infinite recursive auditing is exactly Theorem 3’s irreducible ignorance.** The Audit Sword, pointed at itself, ultimately strikes the shield it was forged to penetrate.

Remark 3.6. *The conjecture has strong directional support:*

- (i) *Both η^* and $H_{\min}$ are fixed points of information-theoretic operators (recursive audit operator f vs. conditional entropy under $\mathcal{I}_A^{(\infty)}$);*
- (ii) *Both represent the minimum achievable uncertainty about noise given all available SCX information;*
- (iii) *In the special case where the recursive audit operator is linear ($f(\eta) = a\eta + b$), the fixed point $\eta^* = b/(1 - a)$ can be explicitly matched to $H_{\min}$ under appropriate parameterization.*

A general proof — establishing f as the audit-information projection operator onto the σ -algebra generated by SCX observables — is **open**. [Open]

3.4 Gödel Analogy

Open Problem 3.7 (SCX Gödel Correspondence). *Does there exist a formal correspondence:*

$$“\rho \geq 1 \text{ in RA-Theorem}” \longleftrightarrow “\text{Existence of an undecidable proposition in PA}”, \quad (6)$$

such that recursive audit divergence implies the existence of an **absolutely unauditabile audit proposition** — a statement about audit quality whose truth value cannot be determined by any finite iteration of the SCX Audit Sword? [\[Open\]](#)

Remark 3.8. *The analogy is philosophically compelling but formally unproven. In Gödel’s construction, the undecidable proposition is “I am not provable” — a self-referential statement. In SCX, the corresponding statement would be “this audit conclusion is not auditable to within precision ε .” Whether such a statement can be constructed within the formal language of SCX audit logic is an **open problem** at the intersection of mathematical logic and information theory.*

4 Discussion

4.1 The Audit Self-Reference Phase Diagram

ρ Range	Phase	Audit Sword Status	Guarantee
$\rho < 1/2$	Strongly convergent	Safe (exponential)	[Rigorous] (Theorem 3.1)
$1/2 \leq \rho < 1$	Weakly convergent	Safe (polynomial)	[Rigorous] (Theorem 3.1)
$\rho = 1$	Critical	Marginal (random walk)	[Rigorous] (Theorem 3.3)
$\rho > 1$	Divergent	Broken	[Rigorous] (Theorem 3.3)

4.2 Engineering the Fidelity Gap

The condition $\rho < 1$ requires that each meta-audit layer be *strictly* more capable than the layer it audits. This can be engineered through:

- (i) **Human escalation:** $\mathcal{A}_0$ = AI auditor (σ large), $\mathcal{A}_1$ = human-aided auditor (σ smaller), $\mathcal{A}_2$ = expert panel (σ smallest) — this naturally achieves $\rho < 1$;
- (ii) **Increasing data:** each meta-auditor has access to strictly more data than the auditor below it;
- (iii) **Stronger models:** each meta-auditor uses a strictly more capable model architecture.

The key insight: **a flat audit hierarchy ($\rho \approx 1$) cannot converge.** SCX deployment requires a *deep* audit hierarchy with strictly increasing fidelity per level.

4.3 Connection to Other Theorems

- **TS-Theorem:** TS’s Sprint self-audit limit (drift rate ceiling) is the temporal version of RA-Theorem’s ρ -condition; the drift rate v_{drift} corresponds to the per-step σ_k^2 in the recursive sequence.
- **AE-Theorem:** Conjecture ?? unifies η^* with $H_{\min}$ — the recursive convergence point with the thermodynamic heat death.
- **HC-Theorem:** Human intervention corresponds to injecting a layer with $\rho \ll 1$ into the audit hierarchy, breaking the equal-fidelity deadlock.
- **FA-Theorem:** Federated audit hierarchies require $\prod_k \rho_k < 1$ across all federated layers for global convergence.

5 Conclusion

RA-Theorem confronts the Audit Sword with its own reflection. The answer is nuanced:

- (1) When meta-auditors are strictly higher-fidelity ($\rho < 1$), recursion converges — the Sword holds. [\[Rigorous\]](#)
- (2) When meta-auditors are equal or lower fidelity ($\rho \geq 1$), recursion diverges — the Sword breaks. [\[Rigorous\]](#)
- (3) The convergence point is conjectured to be Theorem 3’s barrier — $H_{\min}$, the irreducible ignorance. [\[Open\]](#)
- (4) The Gödel correspondence — whether divergent recursion implies absolutely unauditably propositions — is a foundational open problem. [\[Open\]](#)

The practical mandate is clear: **never deploy a self-audit system where the auditor and meta-auditor have equal fidelity.** The Audit Sword requires a hierarchy, not a circle.

References

- [1] SCX Working Group. “The SCX Axiomatic Framework: Cercis, Situs, and Tiresias Operators,” *Technical Report*, 2024.

- [2] K. Gödel. “Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I,” *Monatshefte für Mathematik und Physik*, 38:173–198, 1931.
- [3] A. M. Turing. “On computable numbers, with an application to the Entscheidungsproblem,” *Proceedings of the London Mathematical Society*, s2-42(1):230–265, 1937.
- [4] S. Banach. “Sur les opérations dans les ensembles abstraits et leur application aux équations intégrales,” *Fundamenta Mathematicae*, 3:133–181, 1922.
- [5] J. L. Doob. *Stochastic Processes*. Wiley, 1953.
- [6] G. Pólya. “Über eine Aufgabe der Wahrscheinlichkeitsrechnung betreffend die Irrfahrt im Straßennetz,” *Mathematische Annalen*, 84:149–160, 1921.
- [7] S. P. Meyn and R. L. Tweedie. *Markov Chains and Stochastic Stability*, 2nd ed. Cambridge University Press, 2012.
- [8] J. Hájek and A. Rényi. “Generalization of an inequality of Kolmogorov,” *Acta Mathematica Academiae Scientiarum Hungaricae*, 6(3–4):281–283, 1955.
- [9] R. M. Smullyan. *Gödel’s Incompleteness Theorems*. Oxford University Press, 1992.
- [10] D. R. Hofstadter. *Gödel, Escher, Bach: An Eternal Golden Braid*. Basic Books, 1979.
- [11] T. M. Cover and J. A. Thomas. *Elements of Information Theory*, 2nd ed. Wiley, 2006.
- [12] C. Villani. *Optimal Transport: Old and New*. Springer, 2008.