

Deep Mathematical Connections of the SCX Theory

Date: 2026-06-28 | **Type:** Theoretical analysis **Status:** Contains both proven connections and speculative (marked “conjectural”) extensions **Purpose:** Situate SCX within the broader landscape of pure and applied mathematics

Table of Contents

1. What Mathematical Domain Does SCX Belong To?
 2. What Is the Core Mathematical Object?
 3. Linear Algebra Connections
 4. Group Theory Connections
 5. Topology and Geometry Connections
 6. Category Theory Connections
 7. Analogous Mathematical Structures
 8. Is SCX a Special Case of Something Bigger?
 9. What Is “SCX” Mathematically?
 10. Summary Table
-

1. What Mathematical Domain Does SCX Belong To?

1.1 The Surface Answer

At first glance, SCX draws on:

Area	Role	Theorems
Probability (concentration)	Hoeffding/Chernoff bounds for FPR/FPR control	Thm 1
Information theory	Mutual information, Fano inequality, Pinsker	Thm 2
Statistical decision theory	Minimax lower bounds, hypothesis testing	Thm 3, Thm 4'
Large deviations theory	Cramer's theorem, Bahadur-Rao exact asymptotics	Thm 4'
High- dimensional statistics	k-means consistency, sub-Gaussian concentration	Thm 5, Prop 6
Empirical process theory	Uniform convergence, covering numbers	Thm 5

A reviewer might classify it as “**theoretical machine learning**” or “**statistical learning theory**” – but that is a field label, not a mathematical

domain.

1.2 The Deeper Answer

SCX is fundamentally a problem in structured statistical decision theory with information constraints.

The reason goes beyond which tools are used. The SCX theory studies the following irreducible structure:

Given: (i) M conditionally independent Bernoulli observations (expert errors) per sample, (ii) a latent stratification variable S (state), and (iii) a feature representation $\phi(X)$ that partially reveals S . Goal: Decide whether each sample's label is corrupted, minimizing F1 risk. Result: The optimal decision rule is thresholding the mean of the Bernoulli observations, with the threshold adaptively shifted by $O(1/M)$ to balance FPR and FNR, achieving the exact minimax constant governed by Chernoff information.

This is a **decision problem with a nuisance parameter structure (the state) and an information-constrained side channel (the features)**. No single subfield of statistics claims this exact structure as its own.

1.3 The “Home” Field

If forced to assign a single home field, the strongest claim is **Statistical Decision Theory** (Wald paradigm), because:

- Theorems 1 and 4' are about the optimality/rate-optimality of a specific

decision rule (thresholded consensus score).

- Theorem 2 is about the fundamental limits of decision making under information constraints (by mutual information on feature representations).
- Theorem 3 is about the non-identifiability of two data-generating processes from observable data – a classic decision-theoretic concept (Le Cam deficiency / statistical equivalence).
- All revolve around the minimax framework: worst-case risk over a parameter space.

Second home: Large Deviations Theory. The transition from Theorem 1 (Hoeffding rate $\exp(-2M \Delta^2)$) to Theorem 4' (exact Chernoff rate $\exp(-M \kappa)$ with $\kappa < 2 \Delta^2$) is precisely the passage from a suboptimal concentration inequality to the true large-deviations rate. The Bahadur-Rao theorem provides the exact prefactor. Theorem 4' is, at its heart, a large-deviations refinement of a decision-theoretic bound.

Third home: Information theory. The chain of inequalities in Theorem 2 (mutual information $\rightarrow$ Pinsker $\rightarrow$ TV $\rightarrow$ risk) is an information-theoretic argument. The Chernoff information $\kappa = C(\text{Bern}(p_0), \text{Bern}(p_1))$ is an information-theoretic divergence. However, information theory typically studies communication, not decision-making under latent structure.

Verdict: SCX lives at the intersection of **decision theory, large deviations, and information theory**, with no single field able to claim it entirely. This is precisely what makes it novel.

2. What Is the Core Mathematical Object?

2.1 The Surface Answer

The core object is the **consistency score**:

$$C(x) = \frac{1}{M} \sum_{m=1}^M \mathbf{1}\{f_m(x) \neq y\}$$

A sample mean of M conditionally independent Bernoulli indicators, analyzed as a test statistic for the hypothesis:

$$H_0 : \text{label is clean} \quad \text{vs.} \quad H_1 : \text{label is noisy}$$

under the F1 risk function.

2.2 The Deeper Object

The core mathematical object is a **triple**:

$$(\mathcal{X}, \mathcal{F}_M, \Pi)$$

where: - $\mathcal{X}$ is the input space (a measurable space) - $\mathcal{F}_M = \{f_1, \dots, f_M\}$ is an ensemble of M independent experts (functions $\mathcal{X} \rightarrow \mathcal{Y}$) - $\Pi : \mathcal{X} \rightarrow \mathcal{S}$ is a latent stratification (a measurable partition)

The theory studies the **detectability** of a contamination in the label distribution, where:

- The contamination is a uniform mixture over incorrect labels (rate η)

- The detectability is governed by two parameters per stratum:
 - $p_0(s) = \mu_s$: clean error rate in state s
 - $p_1(s) = 1 - C_{\text{bal}} \cdot \mu_s / (K_y - 1)$: noise error rate in state s
- The separation $\Delta_s = p_1(s) - p_0(s)$ determines the difficulty

2.3 Is It a Special Case of a Known Structure?

Yes, partially:

1. **C(x) as a U-statistic:** The consistency score is a U-statistic of order 1 (degree 1 kernel $h(e) = e$) with kernel size M . Standard U-statistic theory gives asymptotic normality, but SCX needs the non-asymptotic exponential tail, which is why Hoeffding (for U-statistics) enters.
2. **The detection problem as a two-sample testing problem:** Within each state, the problem reduces to testing whether M observations $e_1, \dots, e_M$ come from $\text{Bern}(p_0)$ or from a mixture over classes of $\text{Bern}(p_1^c)$. This is a **composite vs. composite hypothesis test** with a nuisance parameter (the noise label c).
3. **The minimax result as a Bayes-optimal test with calibrated prior:** Theorem 4' shows that the optimal threshold θ_{dagger} corresponds to the Bayes decision rule under prior $(\eta, 1-\eta)$ for the two hypotheses. The $O(1/M)$ threshold shift is the finite- M correction to the log-odds ratio.

2.4 Can the Entire Theory Be Expressed as a Consequence of a Deeper Theorem?

Not a single one, but a chain:

- Theorem 1 = **Hoeffding's inequality** + **union bound** + **F1 algebra**
- Theorem 2 = **Data processing inequality** + **Pinsker** + **Fano** + **TV stability of performance metrics**
- Theorem 3 = **Mixture model equivalence** + **minimax argument on indistinguishable distributions**
- Theorem 4' = **Bahadur-Rao theorem** + **Neyman-Pearson lemma** + **F1 expansion** + **saddlepoint approximation**
- Theorem 5 = **Empirical process theory** + **k-means landscape analysis** + **sub-Gaussian concentration**
- Proposition 6 = **Clustering stability theory** (Shamir & Tishby)

Deeper claim: The entire theory is a **consequence of the large deviations principle for Bernoulli sums**, applied in a structured setting. Specifically:

- Theorem 1 uses the coarse (Hoeffding) bound for the LDP rate
- Theorem 4' refines to the exact LDP (Cramer/Bahadur-Rao)
- Theorem 2 uses the complementary regime (mutual information small $\rightarrow$ states unrecoverable $\rightarrow$ SCX degrades), which is also a consequence of an LDP for the feature distribution
- Theorem 3 shows that even with infinite data, the LDP rate cannot distinguish two generative models

This is the deepest unifying structure: **everything in SCX is about the tension between what is detectable (via concentration of expert consensus) and what is not (due to feature weakness or model equivalence)**.

3. Linear Algebra Connections

3.1 The Consistency Score as a Linear Functional

Let $\mathbf{e}(x) = (e_1(x), \dots, e_M(x))^T \in \{0, 1\}^M$ be the vector of expert error indicators for sample x . Then:

$$C(x) = \frac{1}{M} \mathbf{1}^T \mathbf{e}(x)$$

where $\mathbf{1} = (1, 1, \dots, 1)^T$.

Interpretation: $C(x)$ is the projection of $\mathbf{e}(x)$ onto the direction $\mathbf{1}/\sqrt{M}$ – the normalized all-ones vector – followed by scaling back by $1/\sqrt{M}$. That is:

$$C(x) = \frac{1}{\sqrt{M}} \cdot \langle \mathbf{e}(x), \frac{\mathbf{1}}{\sqrt{M}} \rangle$$

The kernel of the linear map $\mathbf{e} \mapsto C$ is the $(M - 1)$ -dimensional subspace:

$$\ker(C) = \{\mathbf{e} \in \mathbb{R}^M : \sum_{m=1}^M e_m = 0\}$$

This is the hyperplane of error patterns with zero net consensus.

Connection to ANOVA: This is a **between-expert sum of squares** decomposition. Let $\bar{e} = C(x)$. Then the total variation in expert errors is:

$$\sum_{m=1}^M (e_m - \bar{e})^2 = \|\mathbf{e} - \bar{e}\mathbf{1}\|^2 = \|\mathbf{e}\|^2 - M\bar{e}^2$$

The first term $\|\mathbf{e}\|^2$ counts how many experts erred; the second $M\bar{e}^2$ is the “explained” variation captured by the consensus.

3.2 k-means State Discovery as Matrix Factorization

The k-means objective minimized during state discovery is:

$$W_n = \sum_{s \in \mathcal{S}} \sum_{i: \hat{s}(x_i)=s} \|\phi(x_i) - \hat{\mu}_s\|^2$$

Theorem (Ding & He, 2004): This is equivalent to PCA on the Gram matrix $K = \Phi\Phi^T$ where Φ is the $N \times d_\phi$ feature matrix.

Specifically, let $H = [h_1 | \dots | h_K]$ be the $N \times K$ cluster membership matrix (with $h_{ik} = 1/\sqrt{n_k}$ if x_i belongs to cluster k). Then:

$$W_n = \text{tr}(K) - \text{tr}(H^T K H)$$

and minimizing W_n is equivalent to maximizing $\text{tr}(H^T K H)$ – the projection of K onto the subspace spanned by the columns of H . The continuous relaxation of this (allowing H to be any $N \times K$ orthogonal matrix) gives PCA: H is the top K eigenvectors of K .

Connection to spectral clustering: The normalized graph Laplacian of a similarity graph built from $\phi(x)$ is:

$$L = I - D^{-1/2} W D^{-1/2}$$

where $W_{ij} = \kappa(\phi(x_i), \phi(x_j))$ (a kernel). The eigenvectors of L correspond to the spectral embedding, and k-means on this embedding produces the state

partition. Theorem 5's consistency result can thus be interpreted as: under strong separation, spectral clustering on the feature Gram matrix recovers the true state partition with exponentially decaying error.

3.3 Gauge Fixing in EGP (Paper IV) as Projection onto a Quotient

In the Expert Governance Protocol (Paper IV), the expert-level corrections (c_0, c_Z) transform under:

$$c_0 \rightarrow c_0 + g, \quad c_Z \rightarrow c_Z - g \quad \text{for any } g \in \mathbb{R}$$

Group structure: This is an action of the additive group $(\mathbb{R}, +)$ on $\mathbb{R}^{Z+1}$.

Let $\mathbf{c} = (c_0, c_1, \dots, c_Z)^T$. The gauge transformation is $\mathbf{c} \rightarrow \mathbf{c} + g \cdot \mathbf{v}$ where $\mathbf{v} = (1, 0, \dots, 0, -1)^T$ with +1 at position 0 and -1 at position Z.

Gauge condition: The EGP fixes the gauge by imposing:

$$\sum_Z \pi_Z c_Z = 0$$

Let $\pi = (\pi_0, \pi_1, \dots, \pi_Z)^T$. The condition is $\pi^T \mathbf{c} = 0$.

Projection interpretation: The gauge fixing is the orthogonal projection onto the subspace $\{\mathbf{c} : \pi^T \mathbf{c} = 0\}$ under the weighted inner product $\langle \mathbf{c}, \mathbf{c}' \rangle_\pi = \sum_Z \pi_Z c_Z c'_Z$.

The quotient space $\mathbb{R}^{Z+1}/\mathbb{R}\mathbf{v}$ is isomorphic to $\mathbb{R}^Z$, with the isomorphism given by the projection onto the constraint. This is a principal bundle with structure group $\mathbb{R}$ and base $\mathbb{R}^{Z+1}/\mathbb{R} \cong \mathbb{R}^Z$.

Conjectural – deeper group structure: If the experts are organized hierarchically (tree structure), the gauge group could be $\mathbb{R}^{\text{edges}}$ acting by adjusting corrections along parent-child edges. The gauge condition would correspond to fixing a section of this bundle.

3.4 F1 Bound as Matrix Inequality

The F1 lower bound in Theorem 1:

$$\text{F1} \geq 1 - \frac{1}{\eta} \sum_s \rho_s \exp(-2M\Delta_s^2)$$

cannot be naturally expressed as a standard matrix inequality (it is scalar and exponential). However, the **F1-to-error conversion** can be:

$$\begin{pmatrix} \text{FPR} & \text{FNR} \end{pmatrix} \cdot \begin{pmatrix} \frac{1-\eta}{2\eta} & \frac{1}{2} \end{pmatrix}^T \geq 1 - \text{F1}$$

This is a linear inequality relating the (2×1) vector of error rates to the scalar $(1 - \text{F1})$. If we define the confusion matrix for state s :

$$C_s = \begin{pmatrix} \text{TN}_s & \text{FP}_s \\ \text{FN}_s & \text{TP}_s \end{pmatrix}$$

then $\text{FPR}_s = \text{FP}_s / (\text{FP}_s + \text{TN}_s)$, $\text{FNR}_s = \text{FN}_s / (\text{FN}_s + \text{TP}_s)$, and the F1 bound can be expressed as a set of linear constraints on C_s after concentration.

Conjectural: The multi-state F1 bound could be written as:

$$\text{F1} \geq 1 - \frac{1}{\eta} \sum_s \rho_s \cdot \text{tr} \left(\begin{pmatrix} 0 & 0 \\ 0 & 1 \end{pmatrix} \cdot C_s \cdot \exp \left(-2M \begin{pmatrix} (\theta - \mu_s)^2 & 0 \\ 0 & (1 - C_{\text{bal}}\mu_s/(K-1) - \theta)^2 \end{pmatrix} \right) \right)$$

This is not a standard matrix inequality but a trace of a product involving the confusion matrix and an exponential damping matrix.

4. Group Theory Connections

4.1 The Gauge Group (Paper IV) – Proven Connection

As established in Section 3.3, the EGP gauge transformation is an action of $(\mathbb{R}, +)$ on the space of expert corrections. This is a **principal $\mathbb{R}$ -bundle**.

The gauge condition $\sum_Z \pi_Z c_Z = 0$ picks a **global section** of this bundle. The fiber over each equivalence class $[\mathbf{c}]$ is $\{g \cdot \mathbf{v} : g \in \mathbb{R}\}$, a one-dimensional affine subspace.

Conjectural – non-Abelian extension: If the corrections interact non-trivially (e.g., $c_0 \rightarrow c_0 + g$, $c_1 \rightarrow c_1 - g/2$, $c_2 \rightarrow c_2 - g/2$), the gauge group might be non-Abelian, e.g., $\mathbb{R}^k$ with a symplectic constraint. This would map to a more complex Lie group structure, possibly $\text{SO}(k, 1)$ for hierarchical expert structures.

4.2 The Symmetric Group S_M Acting on Experts – Partially Proven

Observation: Under Assumption A2 (conditional independence), the expert errors $e_1, \dots, e_M$ are exchangeable within each state if all experts share the same per-state error probability μ_s (i.e., experts are homogeneous within states).

Proven: The consistency score $C(x)$ is symmetric under any permutation $\sigma \in S_M$ of the experts:

$$C(x) = \frac{1}{M} \sum_m e_m(x) \xrightarrow{\sigma} \frac{1}{M} \sum_m e_{\sigma(m)}(x) = C(x)$$

Thus $C(x)$ is the **maximal invariant** of the S_M action on the space of expert error vectors $\{0, 1\}^M$.

Conjectural: The decomposition of the space of error vectors under S_M :

- The orbit of any $\mathbf{e} \in \{0, 1\}^M$ under S_M consists of all vectors with the same number of 1s (Hamming weight).
- These orbits are in bijection with $\{0, 1, \dots, M\}$, i.e., the possible values of $M \cdot C(x)$.
- The distribution of $C(x)$ under H_0 (clean) is $\text{Binomial}(M, p_0)/M$, which is **uniform over each orbit** (all patterns with k errors are equally likely).
- This S_M -symmetry is the deepest reason why **only the consensus score matters** for optimal detection – any S_M -invariant decision rule is a function of $C(x)$.

4.3 Automorphism Groups of the State Partition

Definition: Let $\Pi = \{s_1, \dots, s_K\}$ be a measurable partition of $\mathcal{X}$. The **automorphism group** $\text{Aut}(\Pi)$ is the set of bijections $T : \mathcal{X} \rightarrow \mathcal{X}$ such that:

1. T is measurable (preserves the σ -algebra)
2. T preserves Π setwise: for each i , $T(s_i) = s_j$ for some j (i.e., T permutes the states)
3. T preserves the probability measure: $P_X \circ T^{-1} = P_X$

If the permutation of states is always the identity (states are not interchangeable), we have the **pointwise stabilizer** $\text{Stab}(\Pi) \leq \text{Aut}(\Pi)$.

Conjectural: For a metric space $\mathcal{X}$ with a state structure, the automorphism group of the empirical partition (from k-means) converges to $\text{Aut}(\Pi)$ as $n \rightarrow \infty$, at a rate governed by Theorem 5's concentration bound. This would establish a **consistency of automorphism groups**, analogous to the consistency of the partition itself.

4.4 Galois-like Correspondence – Conjectural

Setup: Consider the following analogy:

Galois Theory	SCX
Field F	Global expert risk vector $\mathbf{R} = (R_1, \dots, R_M)$

Galois Theory	SCX
Extension field E	State-conditioned risk matrix $\{R_m(s)\}$
Galois group $\text{Gal}(E/F)$	Permutations of experts preserving global ranking
Subfields of E	Sub-partitions of the state space
Fixed field of subgroup H	Experts indistinguishable under state grouping

Conjectural correspondence: If we coarse-grain the state space (merge states), we get a coarser partition Π' . This corresponds to a **reduction** in the expert ranking information: some experts that were distinguished under Π become indistinguishable under Π' . This is analogous to a field extension E/F where taking the fixed field under a subgroup $H \leq \text{Gal}(E/F)$ results in a smaller field.

Caveat: This analogy is structural rather than formal. There is no known way to make the correspondence bijective (the “fundamental theorem of Galois theory” requires specific algebraic conditions not present in SCX).

5. Topology and Geometry Connections

5.1 State Space Decomposition as Topological Partition

The state partition $\Pi = \{s_1, \dots, s_K\}$ of $\mathcal{X}$ induces the **quotient space** $\mathcal{X}/\Pi$ with the quotient topology. The quotient map $s : \mathcal{X} \rightarrow \mathcal{X}/\Pi$ is continuous (by definition of the quotient topology) and its fibers are the states.

Proposition 6 (clustering stability) tests whether this quotient structure is identifiable from finite data: if the quotient is “stable” (high ARI across bootstrap resamples), the quotient topology is well-defined empirically.

Theorem 5 shows that under strong separation ($\Delta_{\min} > 0$), the empirical quotient converges to the true quotient with probability $1 - \exp(-cn_{\min}\Delta_{\min}^2/(\sigma^2d_\phi))$.

5.2 Theorem 3 as a Fiber Bundle – Proven Structure

Model: Let $\mathcal{P}$ be the space of data-generating processes satisfying the SCX assumptions. Let $\mathcal{Q}$ be the space of observable joint distributions over $(X, Y, \{f_m\})$.

Theorem 3 constructs two distinct processes $P_{\text{noise}}, P_{\text{hard}} \in \mathcal{P}$ that map to the same $Q \in \mathcal{Q}$.

Fiber bundle interpretation: The map $\pi : \mathcal{P} \rightarrow \mathcal{Q}$ (forgetting the internal structure) is surjective. The fiber $\pi^{-1}(Q)$ over any Q contains at least the two constructed processes. The fiber dimension (in the sense of parametric models) is at least 1: the parameter $\eta \in (0, 1)$ traces a curve within the fiber.

Conjectural: The fibers are not just two points but continuous manifolds.

For $K_y > 2$, the random-expert construction creates a family of processes parametrized by $(\eta, \varepsilon_1, \varepsilon_2, \dots)$ all mapping to the same observable distribution. The fiber is a simplex of dimension at least $K_y - 1$.

More formally, let the set of processes that are indistinguishable from P_{noise} be:

$$\mathcal{J}(P_{\text{noise}}) = \{P' \in \mathcal{P} : \forall \text{ measurable } A, P_{\text{noise}}(A) = P'(A)\}$$

Theorem 3 shows $|\mathcal{J}(P_{\text{noise}})| \geq 2$. The conjecture is that $\mathcal{J}(P_{\text{noise}})$ is a **compact convex set** in the space of probability measures, with extreme points corresponding to different allocations of “noise vs. hardness” across the ambiguous subset.

5.3 The Detection Boundary Δ_s as Geodesic Distance

On the space of Bernoulli distributions $\{\text{Bern}(p) : p \in (0, 1)\}$, the **Fisher information metric** is:

$$g(p) = \frac{1}{p(1-p)}$$

The geodesic distance under this metric is:

$$d(p_0, p_1) = \left| \int_{p_0}^{p_1} \frac{dp}{\sqrt{p(1-p)}} \right| = 2 |\arcsin \sqrt{p_1} - \arcsin \sqrt{p_0}|$$

This is the **Bhattacharyya angle** (or more precisely, twice the Bhattacharyya angle).

The **Chernoff information** $\kappa = C(\text{Bern}(p_0), \text{Bern}(p_1))$ satisfies:

$$\kappa = -\log \inf_{t \in (0,1)} \int p_0^t p_1^{1-t} d\mu = -\log (p_0^{\theta^*} p_1^{1-\theta^*} + (1-p_0)^{\theta^*} (1-p_1)^{1-\theta^*})$$

Relationship between κ and the geodesic distance:

$$\kappa \leq -\log \left(1 - \frac{H^2}{2} \right) \leq \frac{H^2}{2} \leq 1 - \sqrt{1 - H^2} \leq \frac{d(p_0, p_1)^2}{2}$$

where $H^2 = (\sqrt{p_0} - \sqrt{p_1})^2 + (\sqrt{1-p_0} - \sqrt{1-p_1})^2$ is the squared Hellinger distance. The Chernoff information is **bounded above by the geodesic distance squared** along the statistical manifold.

Connection to Δ_s : For small gaps, by Taylor expansion:

$$\kappa = \frac{(p_1 - p_0)^2}{2p_0(1-p_0)} + O((p_1 - p_0)^3)$$

Since $\Delta_s = \min(\theta - p_0, p_1 - \theta)$ and the optimal $\theta = (p_0 + p_1)/2$ for symmetric thresholds, we have $2\Delta_s^2 = (p_1 - p_0)^2/2$. Thus:

$$\kappa \approx \frac{2\Delta_s^2}{p_0(1-p_0)}$$

For small p_0 , $p_0(1-p_0) \approx p_0$, so $\kappa \gg 2\Delta_s^2$ – the Chernoff rate is significantly faster than the Hoeffding bound. This is exactly the numerical finding in Section 4.3 of the unified document (ratio 2.46-3.41).

5.4 Information Geometry of the Multi-State Problem

Conjectural: The full SCX model (multi-state, multi-expert) lives on a **product of statistical manifolds**:

$$\mathcal{M} = \times_{s \in \mathcal{S}} \mathcal{M}_s$$

where each $\mathcal{M}_s$ is a 2-dimensional manifold parametrized by $(p_0(s), p_1(s))$ with the Fisher metric. The global F1 risk is a function on this product manifold.

The minimax optimal threshold $\theta^\dagger$ corresponds to the point on the product manifold that minimizes the F1 risk. The $O(1/M)$ threshold shift corresponds to a **geodesic deviation** from the symmetric Chernoff point θ^* toward the Bayes-optimal decision boundary.

6. Category Theory Connections

6.1 State-Conditioned Expertise as a Functor – Conjectural

Define two categories:

- **DataCat:** Objects are pairs $(\mathcal{X}, P_X)$ of a measurable space and a probability measure. Morphisms are measure-preserving maps $h : \mathcal{X}_1 \rightarrow \mathcal{X}_2$ such that $P_{X_2} = P_{X_1} \circ h^{-1}$.
- **ExpertCat:** Objects are M -tuples $(\mathcal{F}, \mathcal{Y})$ where $\mathcal{F} = (f_1, \dots, f_M)$ with $f_m : \mathcal{X} \rightarrow \mathcal{Y}$. Morphisms are pointwise transformations.

Training functor $T : \text{DataCat} \rightarrow \text{ExpertCat}$: Given a dataset $D \sim \mathcal{D}^n$, $T(D) = (f_1, \dots, f_M)$ where each f_m is trained on a disjoint subset $D_m \subset D$ (Assumption A1).

State-conditioned restriction functor $R_s : \text{ExpertCat} \rightarrow \text{ExpertCat}$:

$$R_s(\mathcal{F})(x) = \mathcal{F}(x) \cdot \mathbf{1}\{x \in s\}$$

The SCX framework studies the interplay between T and $\{R_s\}_{s \in \mathcal{S}}$.

Conjectural: The state-conditioned expert risk satisfies a **functoriality** property:

If $h : \mathcal{X}_1 \rightarrow \mathcal{X}_2$ is a morphism in DataCat that preserves the state partition ($h(s_i) = s_i$ for all i), then:

$$R_m(s) \text{ computed on } \mathcal{X}_1 = R_m(s) \text{ computed on } \mathcal{X}_2$$

i.e., the risk is invariant under state-preserving transformations. This would make $R_m(s)$ a **functor from the category of state-marked probability spaces to the category of real numbers**.

6.2 Two-Layer Architecture as a Natural Transformation

The SCX two-layer architecture consists of:

- Layer 1: $\phi : \mathcal{X} \rightarrow \mathbb{R}^{d_\phi}$ (feature extractor)
- Layer 2: $\hat{s} : \mathbb{R}^{d_\phi} \rightarrow \mathcal{S}$ (state assigner, via k-means)

Consider the category **Prob** of probability spaces. Let $F : \text{Prob} \rightarrow \text{Prob}$ be the “feature extraction” functor:

$$F(\mathcal{X}, P_X) = (\mathbb{R}^{d_\phi}, P_\phi)$$

where $P_\phi = P_X \circ \phi^{-1}$. Let $G : \text{Prob} \rightarrow \text{Set}$ be the “state assignment” functor:

$$G(\mathcal{X}, P_X) = \mathcal{X}/\Pi$$

(the set of states under the true partition).

Conjectural: The two-layer architecture is a **natural transformation** $\eta : G \Rightarrow G \circ F$ if the following diagram commutes for all morphisms h in Prob:

$$\begin{array}{ccc} G(\mathcal{X}) & \xrightarrow{\quad _X \quad} & G(F(\mathcal{X})) \\ | & & | \\ G(h) & & G(F(h)) \\ \downarrow & & \downarrow \\ G(\mathcal{Y}) & \xrightarrow{\quad _Y \quad} & G(F(\mathcal{Y})) \end{array}$$

This would mean that the state assignment via features is consistent under transformations of the input space. This is not formally provable without additional assumptions, but it captures the intuition that the two-layer architecture should be “representation-agnostic” – the state assignment should depend only on the features, not on the representation of the input space.

6.3 Curation-Exploration Tradeoff as an Adjunction – Speculative

The SCX workflow cycle is:

1. **Curation:** $C : \text{Data} \rightarrow \text{CleanData}$ (filter out noisy samples via SCX detector)

2. **Exploration:** $E : \text{CleanData} \rightarrow \text{Data}$ (add new samples, e.g., via active learning)

Conjectural: (C, E) form an **adjunction** $C \dashv E$ in an appropriate 2-category of datasets:

The unit $\eta : \text{Id}_{\text{Data}} \Rightarrow E \circ C$ reflects that curating a dataset and then exploring adds back some samples (not the same ones).

The counit $\varepsilon : C \circ E \Rightarrow \text{Id}_{\text{CleanData}}$ reflects that exploring and then curating returns a superset of the original clean data.

Triangle identities:

$$\begin{aligned} C &\xrightarrow{C\eta} CEC \xrightarrow{\varepsilon C} C = \text{id}_C \\ E &\xrightarrow{\eta E} ECE \xrightarrow{E\varepsilon} E = \text{id}_E \end{aligned}$$

The first identity says: curating, then exploring+curating again, should not change the curation result. The second says: exploring, then curating+exploring again, should not change the exploration result.

Speculative: The failure of these identities (due to imperfect noise detection) measures the “distance” from a true adjunction, which is precisely the error probability $\exp(-2M\Delta^2)$ in Theorem 1.

7. Analogous Mathematical Structures

7.1 Arrow’s Impossibility Theorem – PROVEN ANALOGY, REMOVED

This analogy was explored and **removed** from the theory files. The mapping was:

Arrow	SCX
Voters i	States s
Candidates a	Experts f_m
Social preference ordering	Global ranking R_m
Pareto optimality	State-level ranking consistency
Independence of irrelevant alternatives	Independence of new experts

Why it was removed: The conditions don’t match. Arrow’s theorem requires all four conditions simultaneously to produce a contradiction. SCX’s Proposition 1 only requires the existence of a ranking crossing – a weaker condition. The Arrow analogy inflated the mathematical depth of the SCX result.

Verdict: False analogy. Notational only.

7.2 No Free Lunch Theorem – GENUINE ANALOGY

The No Free Lunch (NFL) theorem for optimization states: for any algorithm A , the average performance over all possible objective functions is the same as any other algorithm (including random search).

Parallel to SCX Theorem 2:

- NFL: Over all possible problems, no algorithm outperforms any other.
- SCX Thm 2: Over all feature representations with $I(\phi; S) \leq \delta$, no detector can outperform the baseline by more than $C_F \sqrt{\delta/2}$.

Both are **worst-case impossibility results** that establish fundamental limits. However, SCX’s result is quantitatively graded (the bound scales with $\sqrt{\delta}$), whereas NFL is binary (all algorithms are equivalent).

Key difference: NFL averages uniformly over all problems; SCX bounds the worst case over a restricted class (features with bounded mutual information). The SCX bound is tighter because the class is smaller.

7.3 Cramer-Rao Bound – GENUINE ANALOGY (already noted in theory)

The Cramer-Rao lower bound states: for any unbiased estimator $\hat{\theta}$ of a parameter θ ,

$$\text{Var}(\hat{\theta}) \geq \frac{1}{I(\theta)}$$

where $I(\theta)$ is the Fisher information.

Parallel to SCX Theorem 4’:

- C-R: Lower bound on estimator variance given by Fisher information.
- SCX Thm 4’: Lower bound on F1 error given by Chernoff information κ .

Both are **information-theoretic lower bounds** that take the form “no method can surpass this fundamental limit.” Both involve an **information**

measure (Fisher vs. Chernoff) that captures the difficulty of the underlying statistical problem.

However, there are important differences:

Aspect	Cramer-Rao	SCX Thm 4'
Risk	Variance of estimator	$1 - \text{F1}$ (misclassification rate)
Information	Fisher $I(\theta) = \mathbb{E}[(\partial_\theta \log p)^2]$	Chernoff $\kappa = C(P_0, P_1)$
Asymptotic form	$1/(nI(\theta))$	$C_{\min} e^{-M\kappa} / (\eta\sqrt{M})$
Achievability	MLE achieves CR bound under regularity	SCX adaptive threshold achieves κ constant
Regime	$n \rightarrow \infty$ with fixed p	$M \rightarrow \infty$ with fixed p_0, p_1

The C-R bound is polynomial in n , while SCX's bound is exponential in M . This is because C-R is about **estimation** (smooth parameter) while SCX is about **testing** (discrete hypothesis).

7.4 Gelfand-Naimark Theorem (C*-algebras) – SPECULATIVE ANALOGY

Gelfand-Naimark: Every commutative C -algebra is isometrically - isomorphic to $C_0(X)$ for some locally compact Hausdorff space X . This establishes a duality between commutative C*-algebras and topological spaces.

Conjectural analogy in SCX:

- The space of expert error patterns $\{0, 1\}^M$ with the probability distribution P_{clean} and P_{noise} can be seen as a C*-algebra (the algebra of functions on $\{0, 1\}^M$).
- The consistency score C is a **positive linear functional** on this algebra (a state in the C*-algebra sense).
- The SCX decision rule $\mathbf{1}\{C(x) > \theta\}$ is a **spectral projection** of the operator C onto eigenvalues $> \theta$.
- The state space $\mathcal{S}$ corresponds to a **decomposition of the GNS representation** of the algebra into irreducible components.

Is this deep or superficial? It is unlikely to be deep. The Gelfand-Naimark duality is a tool for non-commutative geometry, and the SCX setting is purely classical (commutative). The projection $\mathbf{1}\{C > \theta\}$ is a spectral projection of a classical random variable, not a quantum observable. The analogy is **notational** at best.

Verdict: Weak analogy. The C*-algebra framework adds no new insight.

7.5 Shannon's Source Coding Theorem – NUMERIC ANALOGY

Shannon's theorem: For a source X with rate-distortion function $R(D)$, the optimal compression at distortion D requires rate $R(D) = I(X; \hat{X})$, and there exist codes that achieve $R(D)$ as block length $\rightarrow \infty$.

Parallel to SCX Thm 4':

- Shannon: $R(D) = I(X; \hat{X})$ is the optimal rate; achieved by random coding.
- SCX Thm 4': $\kappa = C(P_0, P_1)$ governs the optimal error; achieved by adaptive threshold.

Both give **exact asymptotic constants** in terms of an information-theoretic divergence. Both involve a **coding/decision scheme** that achieves the bound asymptotically (random codes for Shannon; thresholded consensus for SCX).

However, the mathematical structures differ:

Aspect	Shannon	SCX
Fundamental quantity	Rate-distortion function $R(D)$	Chernoff information κ
Achievability proof	Random coding (existence)	Adaptive threshold (explicit construction)
Optimality proof	Converse: $R(D) \geq I(X; \hat{X})$	Minimax: $C_{\min}/\eta$ lower bound via Bahadur-Rao
Asymptotic form	$R(D) = I + O(1/\sqrt{n})$	$1 - F1 =$ $\frac{C_{\min}}{\eta\sqrt{M}}e^{-M\kappa} +$ $o(1/\sqrt{M})$

7.6 The Ising Model / Statistical Physics – CONJECTURAL ANALOGY

Conjectural: The SCX consensus score $C(x)$ behaves like a **magnetization** in a Curie-Weiss model with M spins.

- Expert m : spin $\sigma_m = 2e_m - 1 \in \{-1, 1\}$

- Consensus: $C = (\sum_m \sigma_m + M)/(2M)$
- Under H_0 (clean): $\mathbb{E}[\sigma_m] = 2p_0 - 1 < 0$ (ferromagnetic with weak field)
- Under H_1 (noise): $\mathbb{E}[\sigma_m] = 2p_1 - 1 > 0$ (ferromagnetic with strong field)

The SCX decision boundary $C = \theta$ corresponds to a **magnetization threshold**: samples with magnetization above threshold are classified as noise.

The **large deviations** of the magnetization (Theorem 4') correspond to the **free energy** of the Curie-Weiss model, where the rate function $I(\theta) = \text{KL}(\theta \| p)$ plays the role of the large-deviations rate for the empirical mean.

Possible extension: If experts have correlations (breaking A2), the spin system becomes an **Ising model** with pairwise interactions. The detection problem becomes harder, and the optimal test may involve the full spin configuration, not just the magnetization.

Verdict: Aesthetic analogy. The Curie-Weiss mapping is exact for independent experts but adds no technical simplification. It becomes potentially useful if A2 is relaxed (correlated experts).

8. Is SCX a Special Case of Something Bigger?

8.1 The Exponential Bound $\exp(-2M\Delta^2)$ as a Special Case of Hoeffding

Proven: The bound $\exp(-2M\Delta^2)$ is a special case of **Hoeffding's inequality** for sums of bounded independent random variables. With $e_m \in [0, 1]$, $\mathbb{P}(|\bar{e} - \mathbb{E}[\bar{e}]| \geq t) \leq 2 \exp(-2Mt^2)$.

The transition from Theorem 1 (Hoeffding rate) to Theorem 4' (Chernoff rate) is the transition from Hoeffding's inequality to **Cramer's theorem** (the full large deviations principle). This is a general phenomenon:

Setting	Hoeffding bound	Exact LDP (Cramer)
Bernoulli	$\exp(-2Mt^2)$	$\exp(-M \cdot \text{KL}(\theta\ p))$
Bounded r.v.	$\exp(-2Mt^2/(b-a)^2)$	$\exp(-M \cdot I(\theta))$ where I is the rate function
Sub-Gaussian	$\exp(-Mt^2/2\sigma^2)$	Sub-Gaussian rate function

Deeper structure: The exponential bound $\exp(-2M\Delta^2)$ is a **coarse estimate** for the true large-deviations rate. The **Pinsker inequality** $\text{KL}(\theta\|p) \geq 2(\theta-p)^2$ gives the relationship: the Hoeffding exponent is always weaker than the true KL exponent.

Thus Theorem 1 is a **corollary of Hoeffding's inequality plus a union bound**, and Theorem 4' is the **sharpening** obtained by replacing Hoeffding with Cramer + Bahadur-Rao.

8.2 Theorem 3 as a Special Case of General Unidentifiability

Proven: Theorem 3 is a special case of the **non-identifiability of mixture models**.

Consider a two-component mixture model:

$$P(y|x) = (1 - \eta)P_{\text{clean}}(y|x) + \eta P_{\text{noise}}(y|x)$$

The noise model $P_{\text{noise}}(y|x)$ is itself a mixture over the observed label $\tilde{y}$ given the true label y^* :

$$P_{\text{noise}}(y|x) = \frac{1}{K-1} \sum_{c \neq y^*} P(y = c | y^* \text{ flipped to } c)$$

The hardness-world is obtained by swapping the interpretation: instead of y^* being deterministic and the label random, we make y^* random and the label deterministic. This is a **reparameterization** of the mixture model.

General principle: Any time a latent variable model has a **symmetry** (here: swapping the roles of noise and ambiguity), the parameters are unidentifiable. This is closely related to:

- **Label switching** in Bayesian mixture models (the posterior is invariant under permuting component labels)
- **Identifiability of finite mixtures** (Titterton, 1985): a necessary condition for identifiability is that the component distributions are linearly independent
- **Hidden Markov model identifiability** (Allman et al., 2009): if emission distributions satisfy a “generic” condition, the HMM is identifiable up to label permutation

Theorem 3 as a concrete instance: The specific construction produces a **minimal** unidentifiability: exactly two data-generating processes map to the same observable distribution. This is the “atomic” unidentifiability from which larger families can be built by convex combination.

8.3 Theorem 4' as a Special Case of Exact Asymptotics Theory

Proven: Theorem 4'’s exact constant result is a special case of **second-order asymptotics for hypothesis testing** (also called “refined large deviations” or “moderate deviations”).

Specifically:

- The **Bahadur-Rao theorem** gives exact asymptotics for the sample mean of i.i.d. random variables.
- The **Neyman-Pearson lemma** shows that thresholded likelihood ratio tests are optimal.
- Theorem 4' combines these for Bernoulli observations under F1 risk.

General family: Theorem 4' belongs to the class of results that give **exact error constants for testing simple vs. simple hypotheses** under a specific loss function. The key elements are:

1. Reduction to Bernoulli (by conditional independence / sufficiency of $C(x)$)
2. Bahadur-Rao approximation of the tail probability
3. F1 risk expansion (first-order) to combine FPR and FNR
4. Adaptive threshold to balance the two contributions optimally
5. Matching lower bound via Neyman-Pearson + Cramer

Is it a special case of a known meta-theorem? Not exactly. While each step is standard, the **combination** of F1 risk with Bernoulli tests and adaptive thresholding is engineered. There is no general “exact minimax theory for F1 under mixture models” that Theorem 4' is a special case of.

Conjectural: The result could be unified under a **general theory of exact asymptotic minimax optimality for composite hypothesis testing**

under weighted error rates, where:

- The weights are given by the prior probabilities $(\eta, 1 - \eta)$
- The loss is a linear combination of FPR and FNR
- The optimal test is LRT with threshold adaptively shifted by $O(1/M) \log((1 - \eta)/\eta)$

Under this general theory, Theorem 4' would be the **Bernoulli special case** with $p_0 = \mu, p_1 = 1 - C_{\text{bal}}\mu/(K - 1)$.

8.4 Summary: What Larger Principles Does SCX Exemplify?

Theorem	Larger Principle	Status
Thm 1	Concentration of U-statistics + union bound	Proven special case
Thm 2	Data processing + Pinsker inequality	Proven special case
Thm 3	Mixture model non-identifiability	Proven special case
Thm 4'	Second-order asymptotics for hypothesis testing	Proven special case (each step), but the combination is novel
Thm 5	k-means consistency under strong separation	Proven special case
Prop 6	Clustering stability as feature informativeness	Proven heuristic

The novel synthesis: What makes SCX more than the sum of its parts is the **unified treatment** of all these regimes in a single theoretical framework, connected by a coherent set of assumptions (A1-A6). The theory shows how:

- Expert consensus concentration drives detection (Thm 1, 4')
- Feature informativeness bounds this detection (Thm 2)
- The state structure renders some distinctions impossible (Thm 3)
- Strong features enable state discovery (Thm 5)
- Clustering stability diagnoses this (Prop 6)

No existing theory captures this full chain.

9. What Is “SCX” Mathematically?

9.1 The One-Sentence Definition

SCX is the minimax-optimal statistical decision theory of detecting label contamination in a multi-expert system, where the experts’ conditional error probabilities are modulated by a latent state variable, the state is discovered via spectral clustering in a feature space with information-constrained fidelity, and the detection threshold is adaptively shifted by $O(1/M)$ to achieve the exact Chernoff constant under F1 risk.

9.2 More Precisely, in Pure Mathematics

Let us define the components formally:

- Let $(\Omega, \mathcal{F}, \mathbb{P})$ be a probability space.
- Let $X : \Omega \rightarrow \mathcal{X}$ be a random input, $Y : \Omega \rightarrow \mathcal{Y}$ be a random label, with $\mathcal{Y} = \{1, \dots, K\}$.
- Let $f_1, \dots, f_M : \mathcal{X} \rightarrow \mathcal{Y}$ be measurable functions (experts).
- Let $\Pi : \mathcal{X} \rightarrow \mathcal{S}$ be a measurable partition into $|\mathcal{S}| = K_{\mathcal{S}}$ states.
- Let $Z = \mathbf{1}\{Y \neq Y^*\}$ be the noise indicator, where Y^* is the true label.

Under A1-A6: - $f_m \perp f_{m'} \mid X$ (experts conditionally independent given input) - $\mathbb{P}(Z = 1) = \eta$ (noise rate) - $\mathbb{P}(Y = c \mid Z = 1, X) = 1/(K-1)$ for $c \neq Y^*$ (uniform noise) - $\mathbb{E}[\mathbf{1}\{f_m(X) \neq Y\} \mid Z = 0, X \in s] \leq \mu_s$ (state-bounded clean error) - $\mathbb{E}[\mathbf{1}\{f_m(X) \neq Y\} \mid Z = 1, X \in s] \geq 1 - C_{\text{bal}}\mu_s/(K-1)$ (state-bounded noise error)

Definition: The SCX noise detector is the function:

$$h_{\text{SCX}}(x) = \mathbf{1} \left\{ \frac{1}{M} \sum_{m=1}^M \mathbf{1}\{f_m(x) \neq y\} > \theta^\dagger \right\}$$

where $\theta^\dagger = \theta^* + \frac{1}{M} \frac{\log((1-\eta)/\eta)}{D^*} + O(1/M^2)$, $\theta^* = \arg \min_{\theta} \{\text{KL}(\theta \parallel \mu_s) = \text{KL}(\theta \parallel 1 - C_{\text{bal}}\mu_s/(K-1))\}$.

SCX theory: A chain of theorems establishing:

$$\lim_{M \rightarrow \infty} e^{M\kappa} \sqrt{2\pi M} (1 - \text{F1}(h_{\text{SCX}})) = \frac{C_{\min}}{\eta}$$

where $\kappa = C(\text{Bern}(\mu_s), \text{Bern}(1 - C_{\text{bal}}\mu_s/(K-1)))$ and $C_{\min}$ is the unique minimax constant. This holds globally across states (Thm 4') and degrades gracefully when features are weakly informative (Thm 2, $\text{F1} \leq \text{F1}_{\text{base}} + O(\sqrt{I(\phi; S)})$), while an unidentifiability gap prohibits perfect detection (Thm 3).

9.3 Alternative Characterizations

As a statistical decision theory: SCX is the characterization of the minimax-optimal Bayes decision boundary between two Bernoulli populations with state-dependent shifts, under a ratio-of-successes risk function, with the state discovered from side information.

As an information theory: SCX is the study of the Chernoff information between two product-Bernoulli distributions with a latent mixing parameter, and the data-processing loss incurred when the mixing parameter must be estimated from information-constrained features.

As a large deviations theory: SCX is the refinement of Cramer’s theorem for the sample mean of Bernoulli variables from exponential-rate asymptotics to exact constant asymptotics, applied to a structured decision problem with an adaptive threshold correction.

As an algebra: SCX is the study of the S_M -invariant function $C(x) = \frac{1}{M} \sum e_m$ as the maximal invariant of the expert ensemble, with the state partition providing a decomposition of $\mathcal{X}$ into regions where $C(x)$ has approximately constant distribution under both hypotheses.

10. Summary Table

Domain	Connection Strength	Key Insight	Theorems Involved
Statistical decision theory	PROVEN (core)	Minimax optimal testing under F1 risk; latent state structure	All 6
Large deviations	PROVEN (core)	Chernoff rate is exact; Bahadur-Rao gives constants	Thm 1, Thm 4'
Information theory	PROVEN (core)	Mutual information bounds SCX's advantage; Chernoff info	Thm 2, Thm 4'
Concentration inequalities	PROVEN (used)	Hoeffding/Chernoff bounds for FPR/FPR	Thm 1
U-statistics	PROVEN (weak)	$C(x)$ is a degree-1 U-statistic	Thm 1
Linear algebra (ANOVA)	PROVEN (weak)	$C(x) = (1/M)1^T e(x)$; projection onto all-ones direction	Thm 1

Domain	Connection Strength	Key Insight	Theorems Involved
Spectral cluster-ing (PCA)	PROVEN (medium)	k-means = PCA on Gram matrix; state discovery via eigenvectors	Thm 5, Prop 6
Gauge theory (R action)	PROVEN (medium)	EGP gauge fixing as principal R-bundle with section	Paper IV
Permutation group S_M	PROVEN (weak)	C(x) is S_M-invariant; maximal invariant of expert ensemble	Thm 1
Automorphic groups	Conjectural	Convergence of empirical automorphism group of partition	Thm 5
Galois correspondence	Conjectural	Coarse-graining of states subgroups of expert symmetries	—

Domain	Connection Strength	Key Insight	Theorems Involved
Fiber bundles	PROVEN (medium)	Thm 3: unidentifiable processes as fibers over observable	Thm 3
Information geometry	PROVEN (medium)	Fisher metric on $\text{Bern}(p)$; Chernoff $<$ geodesic distance	Thm 4'
Statistical physics (Ising)	Conjectural	Magnetization analogy for consensus; Curie-Weiss for independent	Thm 1
Category theory (functors)	Conjectural	Training as functor; state- conditioning as natural	—
Adjunctions	Speculative	Curation- Exploration as $C \dashv E$	—
Arrow's theorem	FALSE (removed)	Conditions don't match; misleading analogy	—

Domain	Connection Strength	Key Insight	Theorems Involved
No Free Lunch	GENUINE analogy	Worst-case bound over feature representations	Thm 2
Cramer-Rao bound	GENUINE analogy	Information-theoretic lower bound on estimation	Thm 4'
Gelfand-Naimark	Weak analogy	Spectral projection of consensus operator	Thm 1
Shannon source coding	GENUINE analogy	Exact asymptotic constant; achievability + converse	Thm 4'
BBP phase transition	PROVEN (medium)	Spectral transition for weak/strong feature regime	Thm 2, Thm 5

References

1. Bahadur, R. R., & Rao, R. R. (1960). On deviations of the sample mean. *Annals of Mathematical Statistics*, 31(4), 1015-1027.
2. Chernoff, H. (1952). A measure of asymptotic efficiency for tests of a hypothesis. *Annals of Mathematical Statistics*, 23(4), 493-507.
3. Cramer, H. (1938). Sur un nouveau theoreme-limite de la theorie des probabilites. *Actualites Scientifiques et Industrielles*.
4. Dembo, A., & Zeitouni, O. (2010). *Large Deviations Techniques and Applications* (2nd ed.). Springer.
5. Ding, C., & He, X. (2004). K-means clustering via principal component analysis. *ICML 2004*.
6. Hoeffding, W. (1963). Probability inequalities for sums of bounded random variables. *JASA*, 58(301), 13-30.
7. Le Cam, L. (1986). *Asymptotic Methods in Statistical Decision Theory*. Springer.
8. Pinsker, M. S. (1964). *Information and Information Stability of Random Variables and Processes*. Holden-Day.
9. Baik, J., Ben Arous, G., & Peche, S. (2005). Phase transition of the largest eigenvalue for nonnull complex sample covariance matrices. *Annals of Probability*, 33(5), 1643-1697.
10. Lei, J., & Zhu, L. (2018). A general spectral method for high-dimensional k-means clustering. *Annals of Statistics*, 46(6B), 3181-3216.
11. Von Luxburg, U. (2007). A tutorial on spectral clustering. *Statistics and Computing*, 17(4), 395-416.
12. Amari, S. (2016). *Information Geometry and Its Applications*. Springer.
13. Neyman, J., & Pearson, E. S. (1933). On the problem of the most

efficient tests of statistical hypotheses. *Philosophical Transactions of the Royal Society A*, 231, 289-337.

14. Wald, A. (1950). *Statistical Decision Functions*. Wiley.
15. SCX Unified Theorems Document. `THEOREMS_UNIFIED.md`
16. SCX Exact Constant Minimax. `exact_constant_minimax.md`
17. SCX Random Matrix Connection. `random_matrix_connection.md`
18. SCX Minimax Optimality Analysis. `minimax_optimality.md`

End of deep mathematical connections analysis.